

Binary Search Trees

Binary Search Trees (BSTs)

Fairly fast **search, insertion, deletion**

maximum, minimum, successor, predecessor

Creators: **Dumey (1952) Wheeler (1957)**

Berners-Lee (1959) Windley (1960)

Booth & Colin (1960) Hibbard (1962)

Complexity (for n data items)

Typically, binary search trees are **fairly fast**:

Avg. case: **search, insert, delete** $\rightarrow O(\log n)$

Worst case: **search, insert, delete** $\rightarrow \Theta(n)$

search, insert, delete $\rightarrow O(h)$ (h BST height)

Storage of binary search trees: $\Theta(n)$

What they are

How they are implemented

Suits

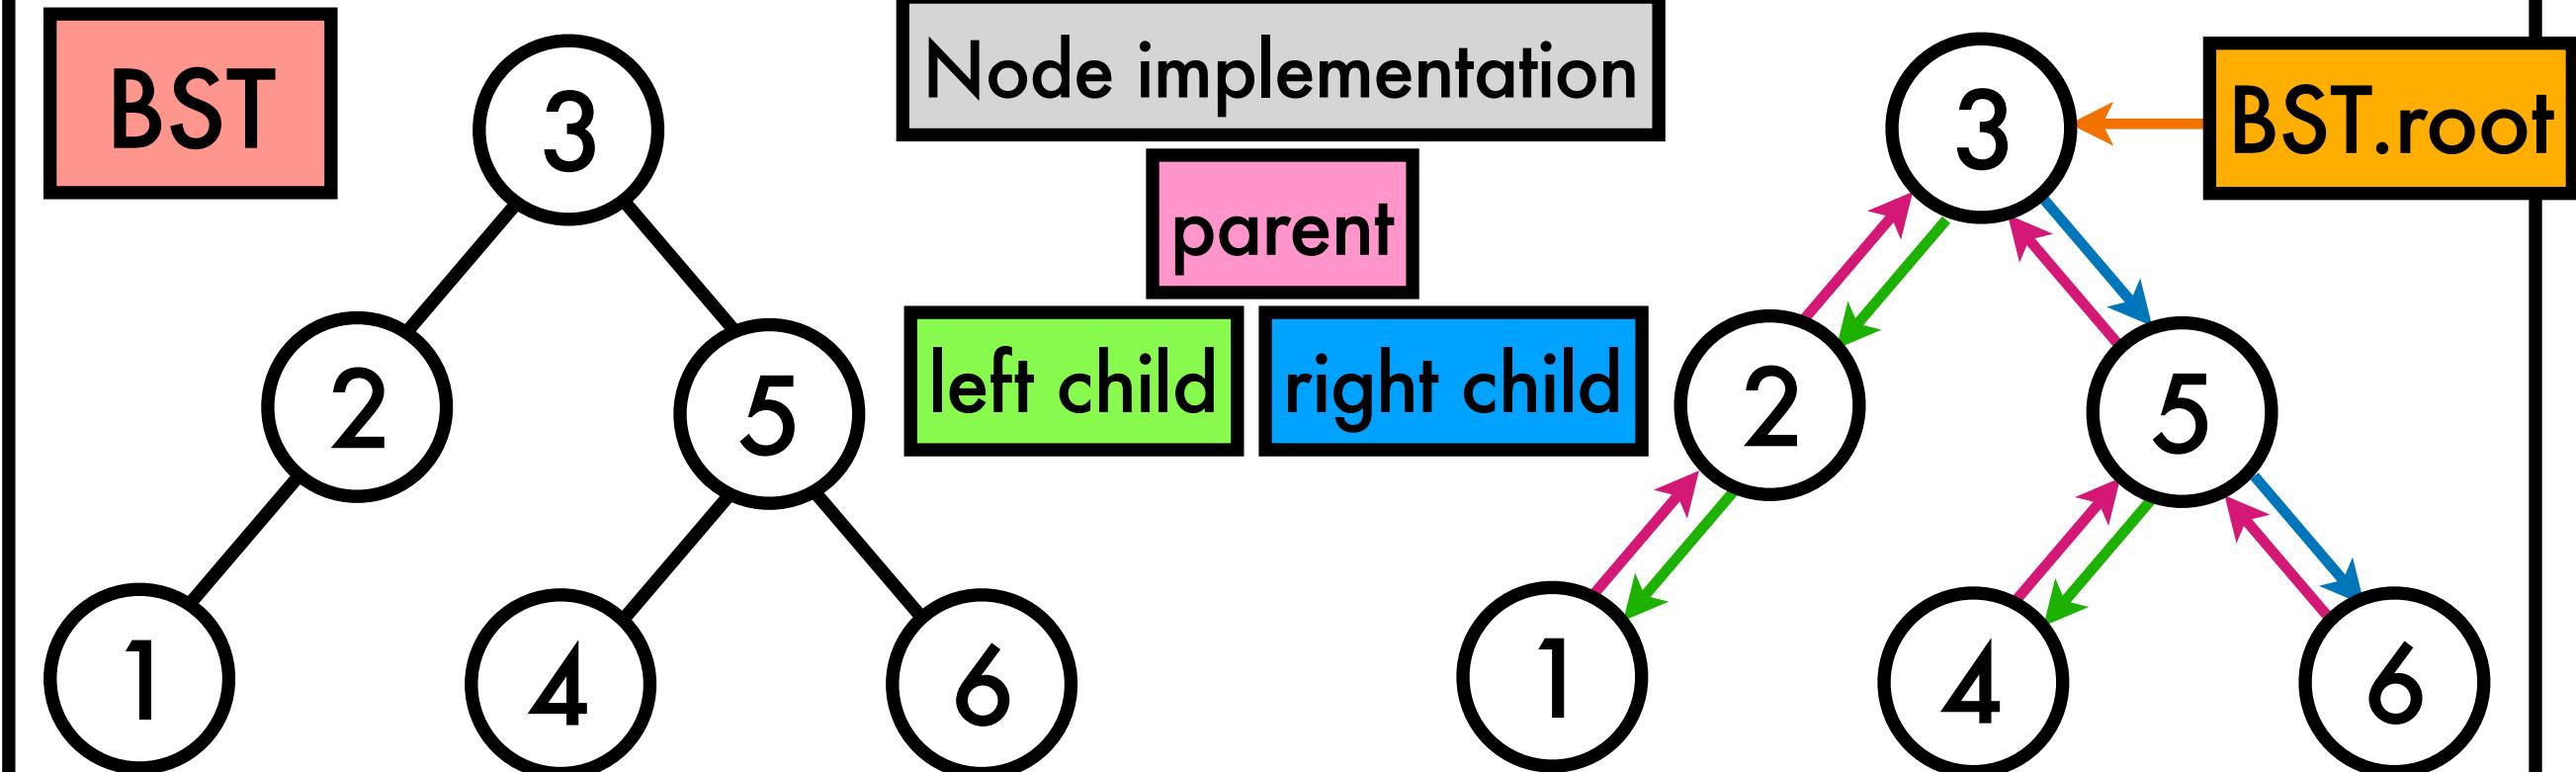
Abstract Data Types

Set

Map

Priority Queue

BST



Binary Search Tree Property: for each node u , any

node l in its **left subtree** satisfies $l.\text{key} \leq u.\text{key}$, any

node r in its **right subtree** satisfies $r.\text{key} \geq u.\text{key}$

Note: this definition allows **duplicate keys**

References/Notes/Image credits:

(History of Binary Search Trees) D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 6.2.2 (1974)

(Wheeler/Berners-Lee) A. Douglas, "Techniques for the recording of, and reference to data in a computer", The Computer Journal (1959)

(D. Wheeler) [https://en.wikipedia.org/wiki/David_Wheeler_\(computer_scientist\)#/media/File:EDSAC_\(14\)_cropped.jpg](https://en.wikipedia.org/wiki/David_Wheeler_(computer_scientist)#/media/File:EDSAC_(14)_cropped.jpg)

(C. Berners-Lee) <https://www.bl.uk/voices-of-science/interviewees/conway-berners-lee>

A. Booth and A. Colin, "On the efficiency of a new method of dictionary construction", Information and Control (1960)

(Booth) https://www.computerhope.com/people/andrew_booth.htm

(Andrew Colin) <https://www.heraldscotland.com/opinion/16998308.obituary-andrew-colin-professor-computer-science/>

T. Hibbard, "Some combinatorial properties of certain trees with applications to searching and sorting", JACM (1962)

(Hibbard) <http://math.oxford.emory.edu/site/cs171/hibbardDeletion/>

(Binary Search Trees) T. Cormen et al., "Introduction to algorithms", Chap 12.1, MIT press, (2022)

Tree traversals

Traversal algorithms

A **traversal algorithm** aims to "process" each node in the tree **exactly once**

The simplest traversals use **depth-first-search** (go **deeper first**, rather than "**breadth first**")

inorder

preorder

postorder

Inorder traversal: process each node **in-between** visiting **left subtree and right subtree**

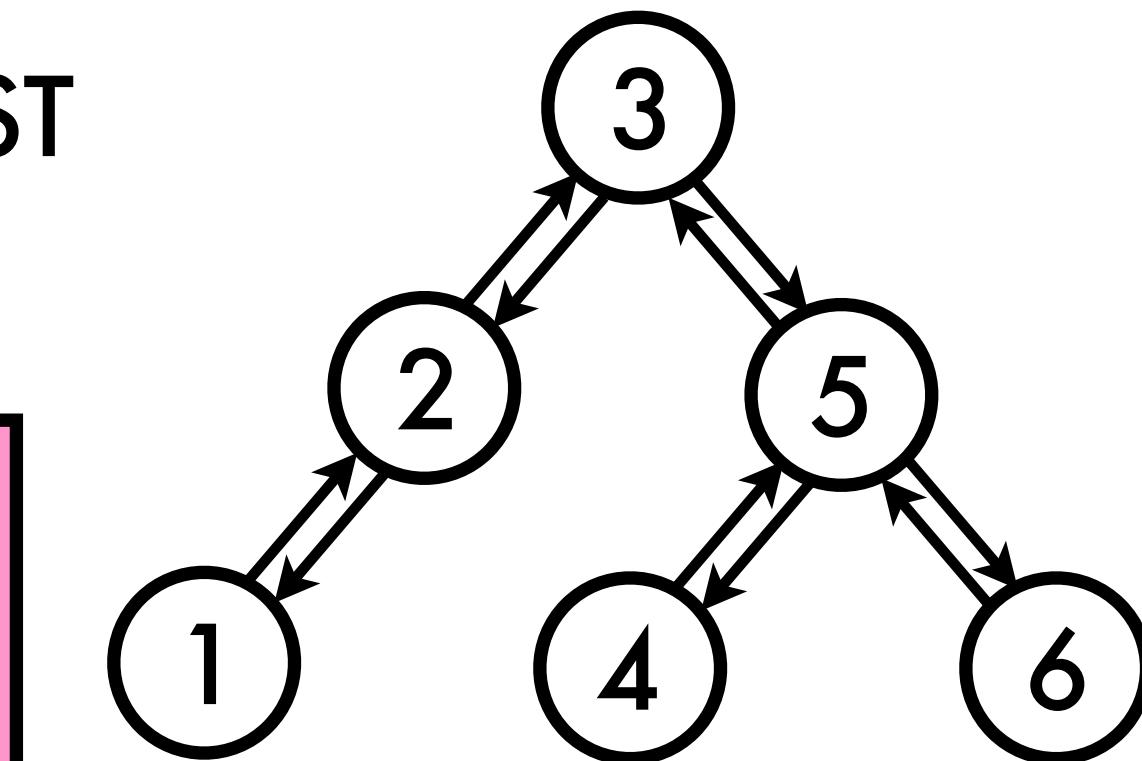
Recursive implementation:

```
def inorder(u):  
    if u is not None:  
        inorder(u.left)  
        print(u.key)  
        inorder(u.right)
```

Inorder traversal of BST

print out: 1 2 3 4 5 6

Note that keys were
printed in order!



```
def preorder(u):  
    if u is not None:  
        print(u.key)  
        preorder(u.left)  
        preorder(u.right)
```

```
def postorder(u):  
    if u is not None:  
        postorder(u.left)  
        postorder(u.right)  
        print(u.key)
```

print out: 3 2 1 5 4 6

print out: 1 2 4 6 5 3

Traversals are $\Theta(n)$ - call themselves **twice at each node** (left child and right child)

References/Notes/Image credits:

(Tree traversals) <http://webdocs.cs.ualberta.ca/~holte/T26/tree-traversal.html>

(Traversal complexity) T. Cormen et al., "Introduction to algorithms", Chap 12.1, MIT press, (2022)

Binary Search Tree Queries

Minimum/Maximum

```
def minimum(u):  
    while u.left:  
        u = u.left  
    return u
```

E.g. argument: root
returns node with key 1

```
def maximum(u):  
    while u.right:  
        u = u.right  
    return u
```

E.g. argument: root
returns node with key 9

Search

```
def search(u, key):  
    while u and key != u.key:  
        if key < u.key:  
            u = u.left  
        else:  
            u = u.right  
    return u
```

query key: 6

returns node

query key: 7

returns None

Complexity: $O(h)$ where h is tree height

Inorder Predecessor

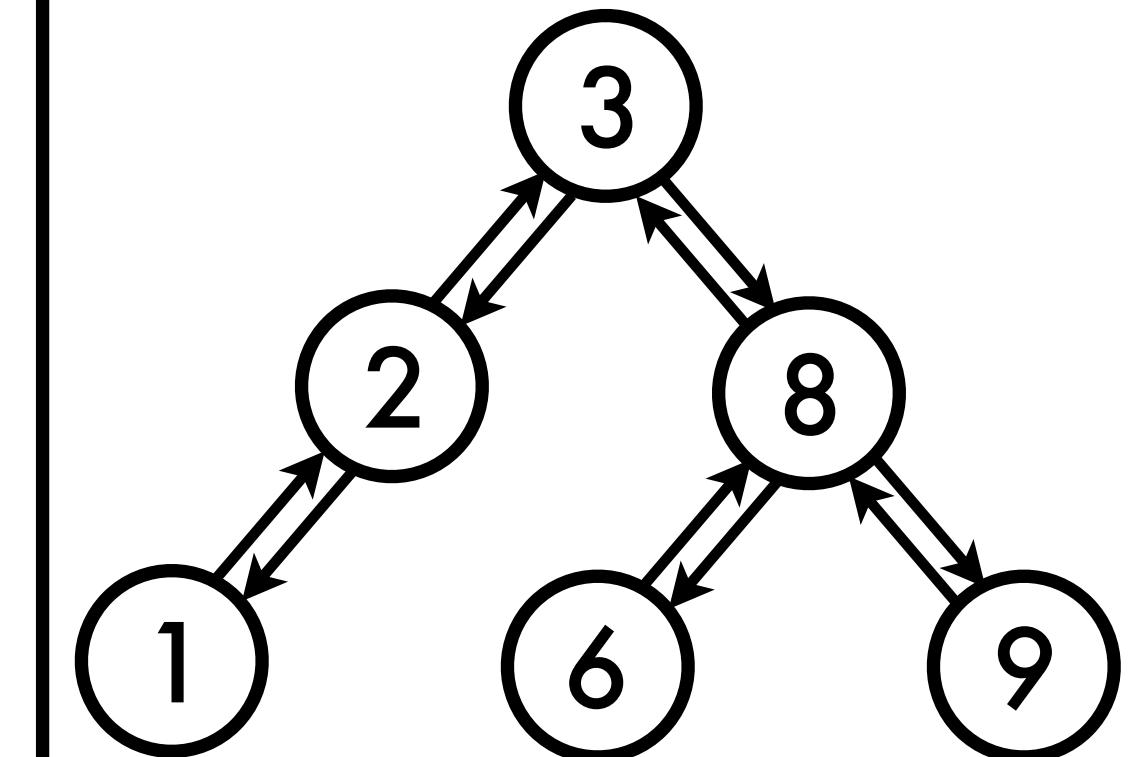
Element **immediately before** node in
the **inorder traversal ordering**

Note: no **key** required!

```
def predecessor(u):  
    if u.left:  
        return maximum(u.left)  
    else:  
        par = u.parent  
        while par and u != par.right:  
            u = par  
            par = par.parent  
        return par
```

Note: **inorder successor** is symmetric

Complexity: $O(h)$ (h is tree height)



Example Binary Search Tree

Predecessor examples

First case

query node with key 3

returns node with key 2

Second case

query node with key 6

returns node with key 3

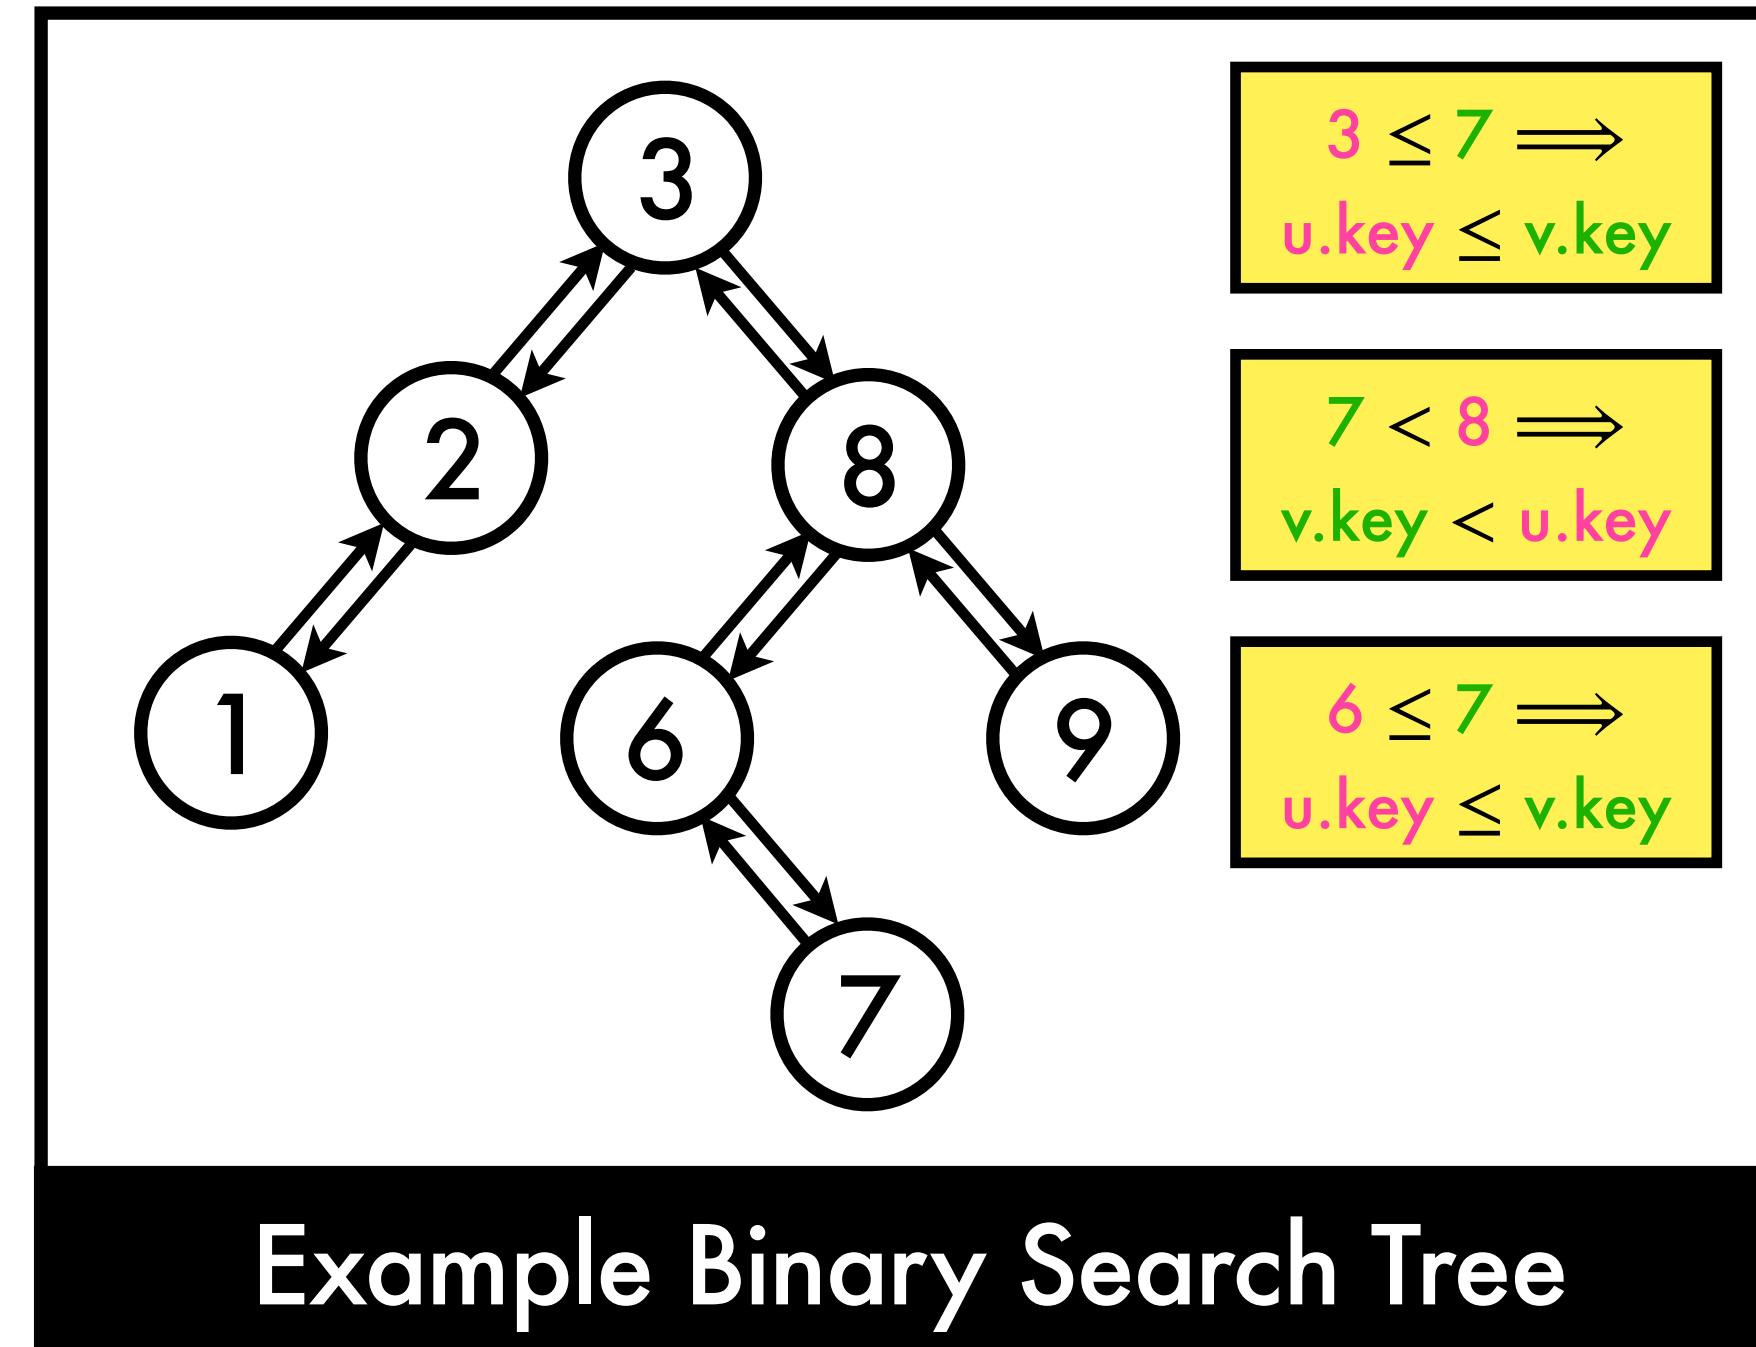
Binary Search Tree Insertion

Insertion

Insert new node v into binary search tree bst

```
def insert(bst, v):
    u = bst.root
    par = None
    while u:
        par = u
        u = u.left if v.key < u.key else u.right
    v.parent = par
    if not par: # handle case when bst was empty
        bst.root = v
    elif v.key < par.key:
        par.left = v
    else:
        par.right = v
```

Complexity: $O(h)$ where h is tree height



Example Binary Search Tree

Insert example node with key 7

We follow Cormen (BST allows duplicate keys)
If not allowed, `insert()` must be modified

References:

(insertion pseudocode) https://en.wikipedia.org/wiki/Binary_search_tree#Insertion

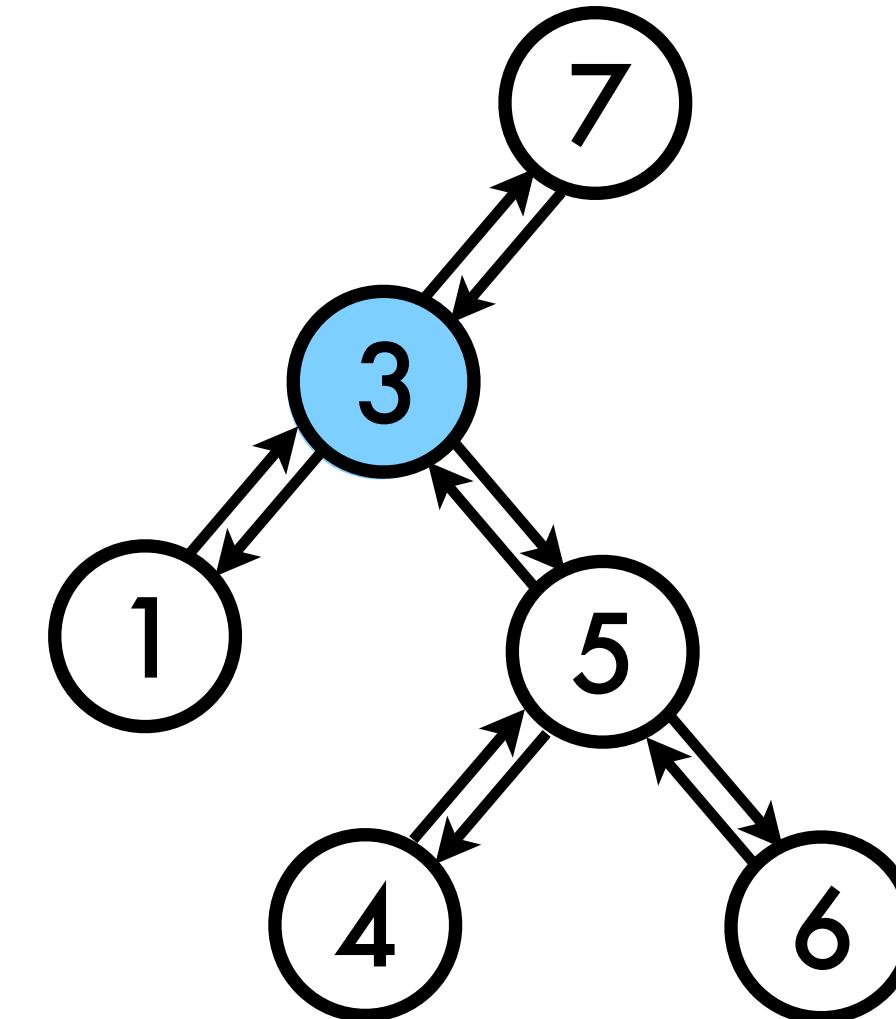
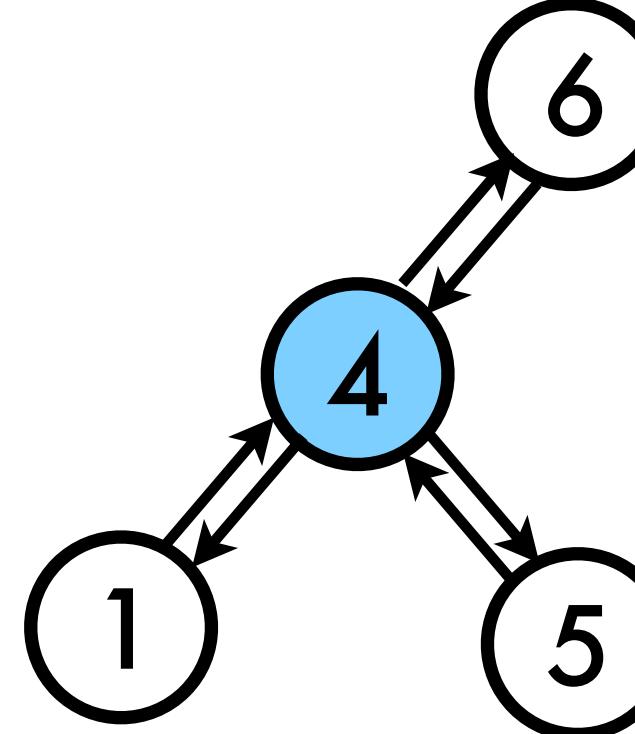
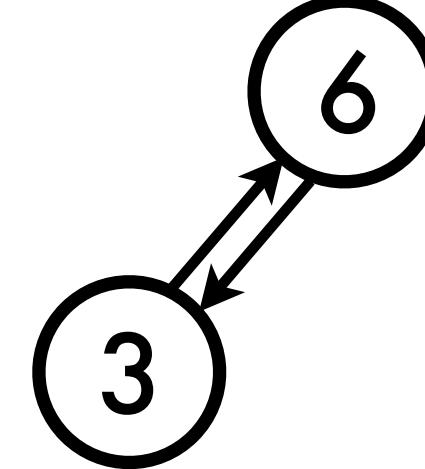
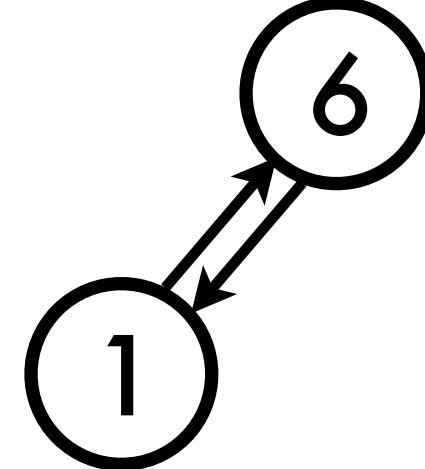
(Insertion) T. Cormen et al., "Introduction to algorithms", Chap 12.3, MIT press, (2022) (a.k.a CLRS due to the authors names)

Binary Search Tree Deletion

Deletion logic

To delete a node from BST, there are four cases to consider goal: preserve the BST property

In the following, we will focus on deleting the node with key 2



case 1:
no right child

case 2:
no left child

case 3: two children
successor is right child

case 4: two children
successor is not right child

References:

https://en.wikipedia.org/wiki/Binary_search_tree#Deletion

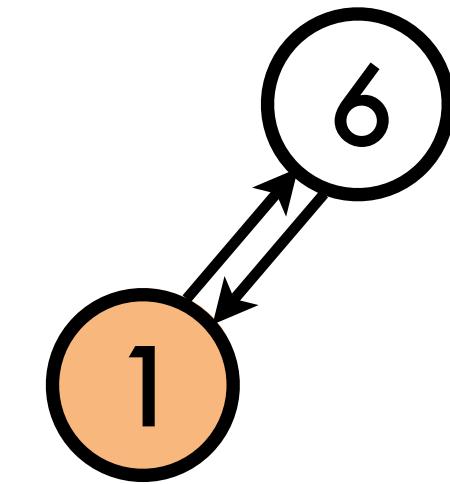
(Deletion) T. Cormen et al., "Introduction to algorithms", Chap 12.3, MIT press, (2022)

Binary Search Tree Deletion

Deletion implementation

Delete node `u` from binary search tree `bst`

```
def delete(bst, u):
    if not u.right:
        shift_nodes(bst, u, u.left)
    elif not u.left:
        shift_nodes(bst, u, u.right)
    else:                      # u has two children
        u_successor = minimum(u.right)
        if u_successor != u.right:
            shift_nodes(bst, u_successor, u_successor.right)
            u_successor.right = u.right
            u_successor.right.parent = u_successor
        shift_nodes(bst, u, u_successor)
        u_successor.left = u.left
        u_successor.left.parent = u_successor
```



```
def shift_nodes(bst, old, src):
    if not old.parent:
        bst.root = src
    elif old == old.parent.left:
        old.parent.left = src
    else:
        old.parent.right = src
    if src:
        src.parent = old.parent
```

Complexity: $O(h)$ where h is tree height

References:

(Deletion pseudocode - we follow this naming convention) https://en.wikipedia.org/wiki/Binary_search_tree#Deletion

(Deletion) T. Cormen et al., "Introduction to algorithms", Chap 12.3, MIT press, (2022)

Red-Black Trees

What they are

How they are implemented

Red-Black Trees (RBTs)

Self-balancing **binary search trees**

Assuredly fast **search, insertion, deletion**

History

Adelson-Velsky & Landis (AVL tree) (1962)



Bayer (Symmetric binary B-trees) (1972)



Guibas & Sedgewick (RB trees) (1978)



Sedgewick (Left-leaning RB trees) (2008)



Complexity (for n data items)

By balancing, red-black trees guarantee **speed**

Worst case: **search, insert, delete** $\rightarrow O(\log n)$

Storage of red-black trees: $\Theta(n)$

Suits

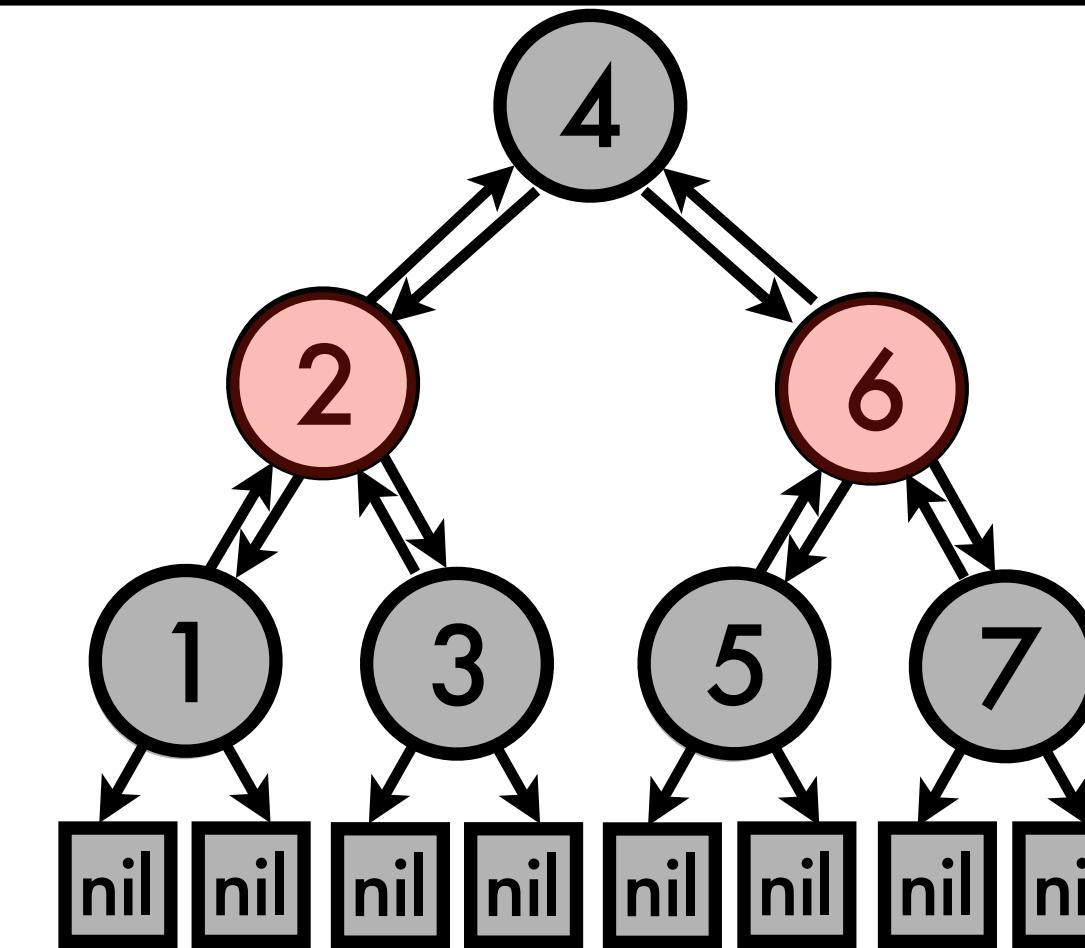
Abstract Data Types

Set

Map

Applications: Container libraries (C++, Java)

Linux CFS



- If node has no child, points to a special `nil` node
- treated as black
 - simplify logic
 - omitted from diagrams

Red-black tree: a **Binary Search Tree** where each node also has a **colour** (which can be **red** or **black**)

Approximately balanced: Tree height is $\Theta(\log n)$

References/Notes/Image credits:

https://en.wikipedia.org/wiki/Red%20black_tree

(AVL Trees) G. M. Adelson-Velsky and E. M. Landis, "An algorithm for the organization of information", Doklady Akademii Nauk (1962)

(G. M. Adelson-Velsky) <https://www.math.toronto.edu/askold/2014-UMN-4-e-Adelson.pdf>

(E. Landis) https://opc.mfo.de/detail?photo_id=2447

R. Bayer, "Symmetric binary B-trees: Data structure and maintenance algorithms", Acta informatica (1972)

(R. Bayer) https://www.computerhope.com/people/rudolf_bayer.htm

L. Guibas and R. Sedgewick, "A dichromatic framework for balanced trees", SFCS (1978)

(L. Guibas) <https://geometry.stanford.edu/member/guibas/>

(R. Sedgewick) <https://sedgewick.io/>

(Red-Black Trees) T. Cormen et al., "Introduction to algorithms", Chap 13.1, MIT press, (2022)

(Linux CFS) https://en.wikipedia.org/wiki/Completely_Fair_Scheduler

Red-Black Tree Properties

Five key properties of Red-Black Trees (CLRS)

RBTs are **Binary Search Trees** with:

Property 1 Every node is **red** or **black**

Property 2 The root node is **black**

Property 3 Every leaf node (nil node) is **black**

Property 4 If a node is **red**, both of its children are **black**

Property 5 Starting from any node, all simple paths down to leaf nodes hold the same number of **black** nodes

Tree height (h) is $\Theta(\log n)$

By **Property 5** all root-leaf paths have the same number of **black** nodes

By **Property 4** **red** nodes have **black** children

\Rightarrow longest path, $h_{\max} \leq 2 \cdot$ shortest path, h_{\min}

Note: **perfect binary tree** has $n = 2^{h+1} - 1$ nodes

-

$$\begin{array}{l} h = 0 \\ n = 1 \end{array}$$

$$\begin{array}{l} h = 1 \\ n = 3 \end{array}$$

$$\begin{array}{l} h = 2 \\ n = 7 \end{array}$$

$$\Rightarrow 2^{h_{\min}+1} - 1 \leq n \leq 2^{h_{\max}+1} - 1$$

$$\Rightarrow h_{\min} \leq \log(n + 1) - 1 \leq h_{\max} \leq 2 \cdot h_{\min}$$

$$\Rightarrow h_{\min} = \Theta(\log n)$$

Key "balanced tree" result

$$h_{\min} \leq h_{\max} \leq 2 \cdot h_{\min} \Rightarrow h_{\max} = \Theta(\log n)$$

\Rightarrow Min Max Search Successor Predecessor $O(\log n)$

References:

(CLRS) T. Cormen et al., "Introduction to algorithms", Chap 13.1, MIT press, (2022)

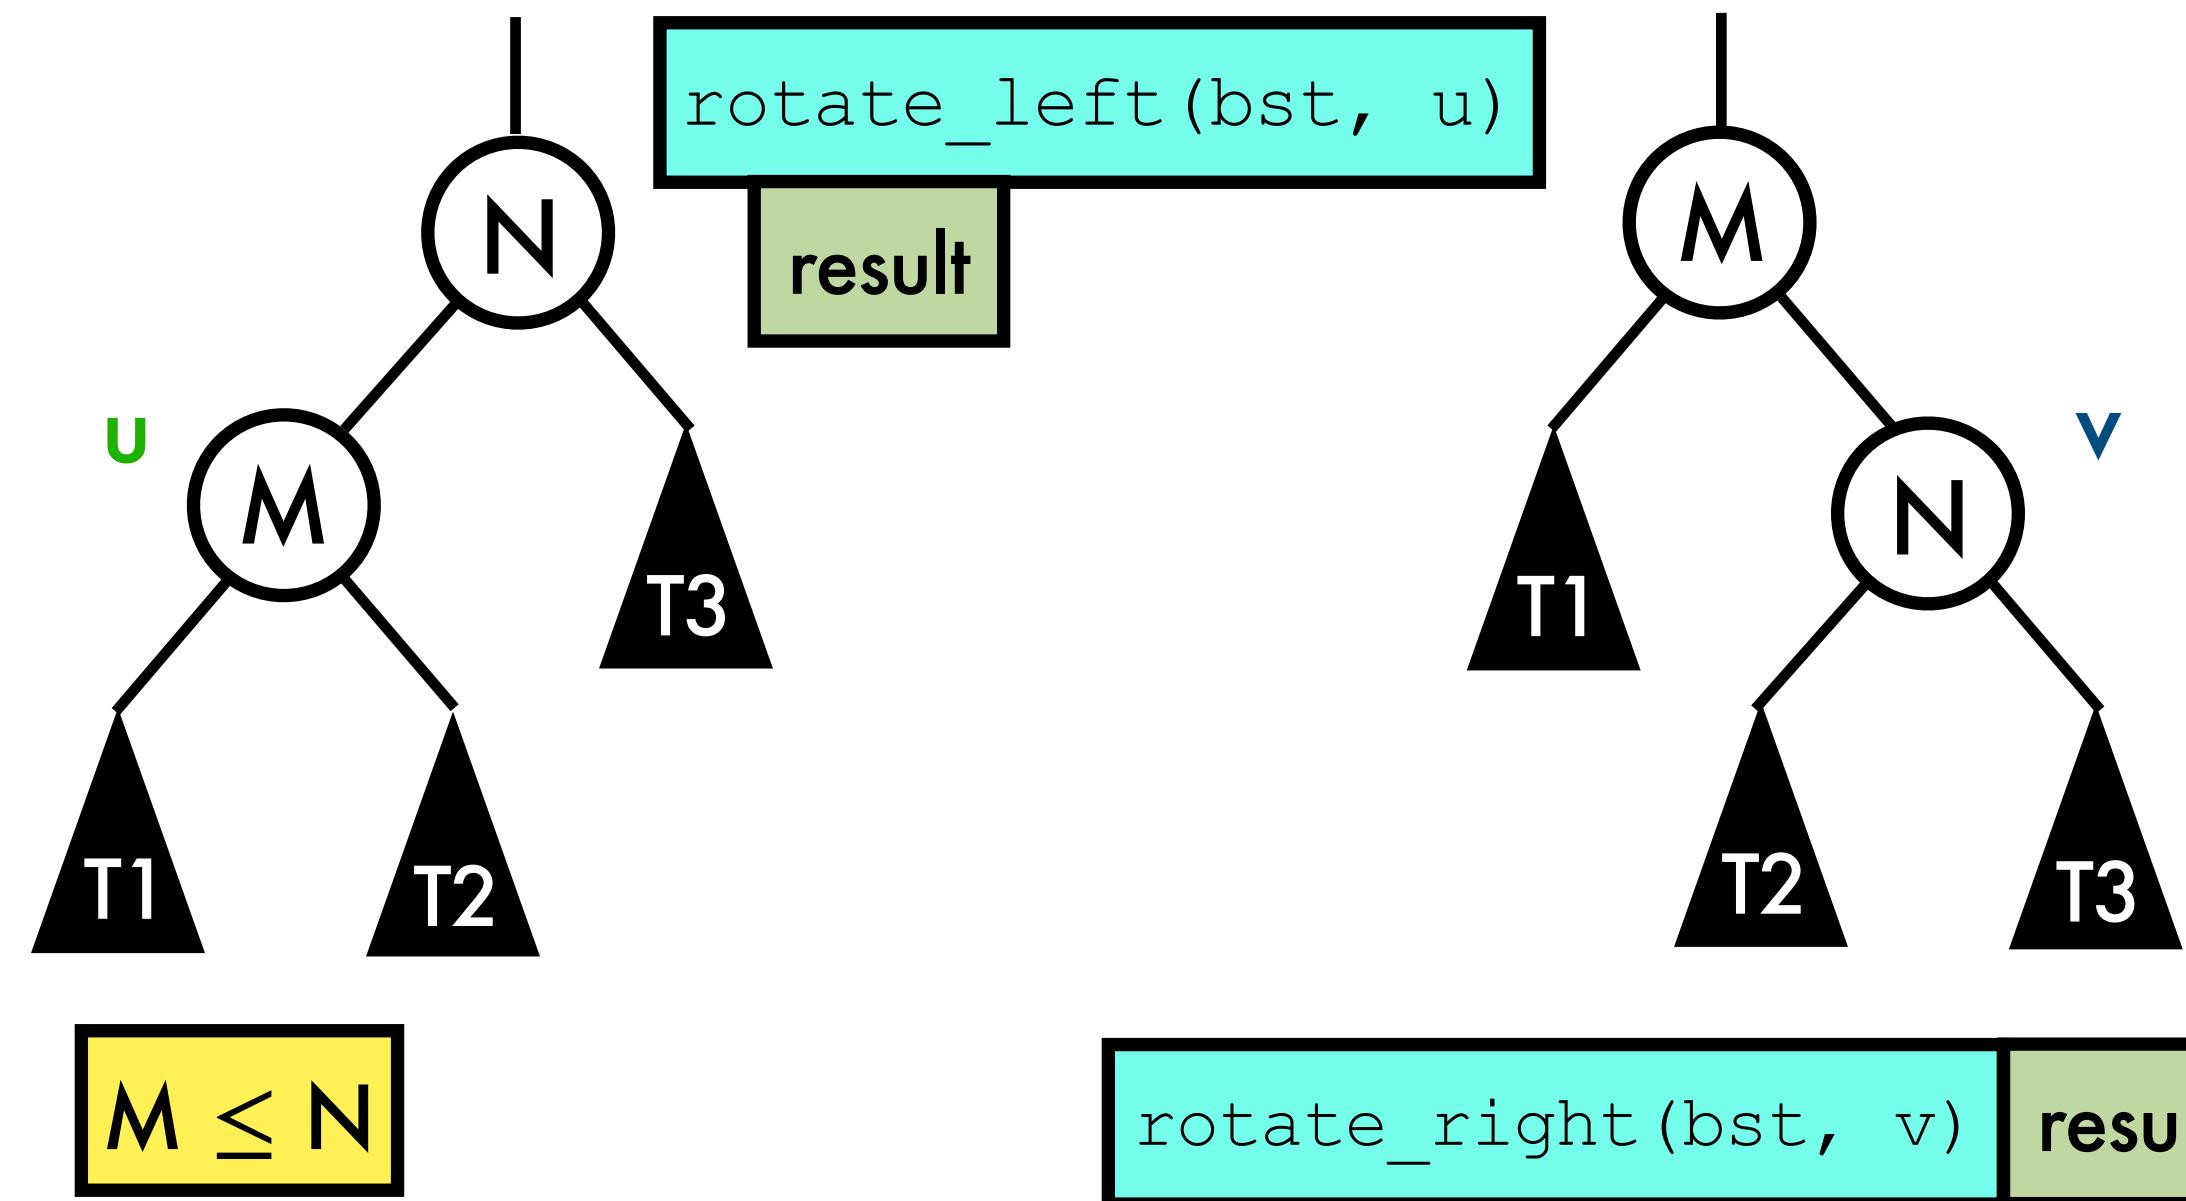
(Proof of tree height) L. Arge and M. Lagoudakis, CPS 230 lecture notes, https://courses.cs.duke.edu/cps130/fall02/fall02lectures/lecture18/long_redblack.pdf (2002)

Rotation Operations

The rotation operation

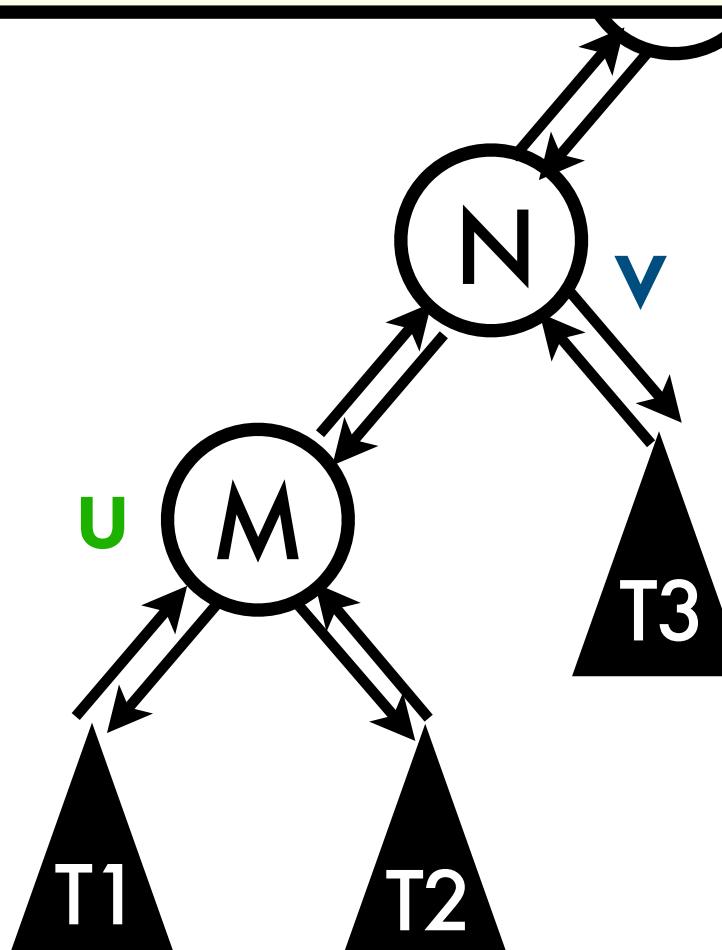
Rotations restructure a tree locally without breaking the Binary Search Tree Property

Rotations change links (not colours) at $O(1)$ cost



Rotate_left implementation in Python

```
def rotate_left(self, u): # self is BST instance
    v = u.right
    u.right = v.left
    if v.left != self.nil:
        v.left.parent = u
    v.parent = u.parent
    if u.parent is None:
        self.root = v
    elif u.parent.left == u:
        u.parent.left = v
    else:
        u.parent.right = v
    v.left, u.parent = u, v
```



References:

(Rotations) L. Arge and M. Lagoudakis, CPS 230 lecture notes, https://courses.cs.duke.edu/cps130/fall02/fall02lectures/lecture18/long_redblack.pdf (2002)

(rotation logic/comments based on CLRS) T. Cormen et al., "Introduction to algorithms", Chap 13.2, MIT press (2022)

Python code snippet reference for `rotate_left` - <https://blog.boot.dev/python/red-black-tree-python/>

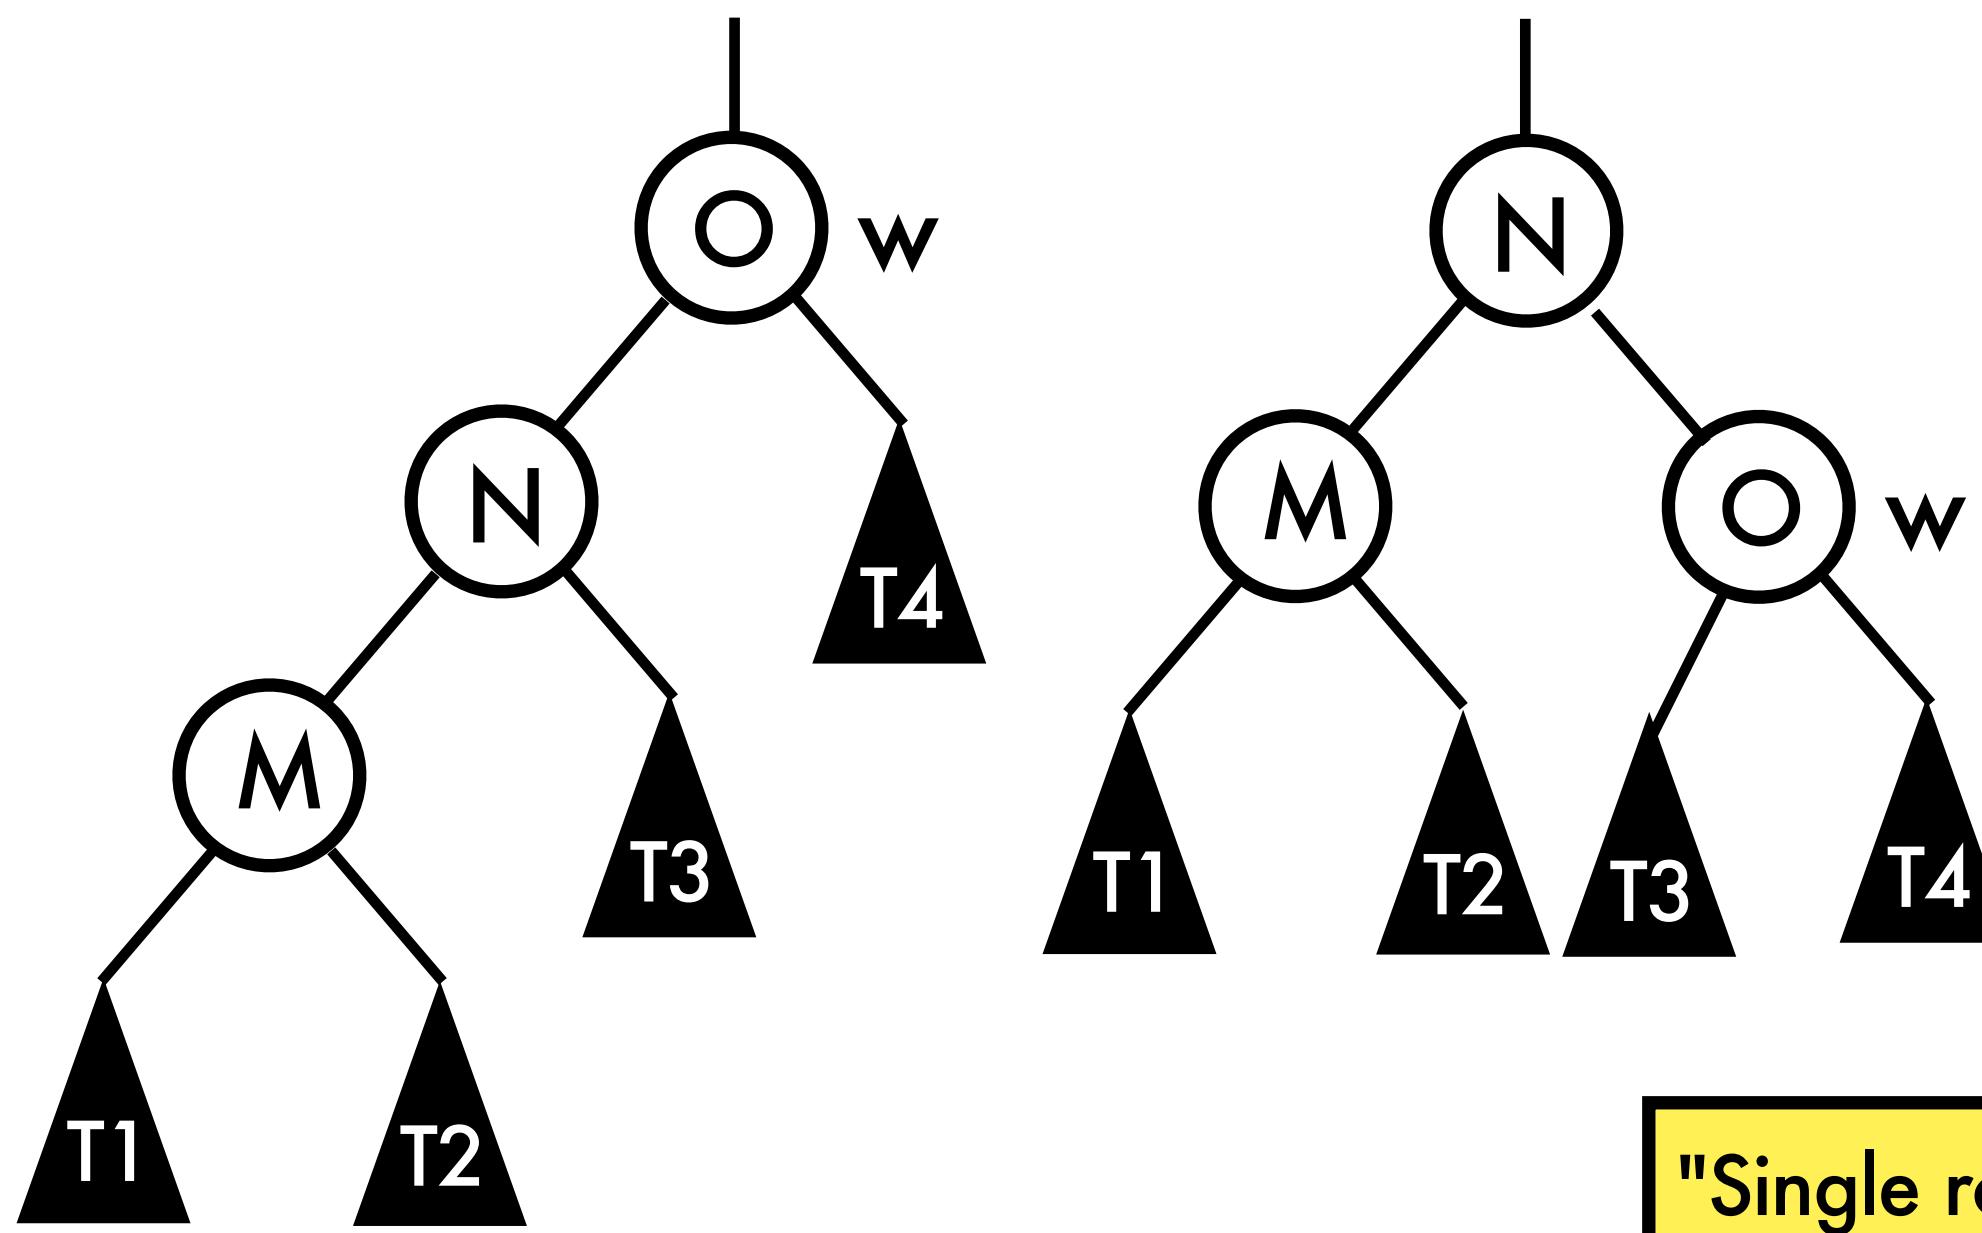
Why Are Rotations Useful?

The benefits of rotation

When too many nodes are on a single path,
rotation distributes them to neighbour paths

This restores **balance** in the tree

It preserves the **Binary Search Tree Property**

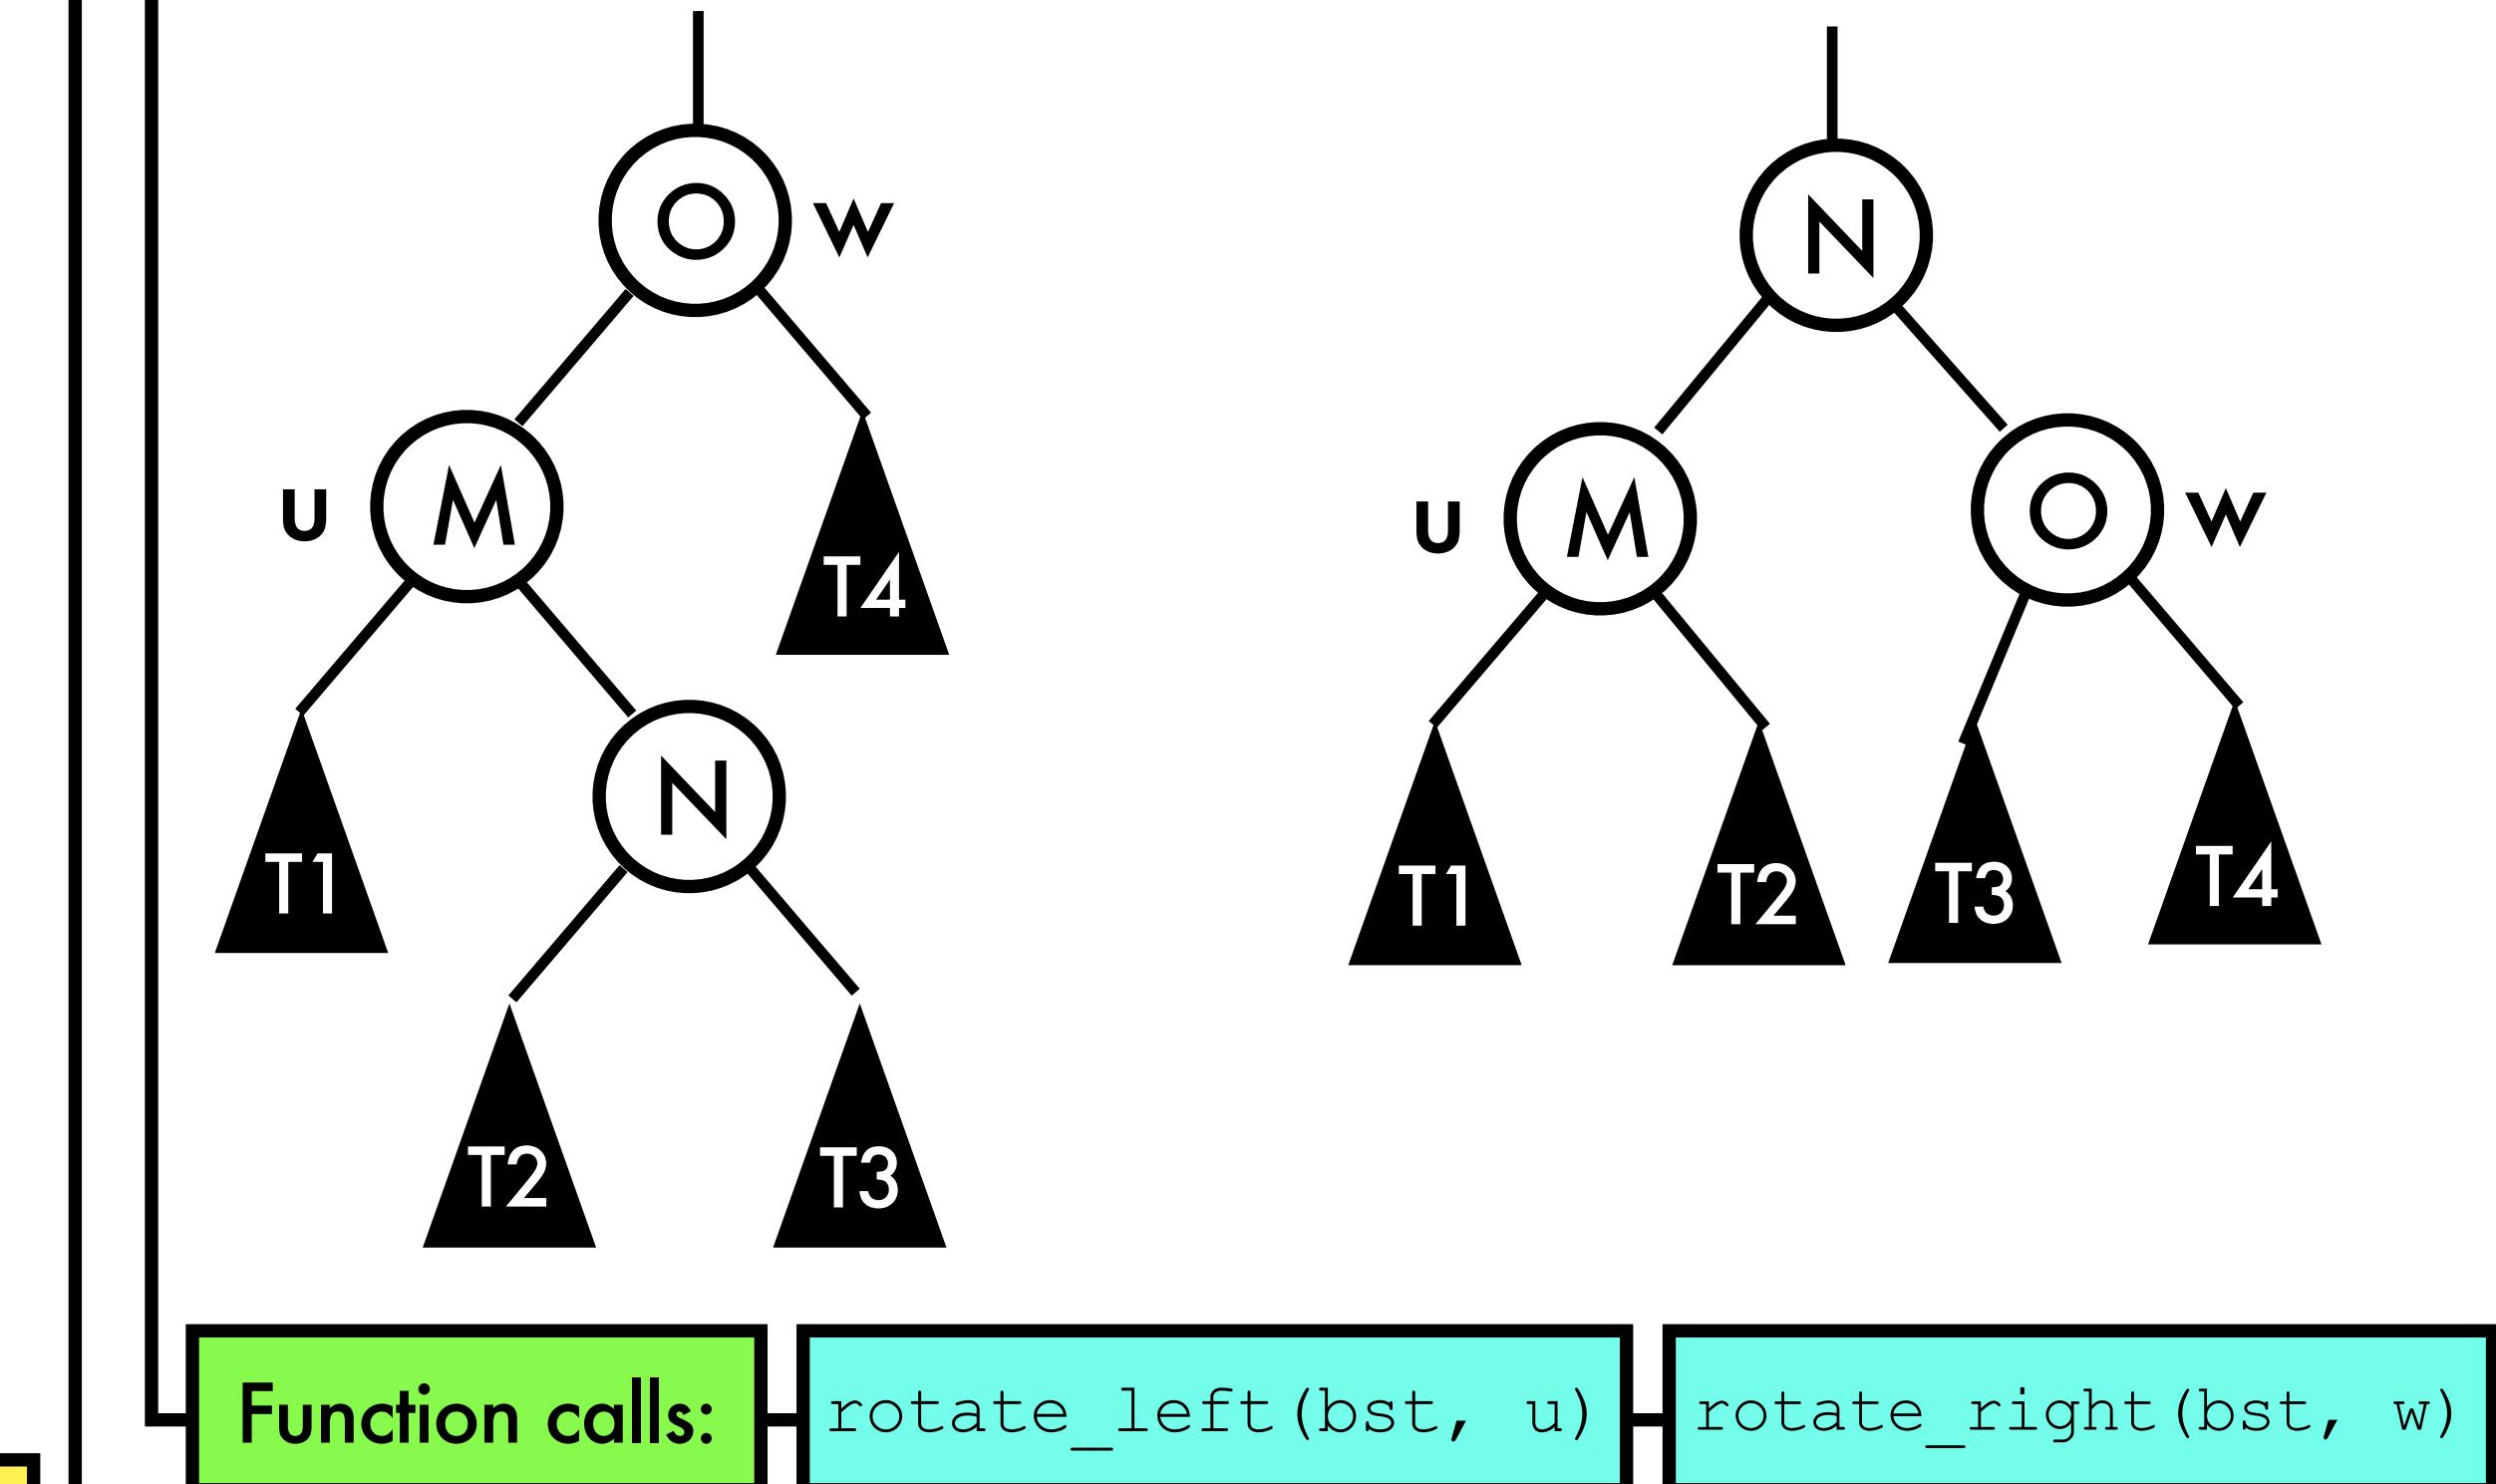


`rotate_right(bst, w)`

Produces more balanced structure

Double rotations

If we have a **chain** with both a **left child** and a **right child**, then a "**double rotation**" is needed



References:

(Rotations) L. Arge and M. Lagoudakis, CPS 230 lecture notes, https://courses.cs.duke.edu/cps130/fall02/fall02lectures/lecture18/long_redblack.pdf (2002)

Red-Black Tree Insertion

Insertion overview

Initially, RBT insertion is the same as **BST** insertion

What **colour** should the **inserted node** be?

We have two options: **red** or **black**

If we colour the **inserted node red** and its parent

is **red**, we violate **Property 4**

If we colour the **inserted node black**, the extra
black node will violate **Property 5**

Strategy: **fix problem locally** with **rotations** and
recolouring

"Escalate" further issues **up the tree**

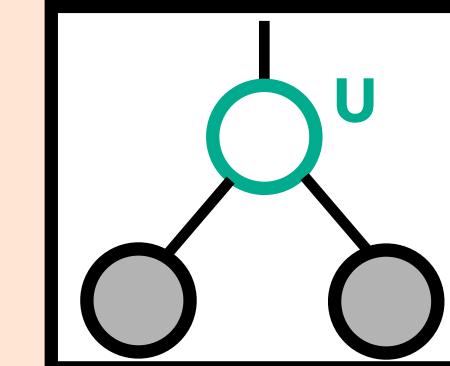
Tree height is $\Theta(\log n)$ so fixing is $O(\log n)$

Total Red-Black Tree insertion $O(\log n)$

Insertion - handling cases

Red-black tree **insertion** involves patiently
handling **several cases**

Precursor: whenever we have a **RBT violation**,
the node has **black** children
True after inserts (**nuls are black**)
Also true **further up the tree**



Case 1 **u** has a black parent

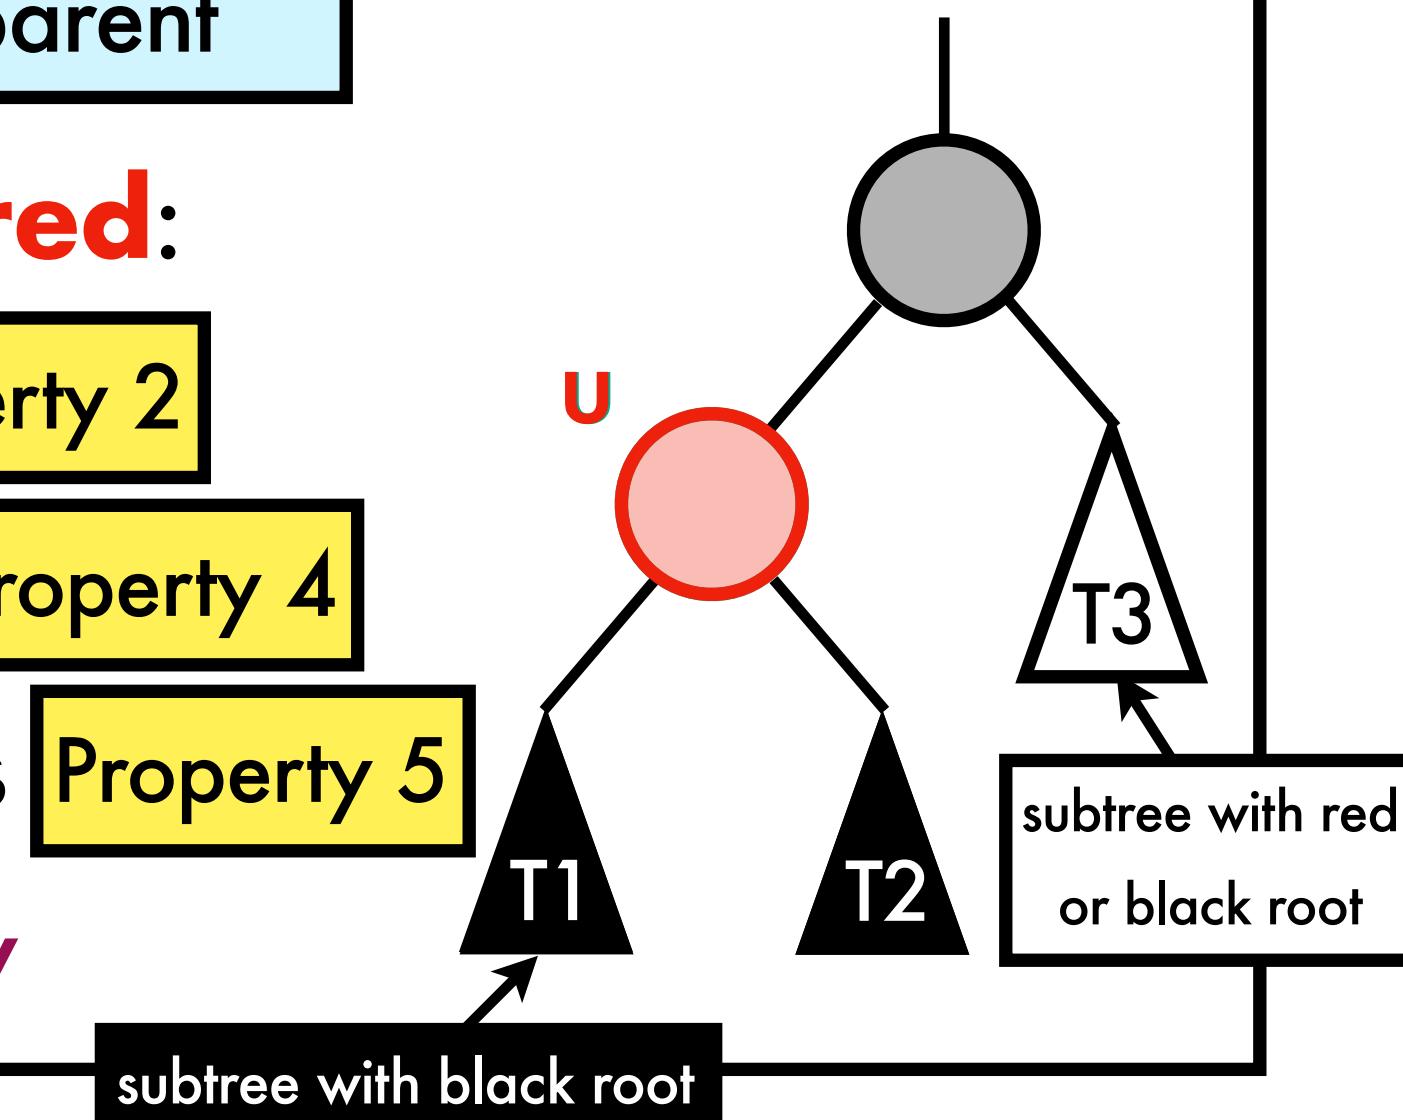
Simply colour node **u red**:

u is not **the root** **Property 2**

Only **black** children **Property 4**

No extra **black** nodes **Property 5**

Problem is fixed **locally**



Reference:

(Rotations) L. Arge and M. Lagoudakis, CPS 230 lecture notes, https://courses.cs.duke.edu/cps130/fall02/fall02lectures/lecture18/long_redblack.pdf (2002)

Red-Black Tree Insertion Cont.

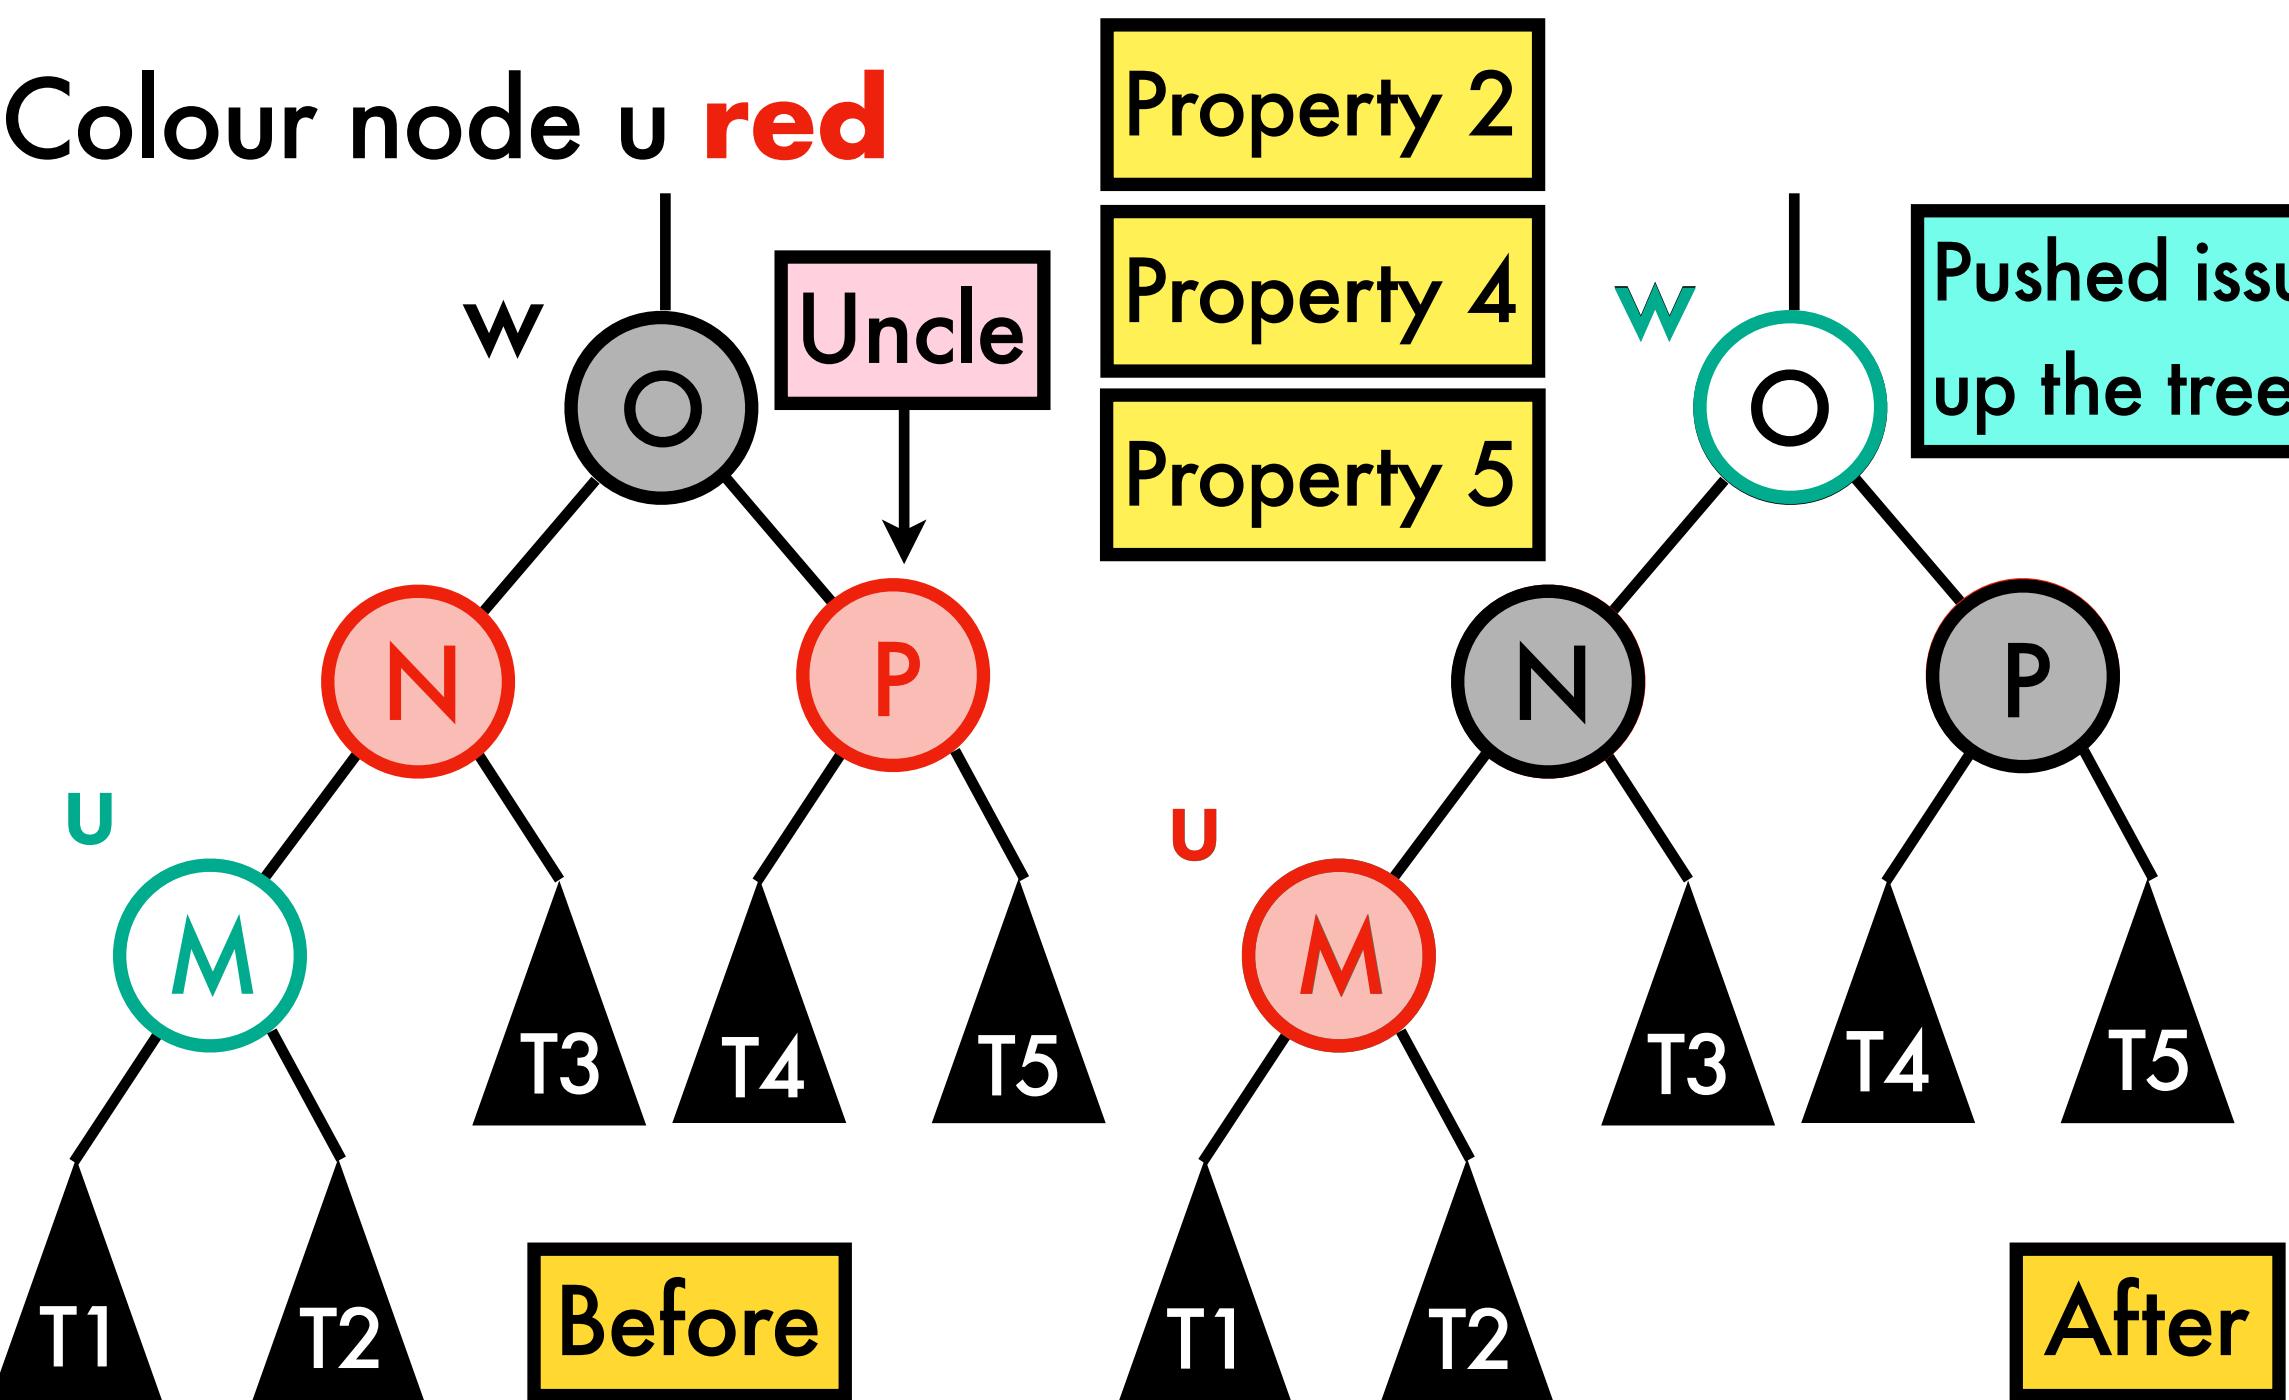
Parent of u is red

When parent of **u** is red, the steps depend on colour of its parent's sibling (its "uncle")

Case 2 **u** has red parent, red uncle

Push black colour of w down to each child

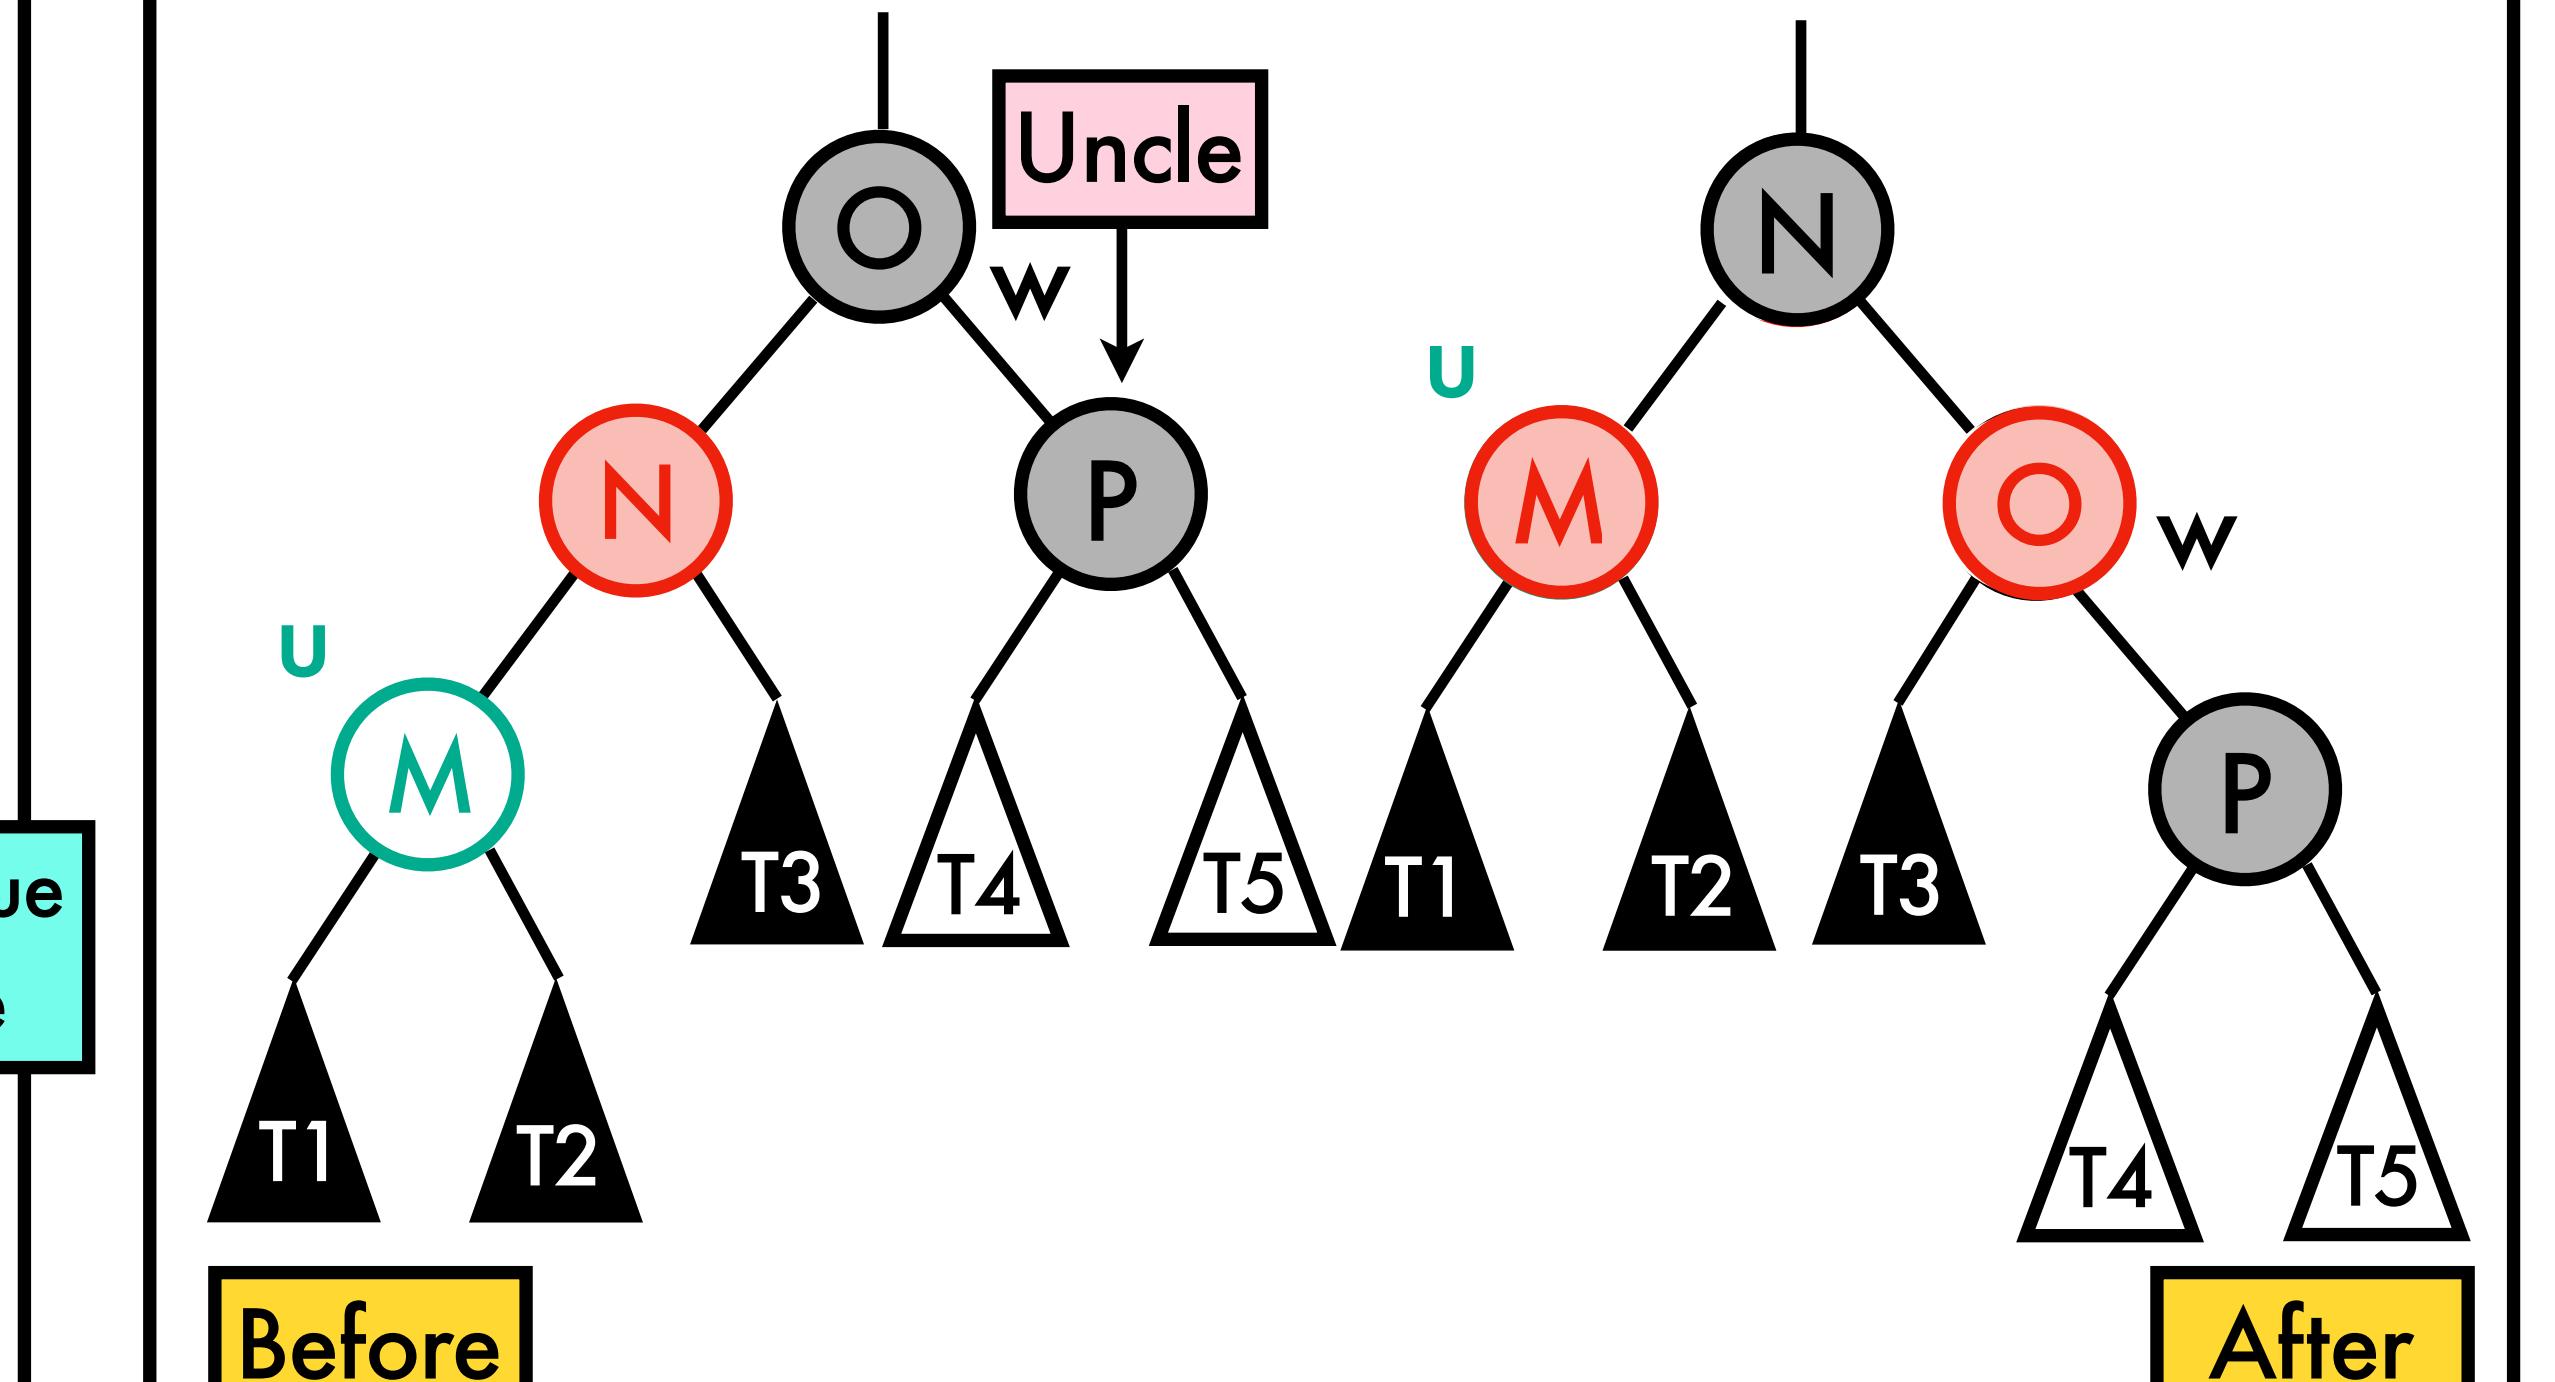
Colour node u red



Parent of u is red

Case 3 **u** has red parent, black uncle

Case 3a **u** is away from its uncle



rotate_right(bst, w)

u → red

u.parent → black

w → red

Problem is fully fixed (no further escalation)

Reference:

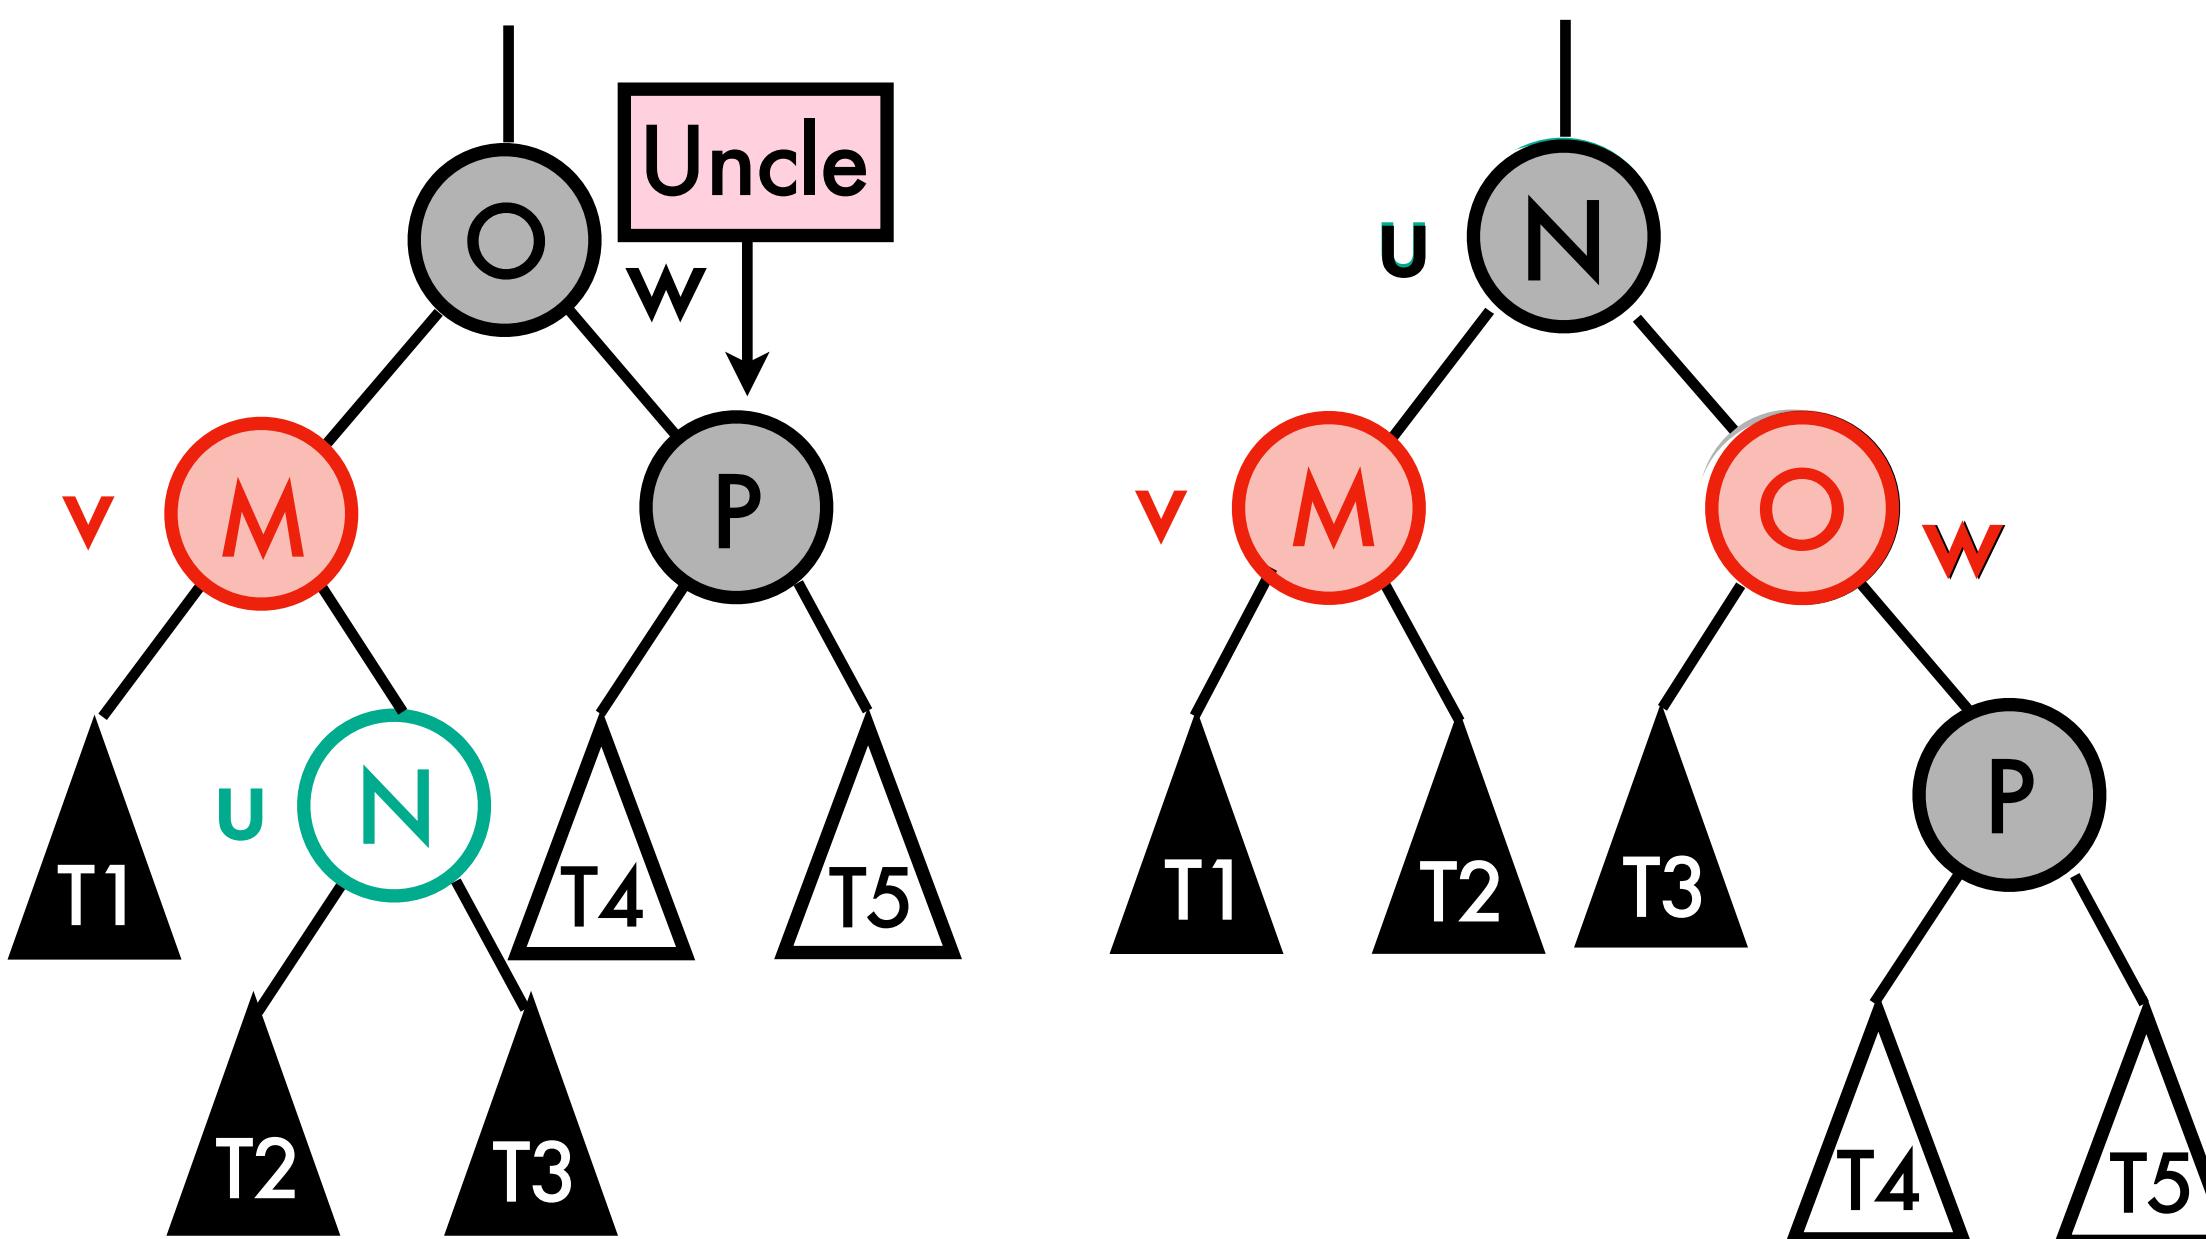
(Rotations) L. Arge and M. Lagoudakis, CPS 230 lecture notes, https://courses.cs.duke.edu/cps130/fall02/fall02lectures/lecture18/long_redblack.pdf (2002)

Red-Black Tree Insertion Cont.

Parent of u is **red**

Case 3 **u** has **red** parent, **black** uncle

Case 3b **u** is towards its uncle



Before

`rotate_left(bst, v)`

`rotate_right(bst, w)`

u → black

w → red

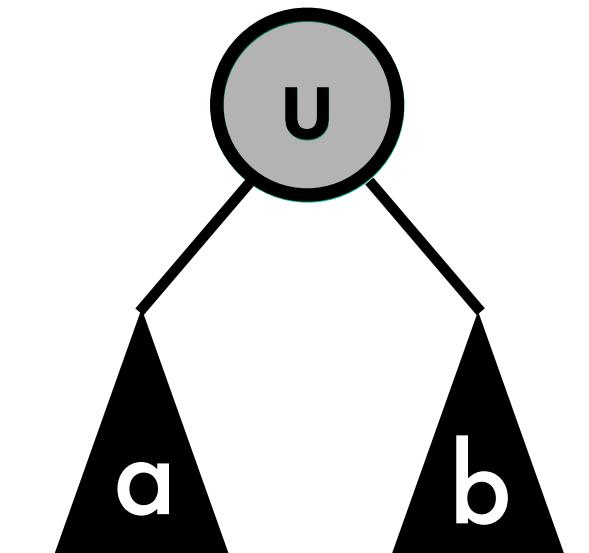
u has no parent

Case 4 **u** has no parent

Colour node **u** **black**

Property 2

Property 5



Problem is fully **fixed** (no further escalation)

Problem is fully **fixed** (no further escalation)

Red-Black Tree Deletion



Widely feared

Deletion overview

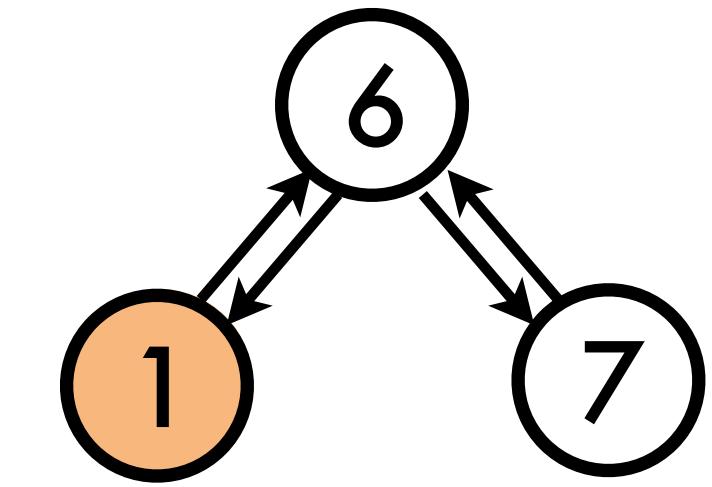
We build on top of **Binary Search Tree deletion**

Red-Black Tree deletion is $O(\log n)$

It is also quite **complicated** to implement

Shifting subtrees

```
def shift_nodes(self, old, new):
    if not old.parent:
        self.root = new
    elif old == old.parent.left:
        old.parent.left = new
    else:
        old.parent.right = new
    new.parent = old.parent
```



Reference:

T. Cormen et al., "Introduction to algorithms", Chap 13.4, MIT press (2022)

Red-Black Tree Deletion

```

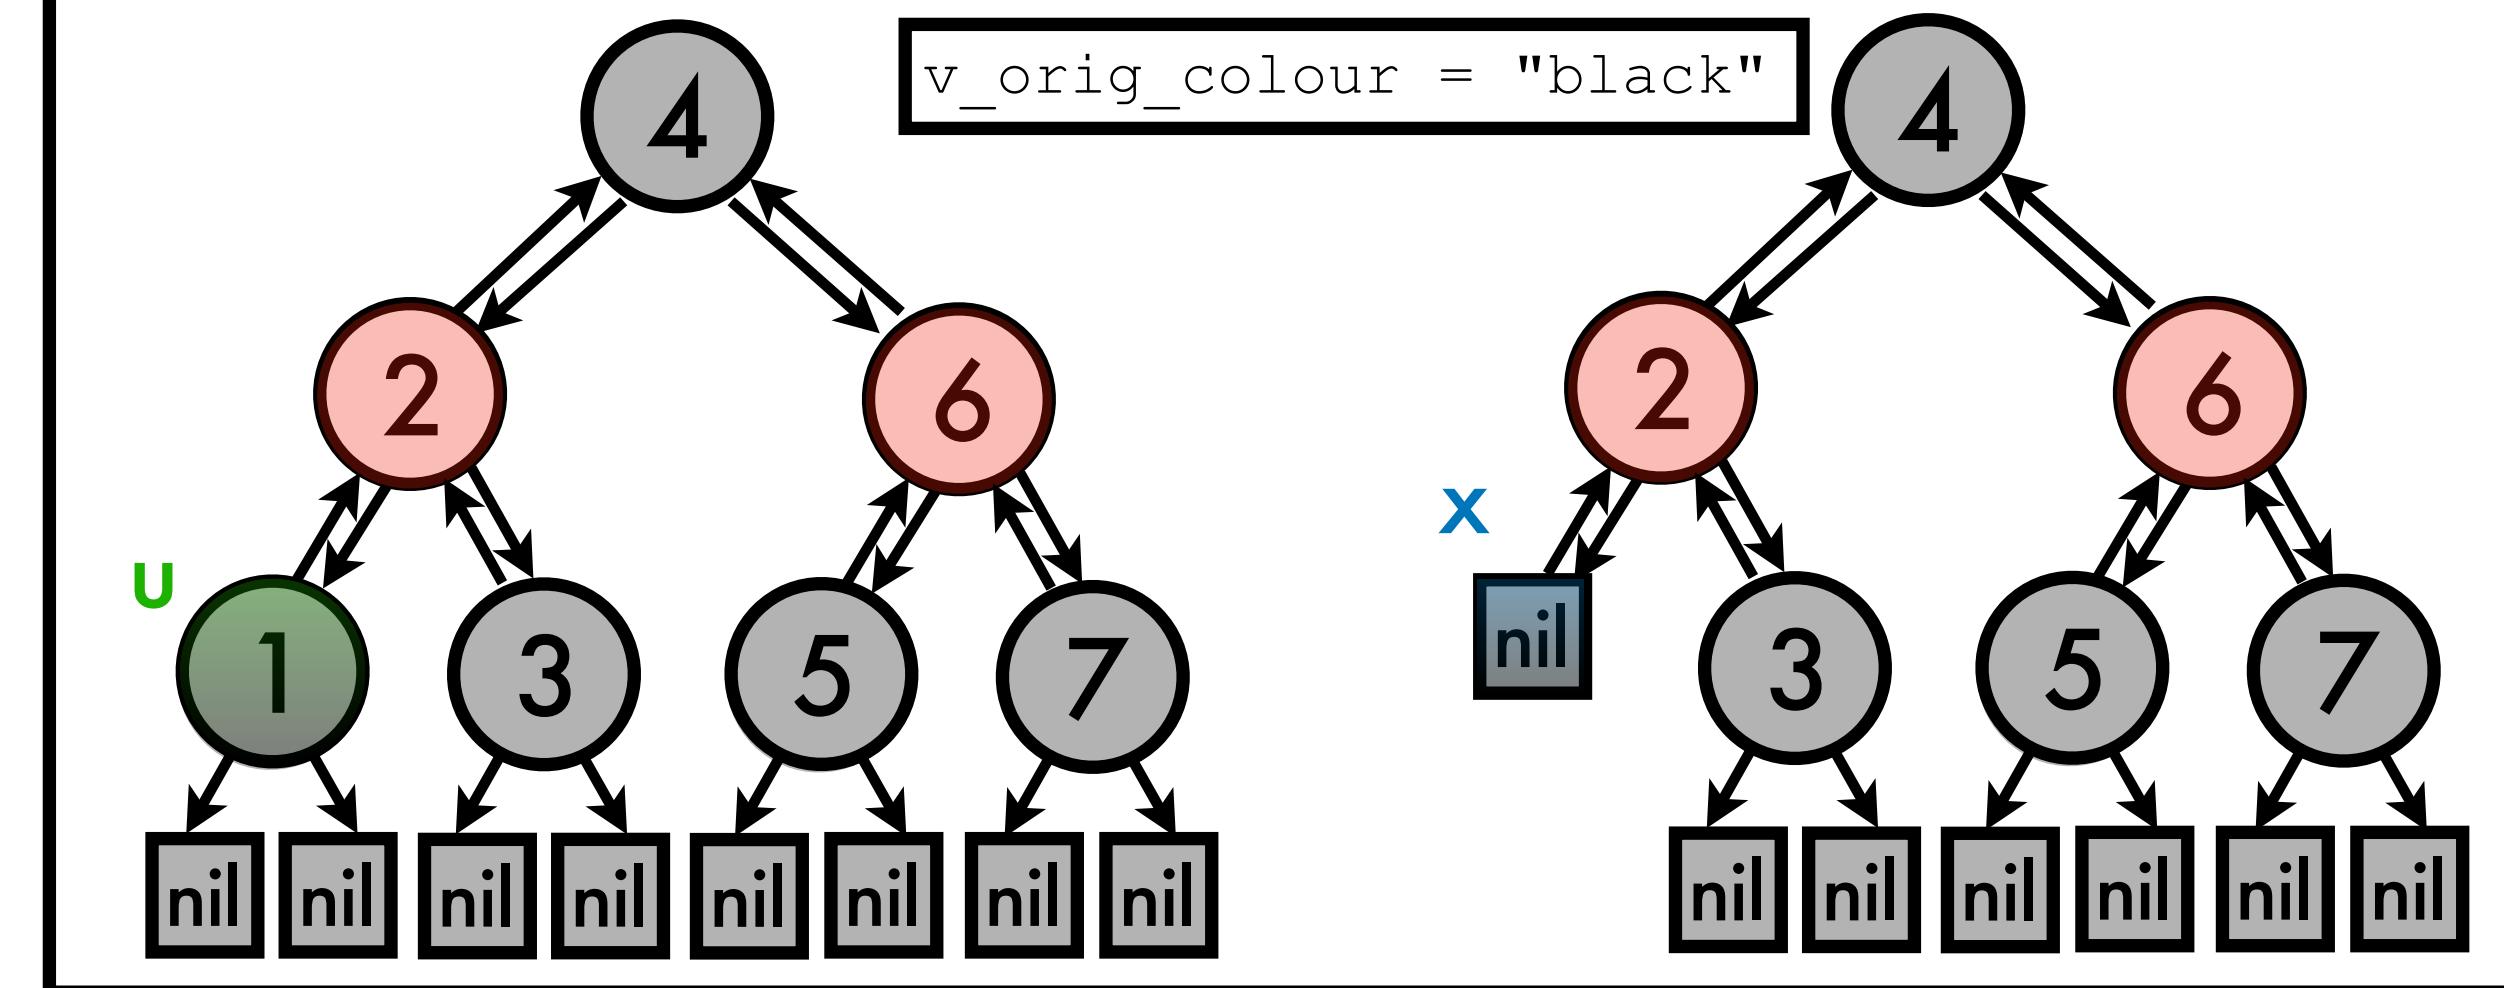
def delete(self, u): # self is an instance of a red-black tree
    v = u # assign v to the node to be deleted
    v_orig_colour = v.colour # track v's original colour
    if u.left == self.nil: # u's left child is nil
        x = u.left
        shift_nodes(self, u, x) # shift up x into u's place
    elif u.right == self.nil: # u's right child is nil
        x = u.right
        shift_nodes(self, u, x) # shift up x into u's place
    else: # u has two children
        v = minimum(u.right)
        v_orig_colour = v.colour
        x = v.right
        if v != u.right:
            shift_nodes(self, v, v.right)
            v.right = u.right
            v.right.parent = v
        else:
            x.parent = v
            shift_nodes(self, u, v)
            v.left = u.left
            v.left.parent = v
            v.colour = u.colour
    if v_orig_colour == "black":
        fix_delete_violations(self, x)

```

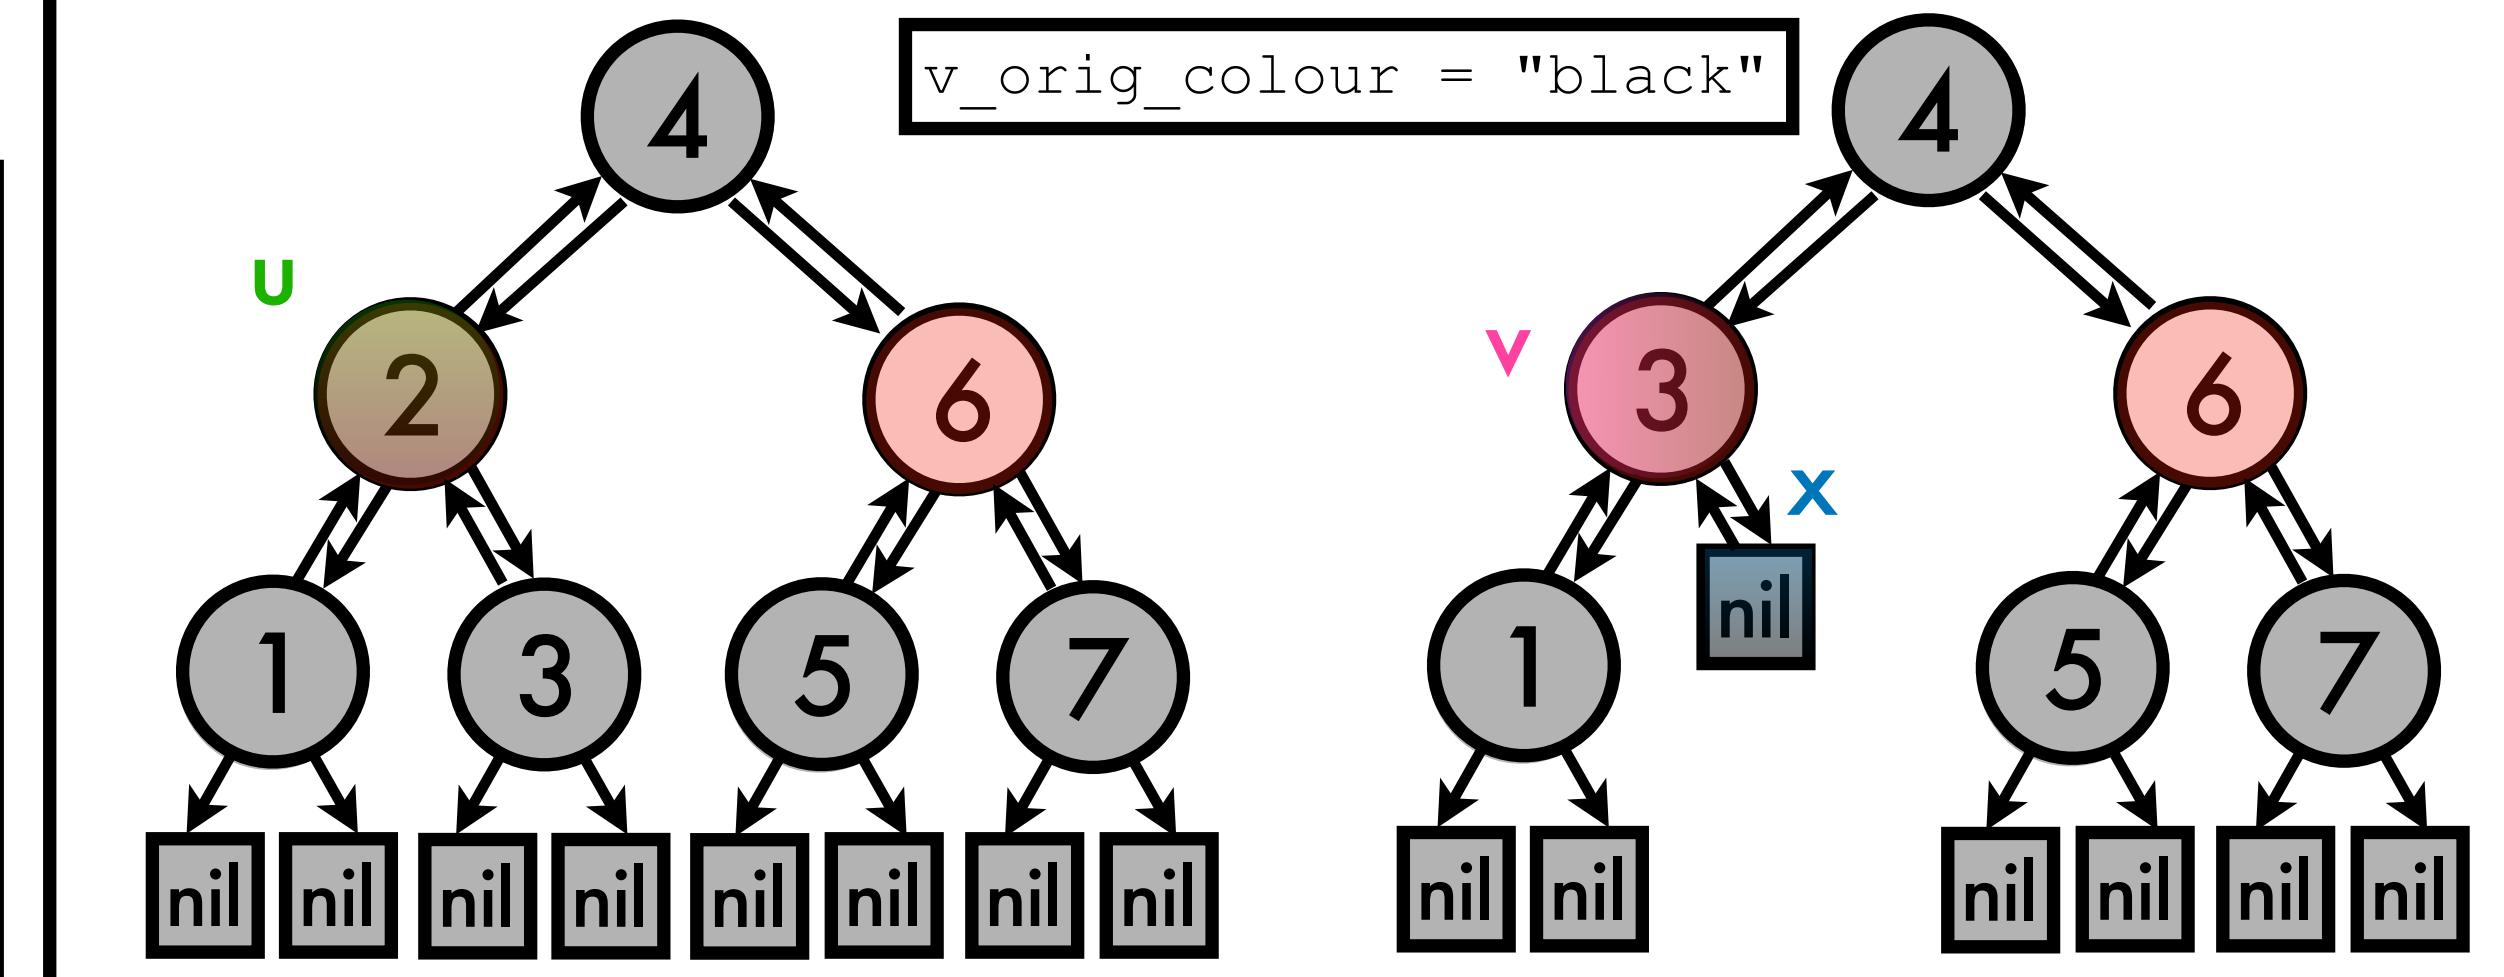
If $v_{\text{orig_color}}$ was red:

- v wasn't the root **Property 2**
- no red nodes have become neighbours **Property 4**
- num black nodes on paths haven't changed **Property 5**

u 's left child is nil (right case is symmetric)



u has two children



Reference:

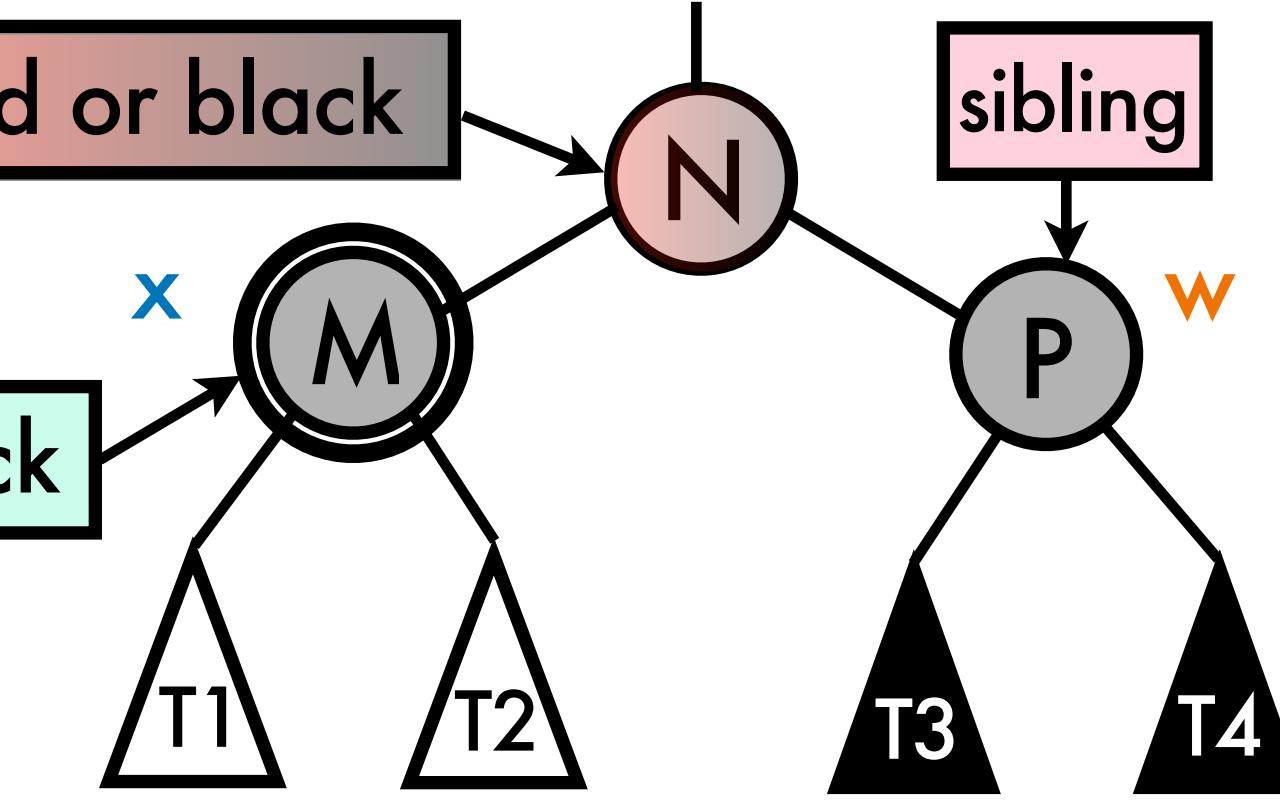
Code adapted from T. Cormen et al., "Introduction to algorithms", Chap 13.4, MIT press (2022)

Fixing RBT Violations

w (x's sibling) is black, w's children are black

Case 2

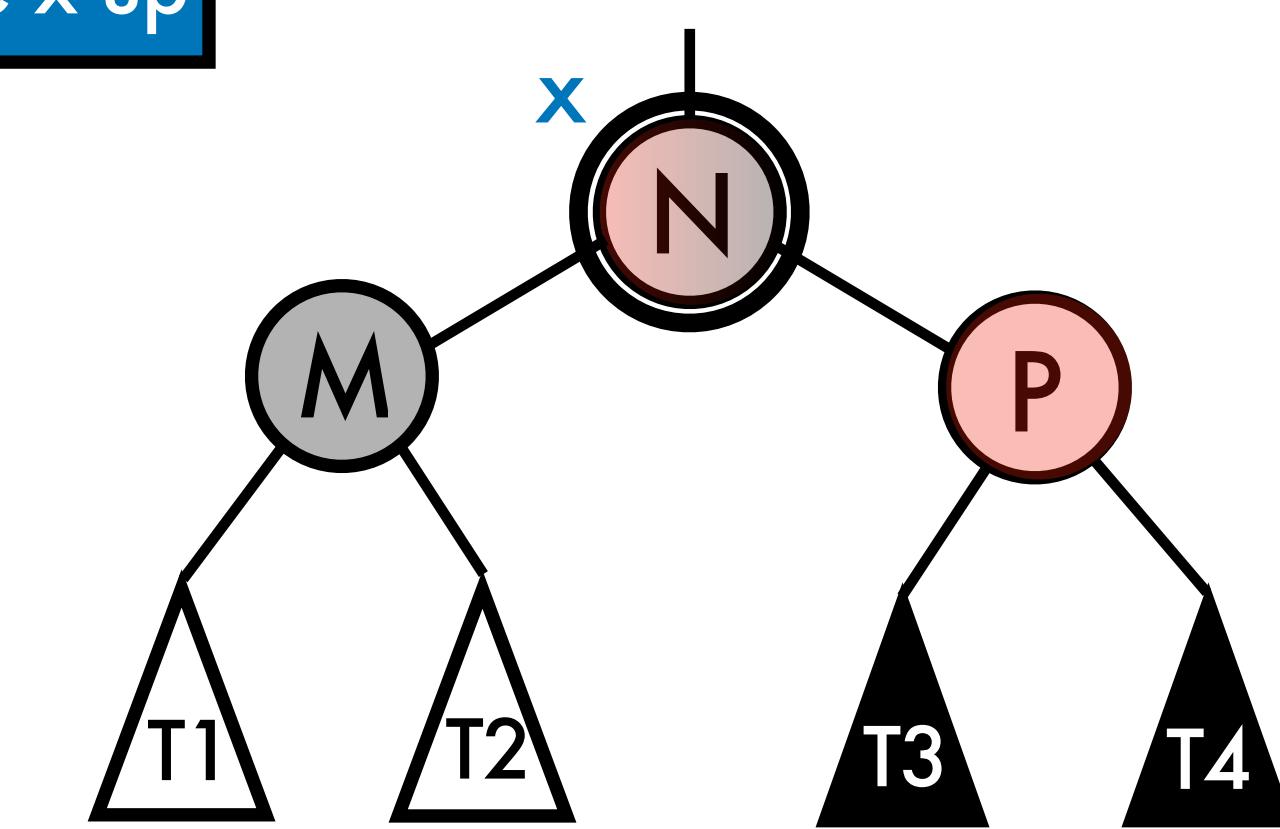
red or black



double black

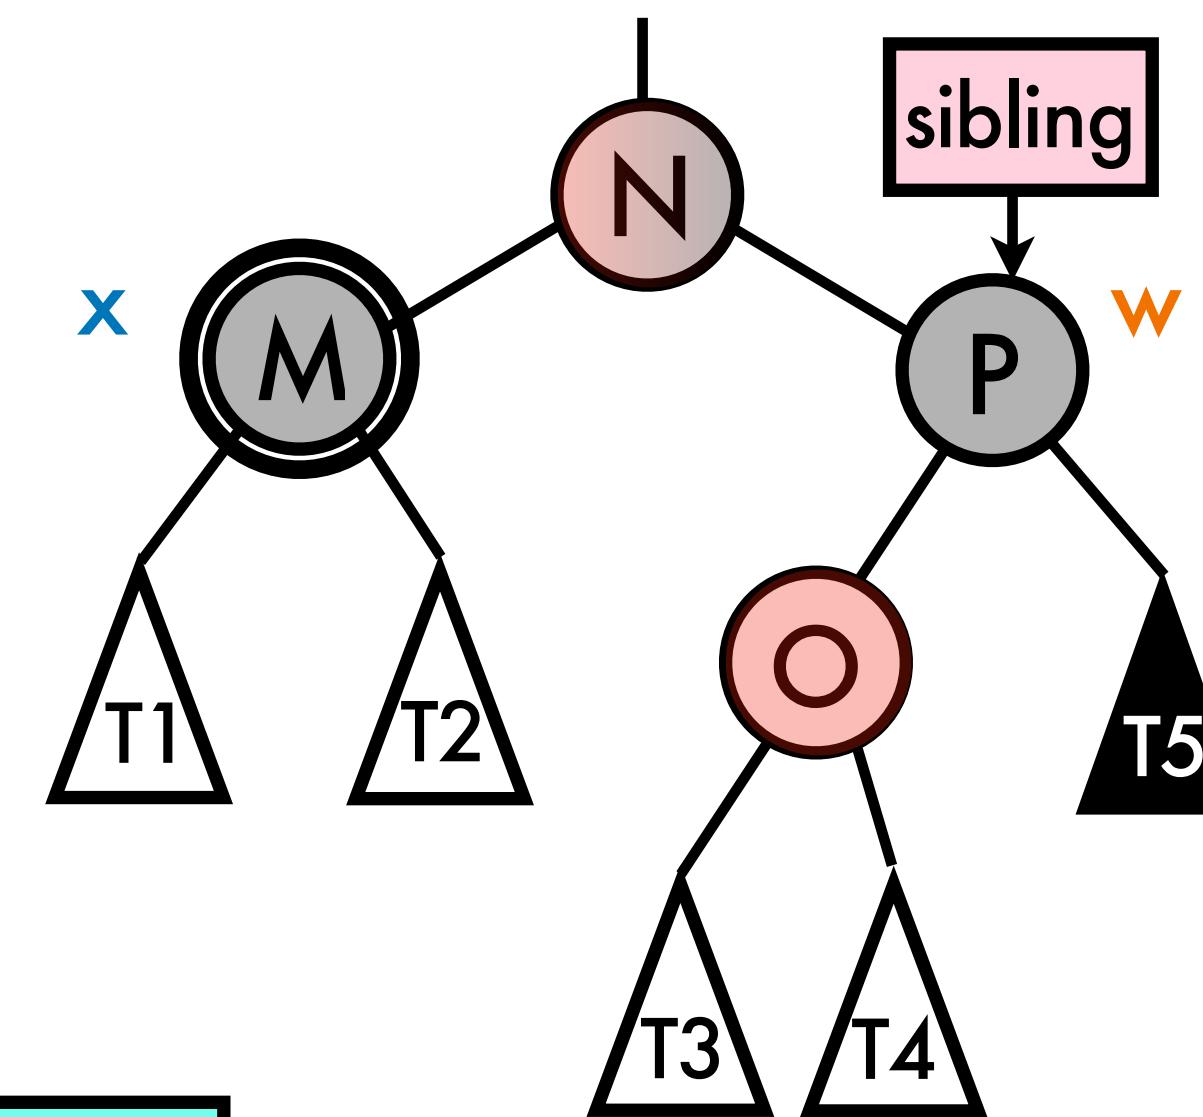
w → red

Move x up



w is black with red left and black right child

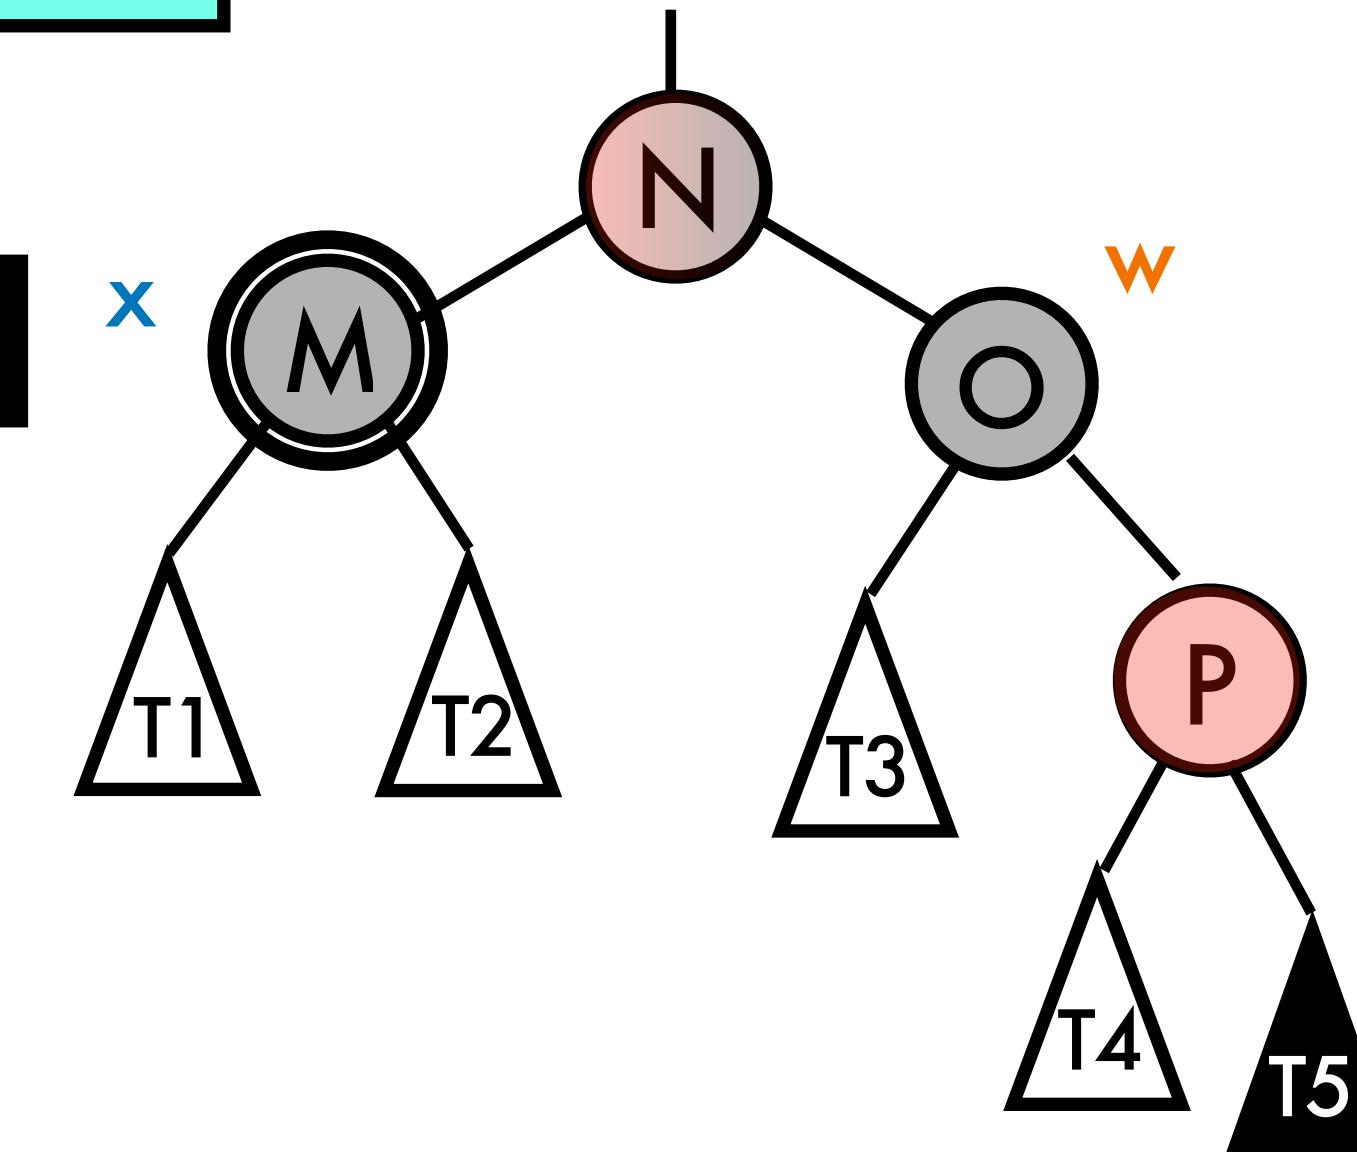
Case 3



rotate_right(bst, w)

w → red

w.parent → black

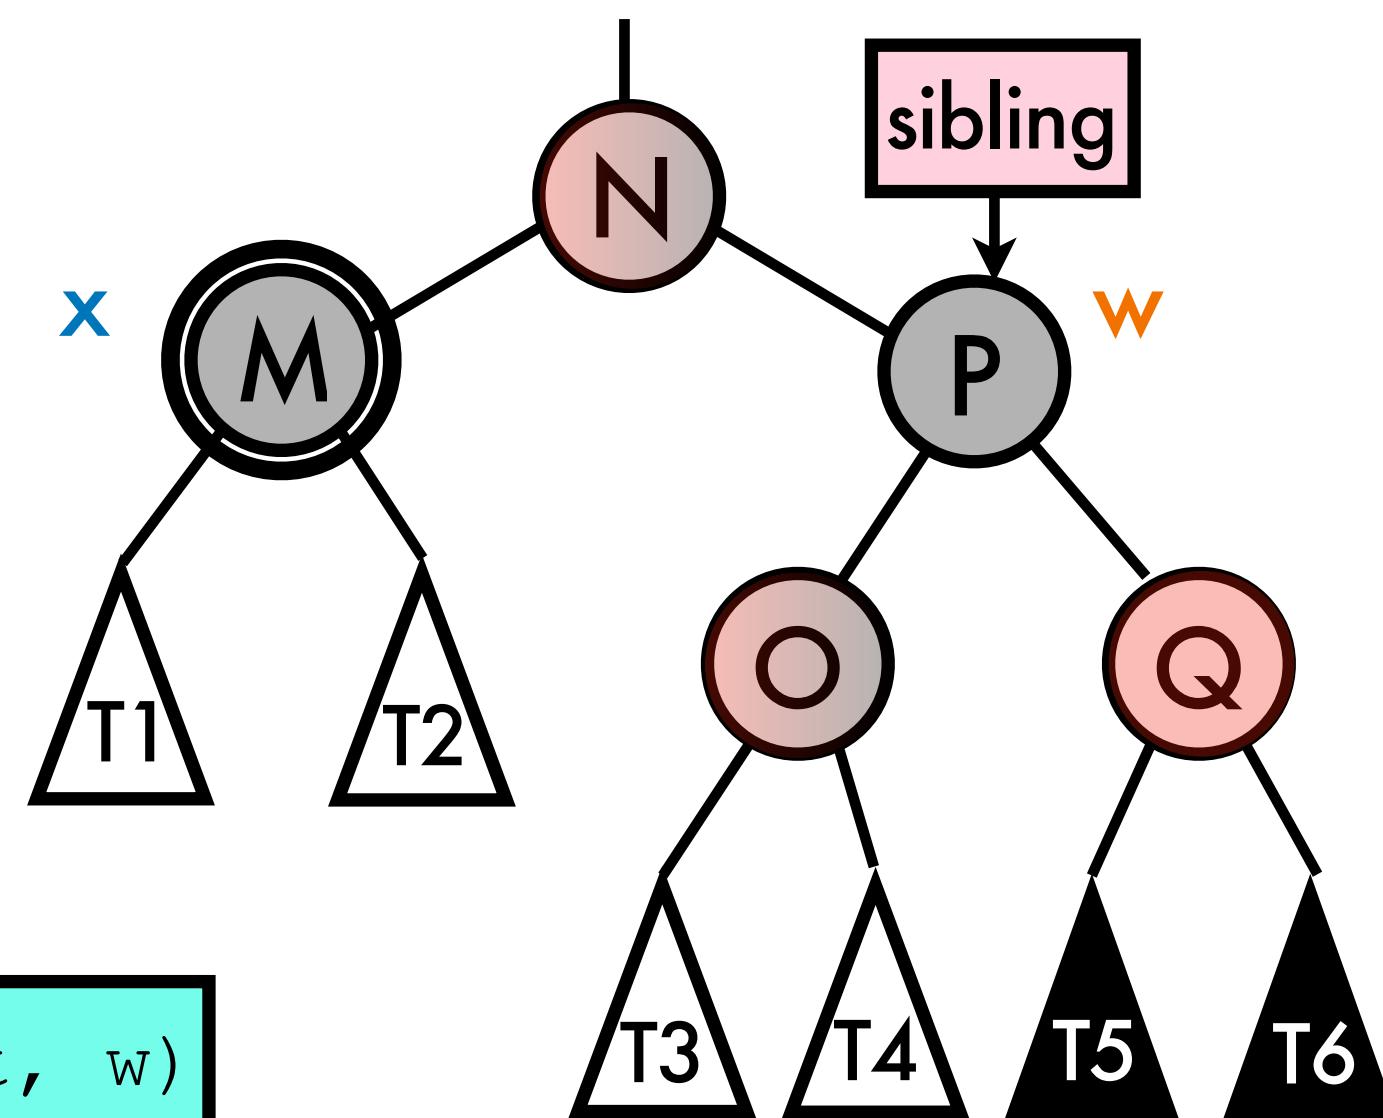


Reference:

T. Cormen et al., "Introduction to algorithms", Chap 13.4, MIT press (2022)

w is black and w's right child is red

Case 4



rotate_left(bst, w)

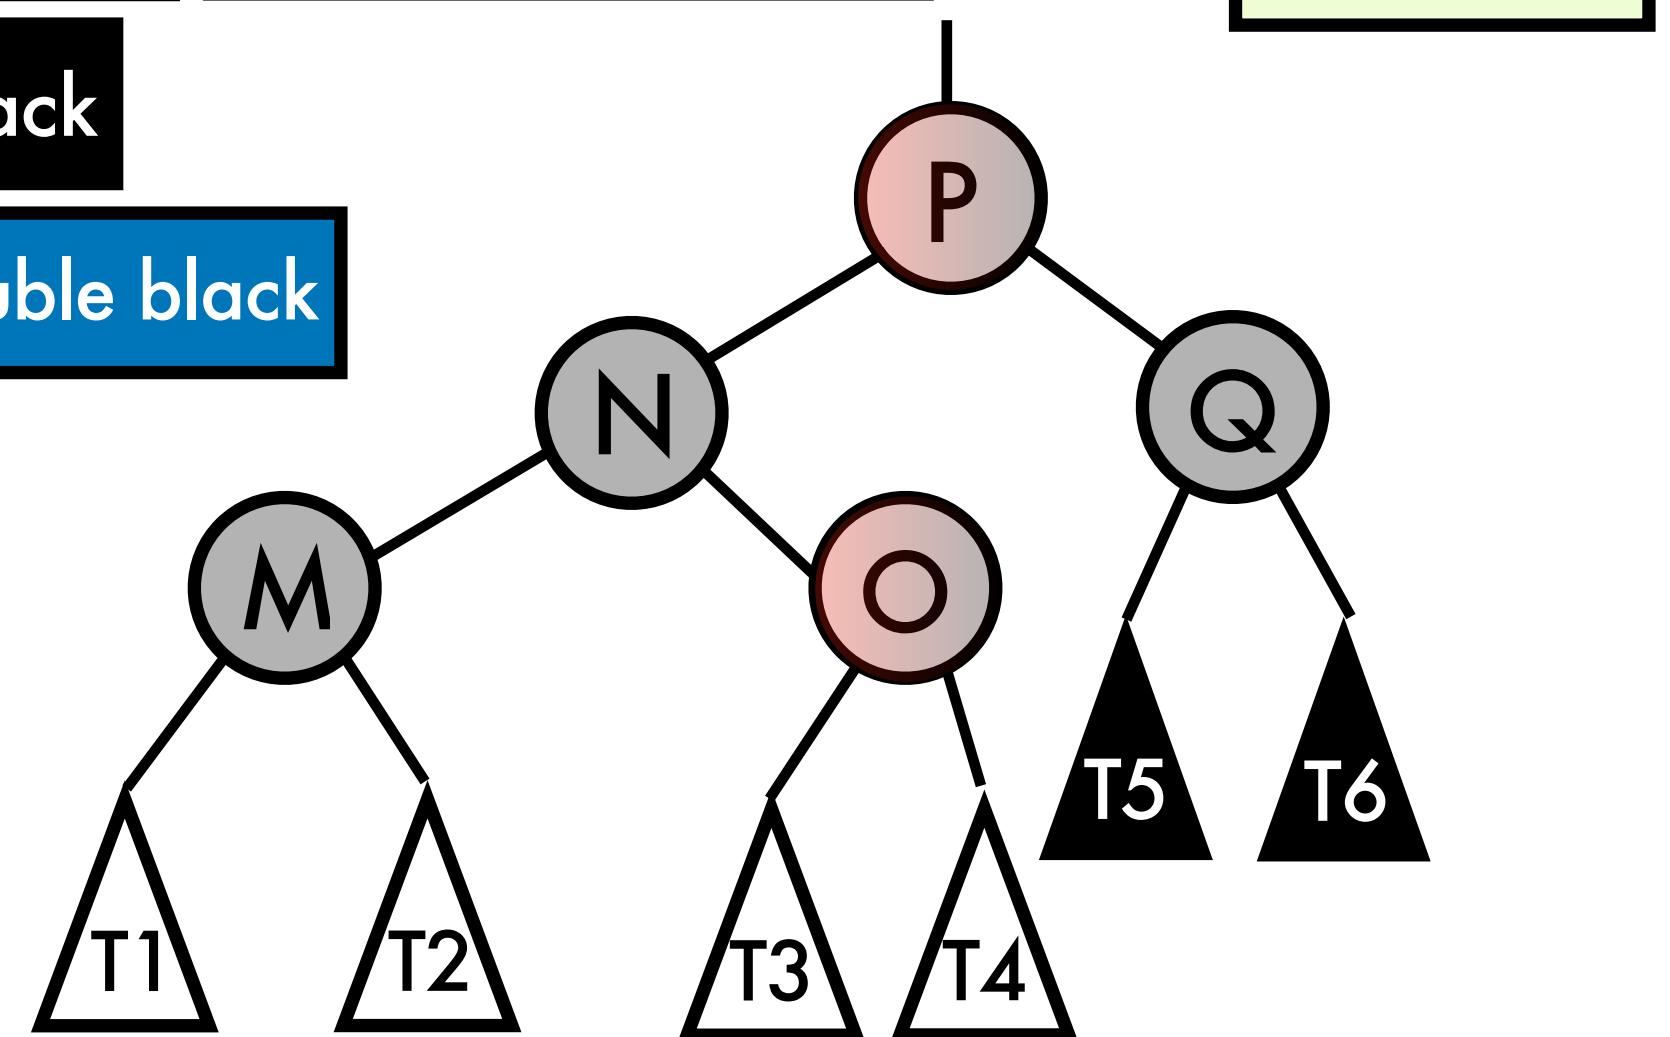
w.right → black

w → w.left's colour

x = root

w.left → black

Remove double black



Fixing RBT Violations

Case flow

Case 1

Case 2

resolves (x red-and-black)

Case 3

Case 4

Case 2

Case 2

loops at most $O(\log n)$ times

Case 3

Case 4

Case 4

resolves (x set to root)

All paths through case handling are $O(\log n)$

⇒ fix_delete_violations () is $O(\log n)$

⇒ Red-Black Tree deletion is $O(\log n)$

Reference:

T. Cormen et al., "Introduction to algorithms", Chap 13.4, MIT press (2022)