

Heapsort and Binary Heaps

Heapsort

Efficient **in-place** **sorting** with a **binary heap**

Introduced by **J. Williams** (1964) 

In-place adaptation by R. Floyd (1964) 

Note: Heapsort not a **stable sort** (sorting does **not preserve** input order of identical keys)

Heapsort complexity (for n data items)

Average case: $\rightarrow O(n \log n)$ **same as quicksort**

Worst case: $\rightarrow O(n \log n)$ **better than quicksort**

Storage: $\Theta(n)$ for items, $O(1)$ for sort (**in-place**)

Example **heapsort** applications:

Linux kernel avoid quicksort $O(n^2)$ **worst-case**

Introsort hybrid **quicksort** **heapsort** **insertion sort**

References/Notes/Image credits:

J. Williams, "Algorithm 232 - Heapsort", Communications of the ACM (1964)

(J. Williams photo) <https://ottawacitizen.remembering.ca/obituary/j-w-j-williams-1065926154>

(adaptation of heapsort) R. Floyd, "Algorithm 245 - Treesort 3", Communications of the ACM (1964)

(R. Floyd photo) https://wiki.alquds.edu/?query=File:Robert_W._Floyd.jpg

(Linux kernel sort.c) <https://github.com/torvalds/linux/blob/master/lib/sort.c> (last accessed 27/12/2022)

<https://en.wikipedia.org/wiki/Introsort>

Max and Min Binary Heaps


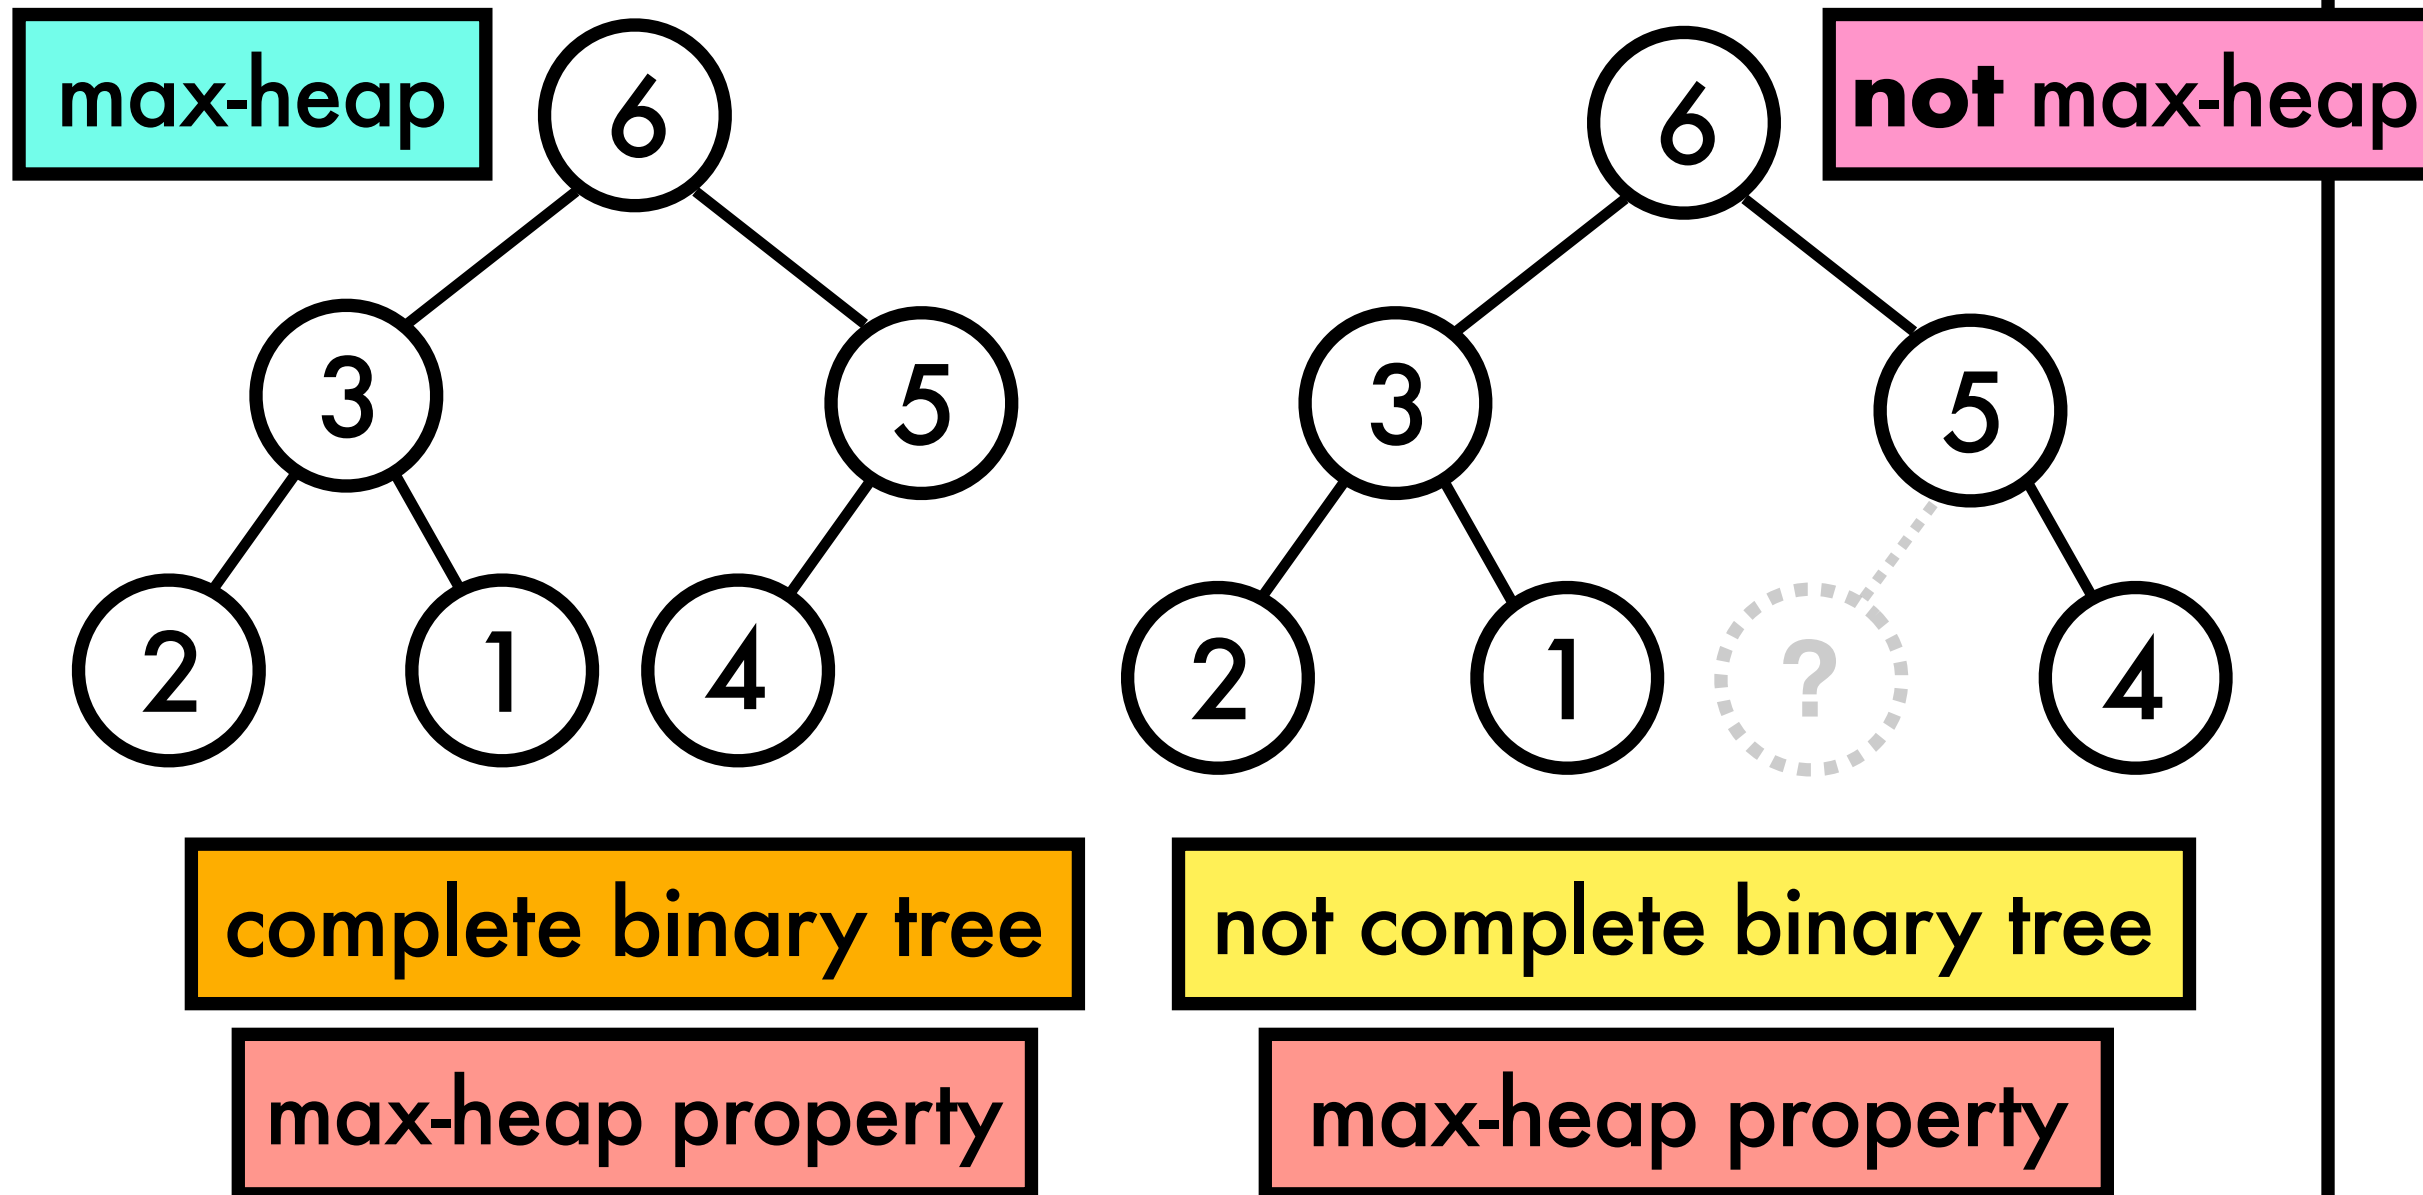
A (binary) **max-heap** is a **binary tree** satisfying:

Shape property: the binary tree is **complete** - all levels except the last are full. If the last is not full, it is filled from **left to right**

Max-heap property: the key of each node is \geq to the keys of its children

A **min-heap** is similar - it satisfies the same **shape property** but with a **min-heap property**:

Min-heap property: the key of each node is \leq the keys of its children



Representing Binary Heaps with Arrays

A heap of `heap_size` nodes can be represented

as an **array**, A of length n with $n \geq \text{heap_size}$

We'll focus on the **max-heap** (used in **heapsort**)

Parent-child relations represented with **indices**

The heap **root** corresponds to array index 0


For node with index i :

- the **parent** has index $\lfloor (i - 1) / 2 \rfloor$
- the **left child** has index $2i + 1$
- the **right child** has index $2i + 2$


Max-heap property:

$$A[i] \leq A[\text{parent}(i)] \text{ for } 1 \leq i < \text{heap_size}$$


max-heap



binary tree representation



array representation



hybrid visualisation

References:

(0-indexed binary heap) https://en.wikipedia.org/wiki/Binary_heap

(array representation) T. Cormen et al., "Introduction to algorithms", Chap 6.1, MIT press (2022)

Binary Heap Height

Binary heap height is $\Theta(\log n)$

The height, h , of a heap with n keys is $h = \lfloor \log n \rfloor$ i.e. $\Theta(\log n)$

Proof:

Due to the **shape property**, the heap is a **complete** binary tree

A heap of height h has **at least**:

$$1 + 2 + 4 + \dots + 2^{h-1} + 1 = 2^h - 1 + 1 = 2^h \text{ nodes}$$

A heap of height h has **at most**:


$$2^{h+1} - 1 \text{ nodes}$$

$$\Rightarrow 2^h \leq n \leq 2^{h+1} - 1$$

Note:


$$\log n < \log(n + 1)$$

$$\Rightarrow h \leq \log n \text{ and } \log(n + 1) \leq h + 1$$



at least 2^h nodes

one node at lowest level



at most $2^{h+1} - 1$ nodes

lowest level is full

$$x \leq y < x + 1, x \in \mathbb{Z}, y \in \mathbb{R} \\ \Rightarrow x = \lfloor y \rfloor$$

$$\Rightarrow h = \lfloor \log n \rfloor$$

References:

M. T. Goodrich et al., "Algorithm design and applications", Chap. 5 (2015)

Fixing Max-Heap Violations

To maintain the **max-heap property**, we fix **heap violations** ("bubbling" violating keys down heap)

Achieved with `max_heapify()` helper function


Assumes left and right child are valid **max-heaps**

```
def max_heapify(A, heap_size, i):  
    left = left_child(i)  
    right = right_child(i)  
    max_i = i  
    if left < heap_size and A[left] > A[max_i]:  
        max_i = left  
    if right < heap_size and A[right] > A[max_i]:  
        max_i = right  
    if max_i != i:  
        A[i], A[max_i] = A[max_i], A[i]  
        max_heapify(A, heap_size, max_i)
```

```
def left_child(i):  
    return 2 * i + 1
```


```
def right_child(i):  
    return 2 * i + 2
```

Before



`max_heapify(A, 6, i=0)`

After



Complexity $O(h) = O(\log n)$ (h is heap height)

References:

T. Cormen et al., "Introduction to algorithms", Chap 6.2, MIT press (2022)





Building a Max-Heap

Somewhat remarkably, we can build a **max-heap** from an unsorted array in **linear time**

Note that each **leaf node** is a valid max-heap

Strategy: call **max_heapify** on left half of array

```
def build_max_heap(A):  
    heap_size = len(A)  
    for i in range(heap_size // 2 - 1, -1, -1):  
        max_heapify(A, heap_size, i)
```



References:

L. Xinyu, "Elementary Algorithms", Chap. 8 (2022)

T. Cormen et al., "Introduction to algorithms", Chap 6.2, MIT press (2022)

Max-Heap Construction has Linear Complexity

Heap construction is $\Theta(n)$ (CLRS)

`build_max_heap()` contains a for loop from $\lfloor n/2 \rfloor - 1$ to 0 so construction is $\Omega(n)$

To show construction is $O(n)$, we can write the **total cost** as:

$$\sum_{h \in \text{node heights}} \text{num. nodes at height } h \times \text{cost of max_heapify at height } h$$

We then use **three observations**: a heap with n keys has height $\lfloor \log n \rfloor$

a heap with n keys has at most $\lceil n/2^{h+1} \rceil$ nodes at height h

max_heapify is $O(h)$ for node at height h

$$\begin{aligned} \sum_{h=0}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil \cdot Ch &\leq \sum_{h=0}^{\lfloor \log n \rfloor} \frac{n}{2^h} \cdot Ch \quad \text{Using that } \lceil n/2^{h+1} \rceil \leq n/2^h \quad \text{since } \lceil u \rceil \leq 2u \text{ for } u \geq 1/2 \\ &= Cn \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h} \leq Cn \sum_{h=0}^{\infty} \frac{h}{2^h} \leq Cn \cdot \text{constant} = O(n) \end{aligned}$$

References:

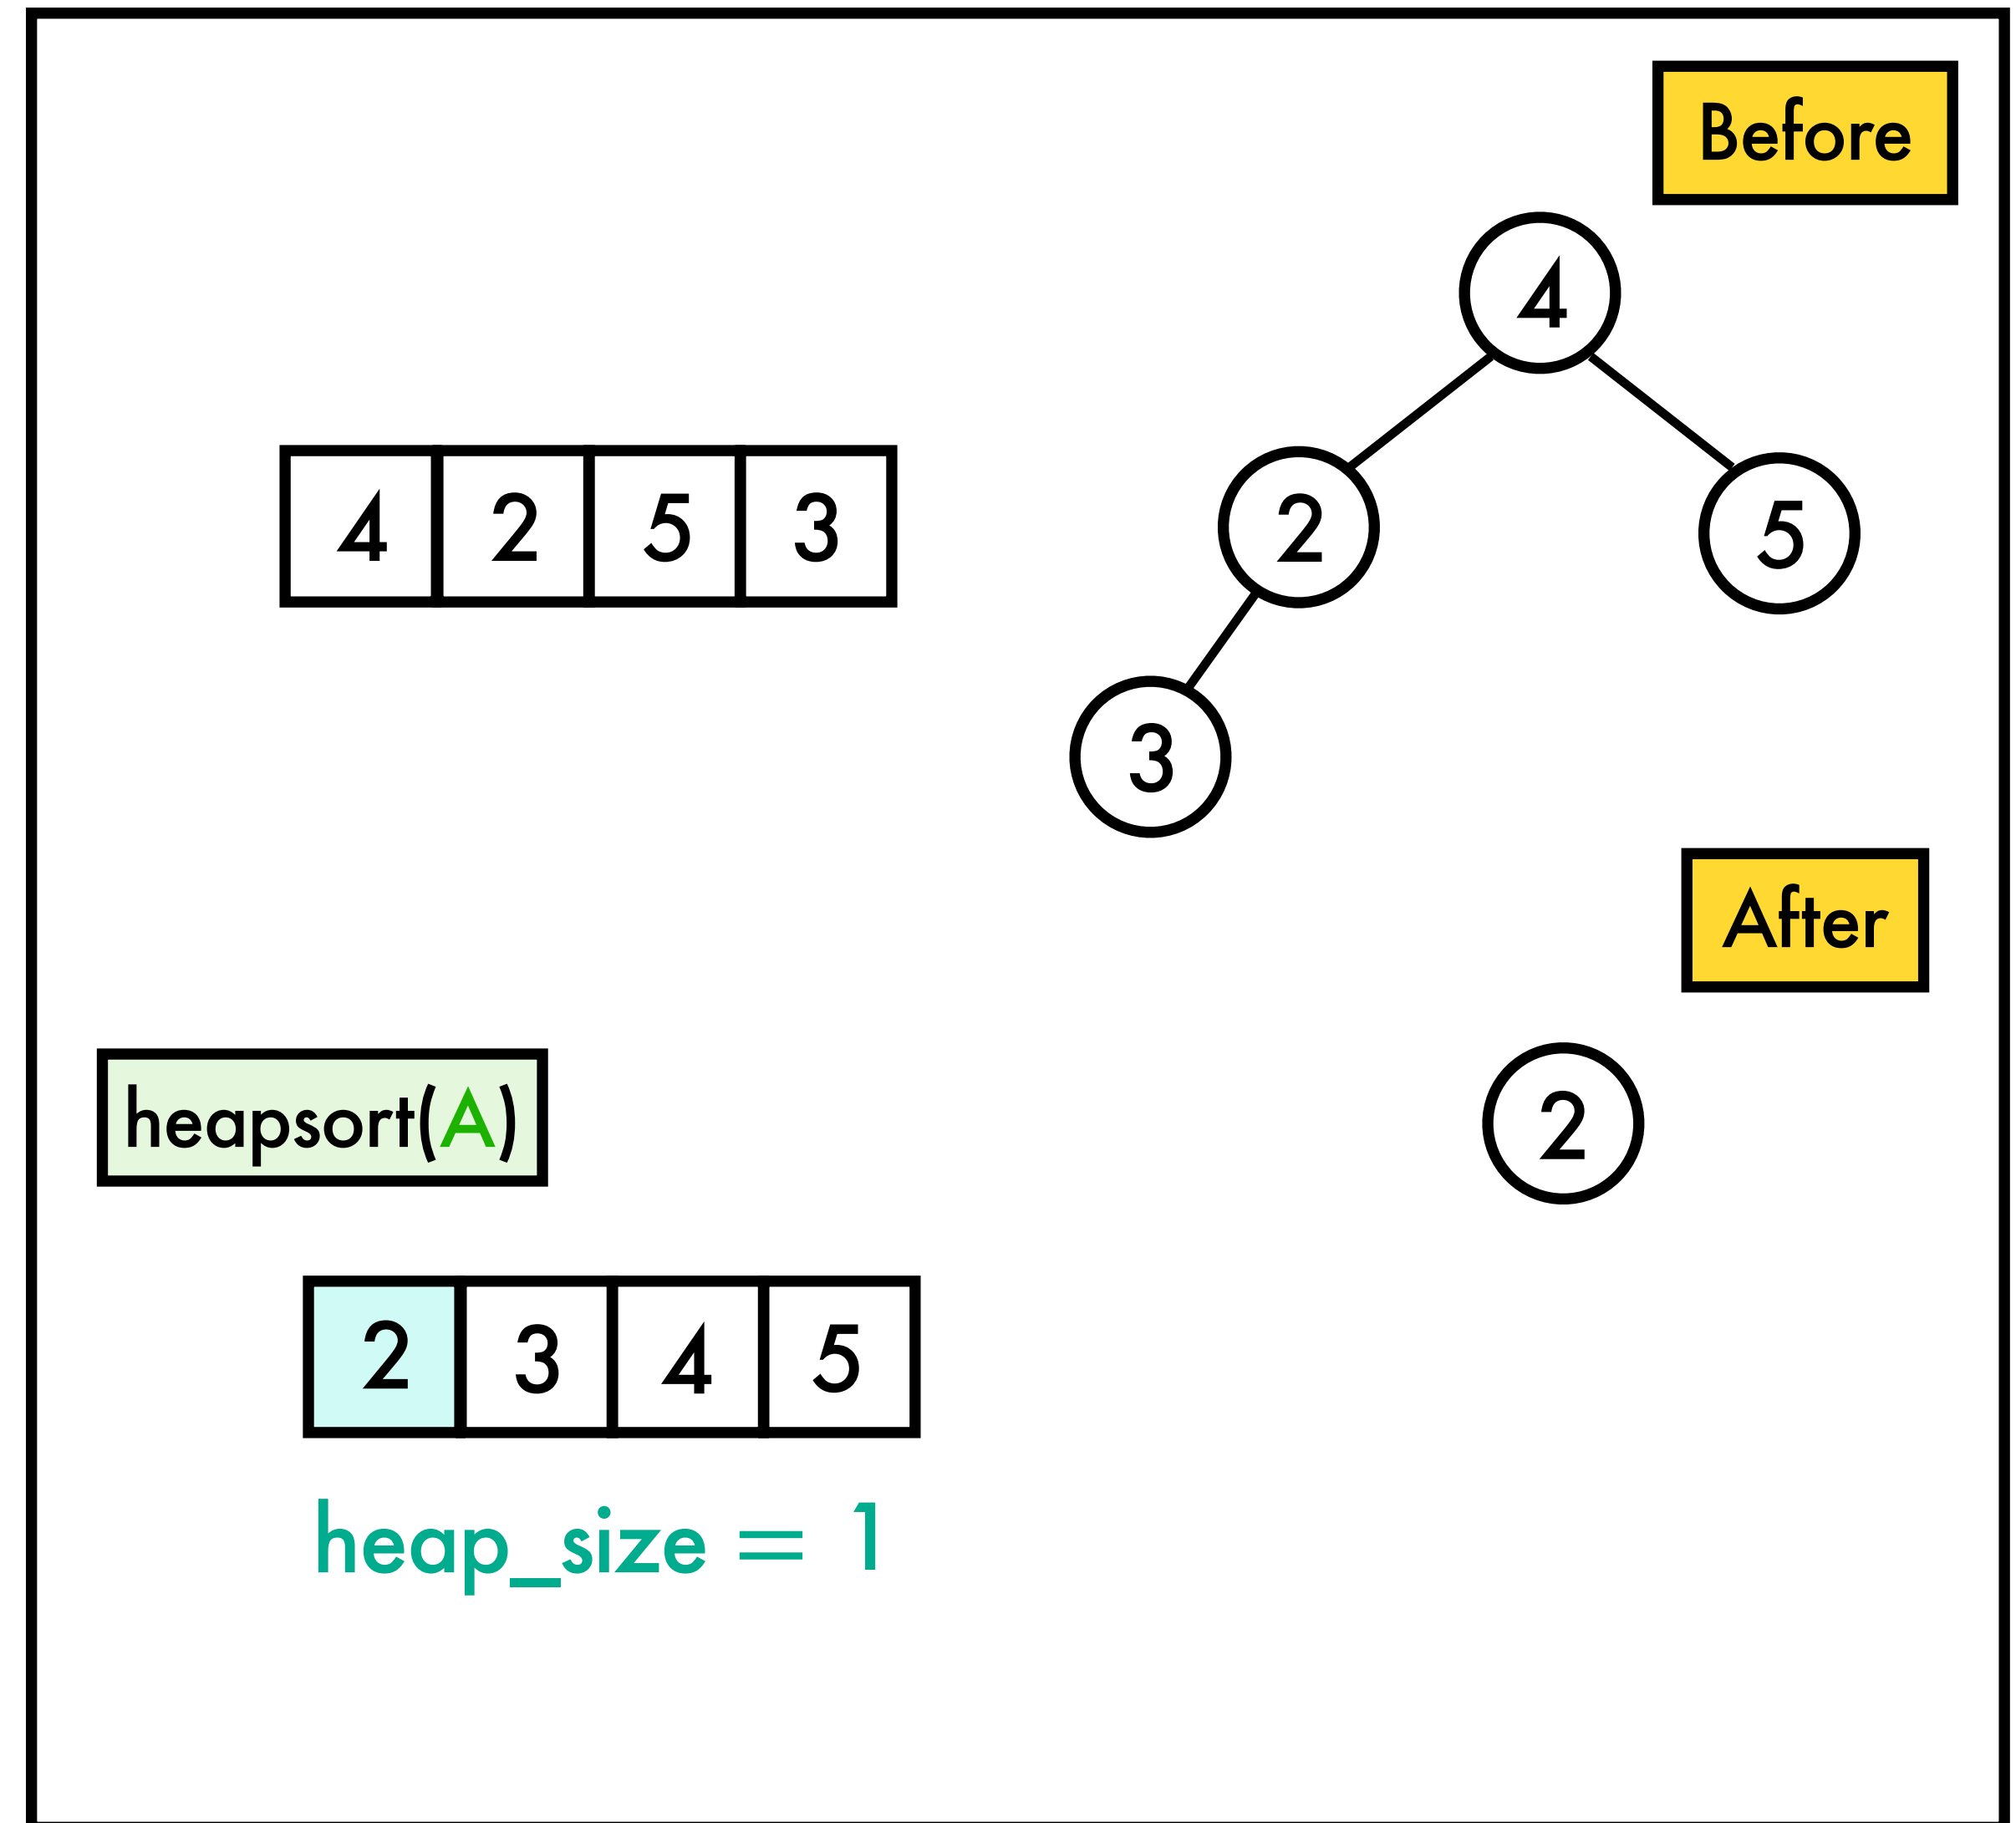
H. Ding, "Proof for The Complexity of Building A Heap", https://www.cse.msu.edu/~huding/331material/timecomplexity_for_heap.pdf

(CLRS) T. Cormen et al., "Introduction to algorithms", Chap 6.3, MIT press (2022)

Heapsort

```
def heapsort(A): # Floyd variant (in-place)
    build_max_heap(A)
    heap_size = len(A)
    while heap_size > 1:
        A[heap_size - 1], A[0] = A[0], A[heap_size - 1]
        heap_size = heap_size - 1
        max_heapify(A, heap_size, 0)
```

Complexity $O(n \log n)$



References:

L. Xinyu, "Elementary Algorithms", Chap. 8 (2022)

(CLRS) T. Cormen et al., "Introduction to algorithms", Chap 6.4, MIT press (2022)

Application of Heaps: Priority Queues

Priority Queues

Beyond **heapsort**, binary heaps are also often used to implement **Priority Queues (PQs)**

Priority Queue is an **Abstract Data Type**

A **Priority Queue** contains a collection of values each associated with a key

As with **heaps** there are **max-PQs** and **min-PQs**

Max-Priority Queue Interface

`get_max()` - return value with largest key

`pop_max()` - pop & return value with largest key

`insert(value, key)` - insert **value** with a key of **key** into the priority queue

`increase_key(value, key)` increase the key associated with **value** to **key**

Min-Priority Queue

insert

get_min

pop_min

decrease_key

Some applications of Priority Queues

Bandwidth manager (max-PQ)

Job scheduler (max-PQ)

Dijkstra's algorithm (min-PQ)

Huffman Coding (min-PQ)

References:

T. Cormen et al., "Introduction to algorithms", Chap 6.5, MIT press (2022)

(Applications of priority queues) https://en.wikipedia.org/wiki/Priority_queue#Applications

Max Priority Queues - get_max()/pop_max()

Assumptions

Keys/values stored in the **priority queue** as dicts

`{"key": key, "value": value}`

Values stored in the **priority queue** are **unique**

Class `MaxPQ` has attributes `heap_size` and `A`

```
def get_max(self): # self is a MaxPQ
    return self.A[0]["value"]
```

Complexity $\Theta(1)$

```
def pop_max(self): # self is a maxPQ
    max_value = self.get_max()
    self.A[0] = self.A[self.heap_size - 1]
    self.heap_size -= 1
    max_heapify(self.A, self.heap_size, 0)
    return max_value
```


Complexity $O(\log n)$

Before

Keys

6	3	5	2	0
B	A	C	L	M

Values




After

Keys

5	3	0	2	6
M	A	C	L	B

Values



max_value = B

References:

L. Xinyu, "Elementary Algorithms", Chap. 8 (2022)

T. Cormen et al., "Introduction to algorithms", Chap 6.5, MIT press (2022)

Max Priority Queues - increase_key()

```
def increase_key(self, key, value): # self is a MaxPQ
    # locate index of `value` in underlying array using an
    # auxiliary data structure (e.g. a red-black tree)
    i = self.value2index[value]
    assert key >= self.A[i]["key"], "requested to decrease key"
    self.A[i]["key"] = key # increase the key
    while i > 0 and self.A[i]["key"] > self.A[parent(i)]["key"]:
        self.A[i], self.A[parent(i)] = self.A[parent(i)], self.A[i]
        i = parent(i)
```


Complexity $O(\log n)$

Before

Keys

6	3	5	2	0
B	A	C	L	M

Values




After

Keys

7	6	5	3	0
L	B	C	A	M

Values



`self.increase_key(7, "L")`

References:

L. Xinyu, "Elementary Algorithms", Chap. 8 (2022)

T. Cormen et al., "Introduction to algorithms", Chap 6.5, MIT press (2022)

Max Priority Queues - insert()

```
def insert(self, key, value): # self is a MaxPQ
    if self.heap_size == len(self.A):
        # expand underlying array to avoid heap overflow
        self.A.append(None)
    tmp_key = float("-inf")
    self.A[self.heap_size] = {"key": tmp_key, "value": value}
    self.heap_size += 1
    self.increase_key(key, value)
```

Complexity $O(\log n)$

Important **Max** Priority Queue operations

pop_max	$O(\log n)$	insert	$O(\log n)$
---------	-------------	--------	-------------

Important **Min** Priority Queue operations


pop_min	$O(\log n)$	insert	$O(\log n)$
---------	-------------	--------	-------------

Before

Keys

6	3	5	2
B	A	C	L

Values




After

Keys

8	6	5	2	3
G	B	C	L	A

Values



self.insert(8, "G")

References:

L. Xinyu, "Elementary Algorithms", Chap. 8 (2022)

T. Cormen et al., "Introduction to algorithms", Chap 6.5, MIT press (2022)

M. T. Goodrich et al., "Algorithm design and applications", Chap. 5 (2015)

Quicksort

Quicksort

Widely-used, fast, in-place **sorting algorithm**

Devised by **Tony Hoare** in Moscow (1959) 

Won wager with boss (who advocated **shellsort**)

Published in ALGOL using **recursion** (1961)

Note: Quicksort not a **stable sort** (sorting does **not preserve** input order of identical keys)

Quicksort complexity (for n data items)

Average case: $\rightarrow O(n \log n)$

Worst-case: $\rightarrow O(n^2)$

Storage: $\Theta(n)$ for items, sort: $O(\log n)$ avg. case,

$O(n)$ worst case **stack usage** $O(\log n)$ with tail call optim.

Quicksort implementations are typically **cache-friendly** (linear scans suit caches)

Employed in **hybrids** (e.g. **Introsort** - falls back to heapsort to avoid $O(n^2)$ worst-case behaviour)

References/Notes/Image credits:

C. A. R. Hoare, "Algorithm 64: quicksort", *Communications of the ACM* (1961)

(Portrait of Hoare) <http://curation.cs.manchester.ac.uk/Turing100/www.turing100.manchester.ac.uk/speakers/invited-list/11-speakers/39-hoare.html>

(Quicksort history) <https://anothercasualcoder.blogspot.com/2015/03/my-quickshort-interview-with-sir-tony.html>

(Elliot Brothers Ltd image) [https://en.wikipedia.org/wiki/Elliott_Brothers_\(computer_company\)#/media/File:Elliott_Sector.jpg](https://en.wikipedia.org/wiki/Elliott_Brothers_(computer_company)#/media/File:Elliott_Sector.jpg)

Storage complexity of quicksort: <https://stackoverflow.com/a/29746572> (technically you can guarantee $O(\log n)$ worst case memory if you have tail call optimisation)

(Introsort) <https://en.wikipedia.org/wiki/Introsort>

Quicksort Overview

Quicksort

Uses a **recursive**, **divide-and-conquer** strategy:

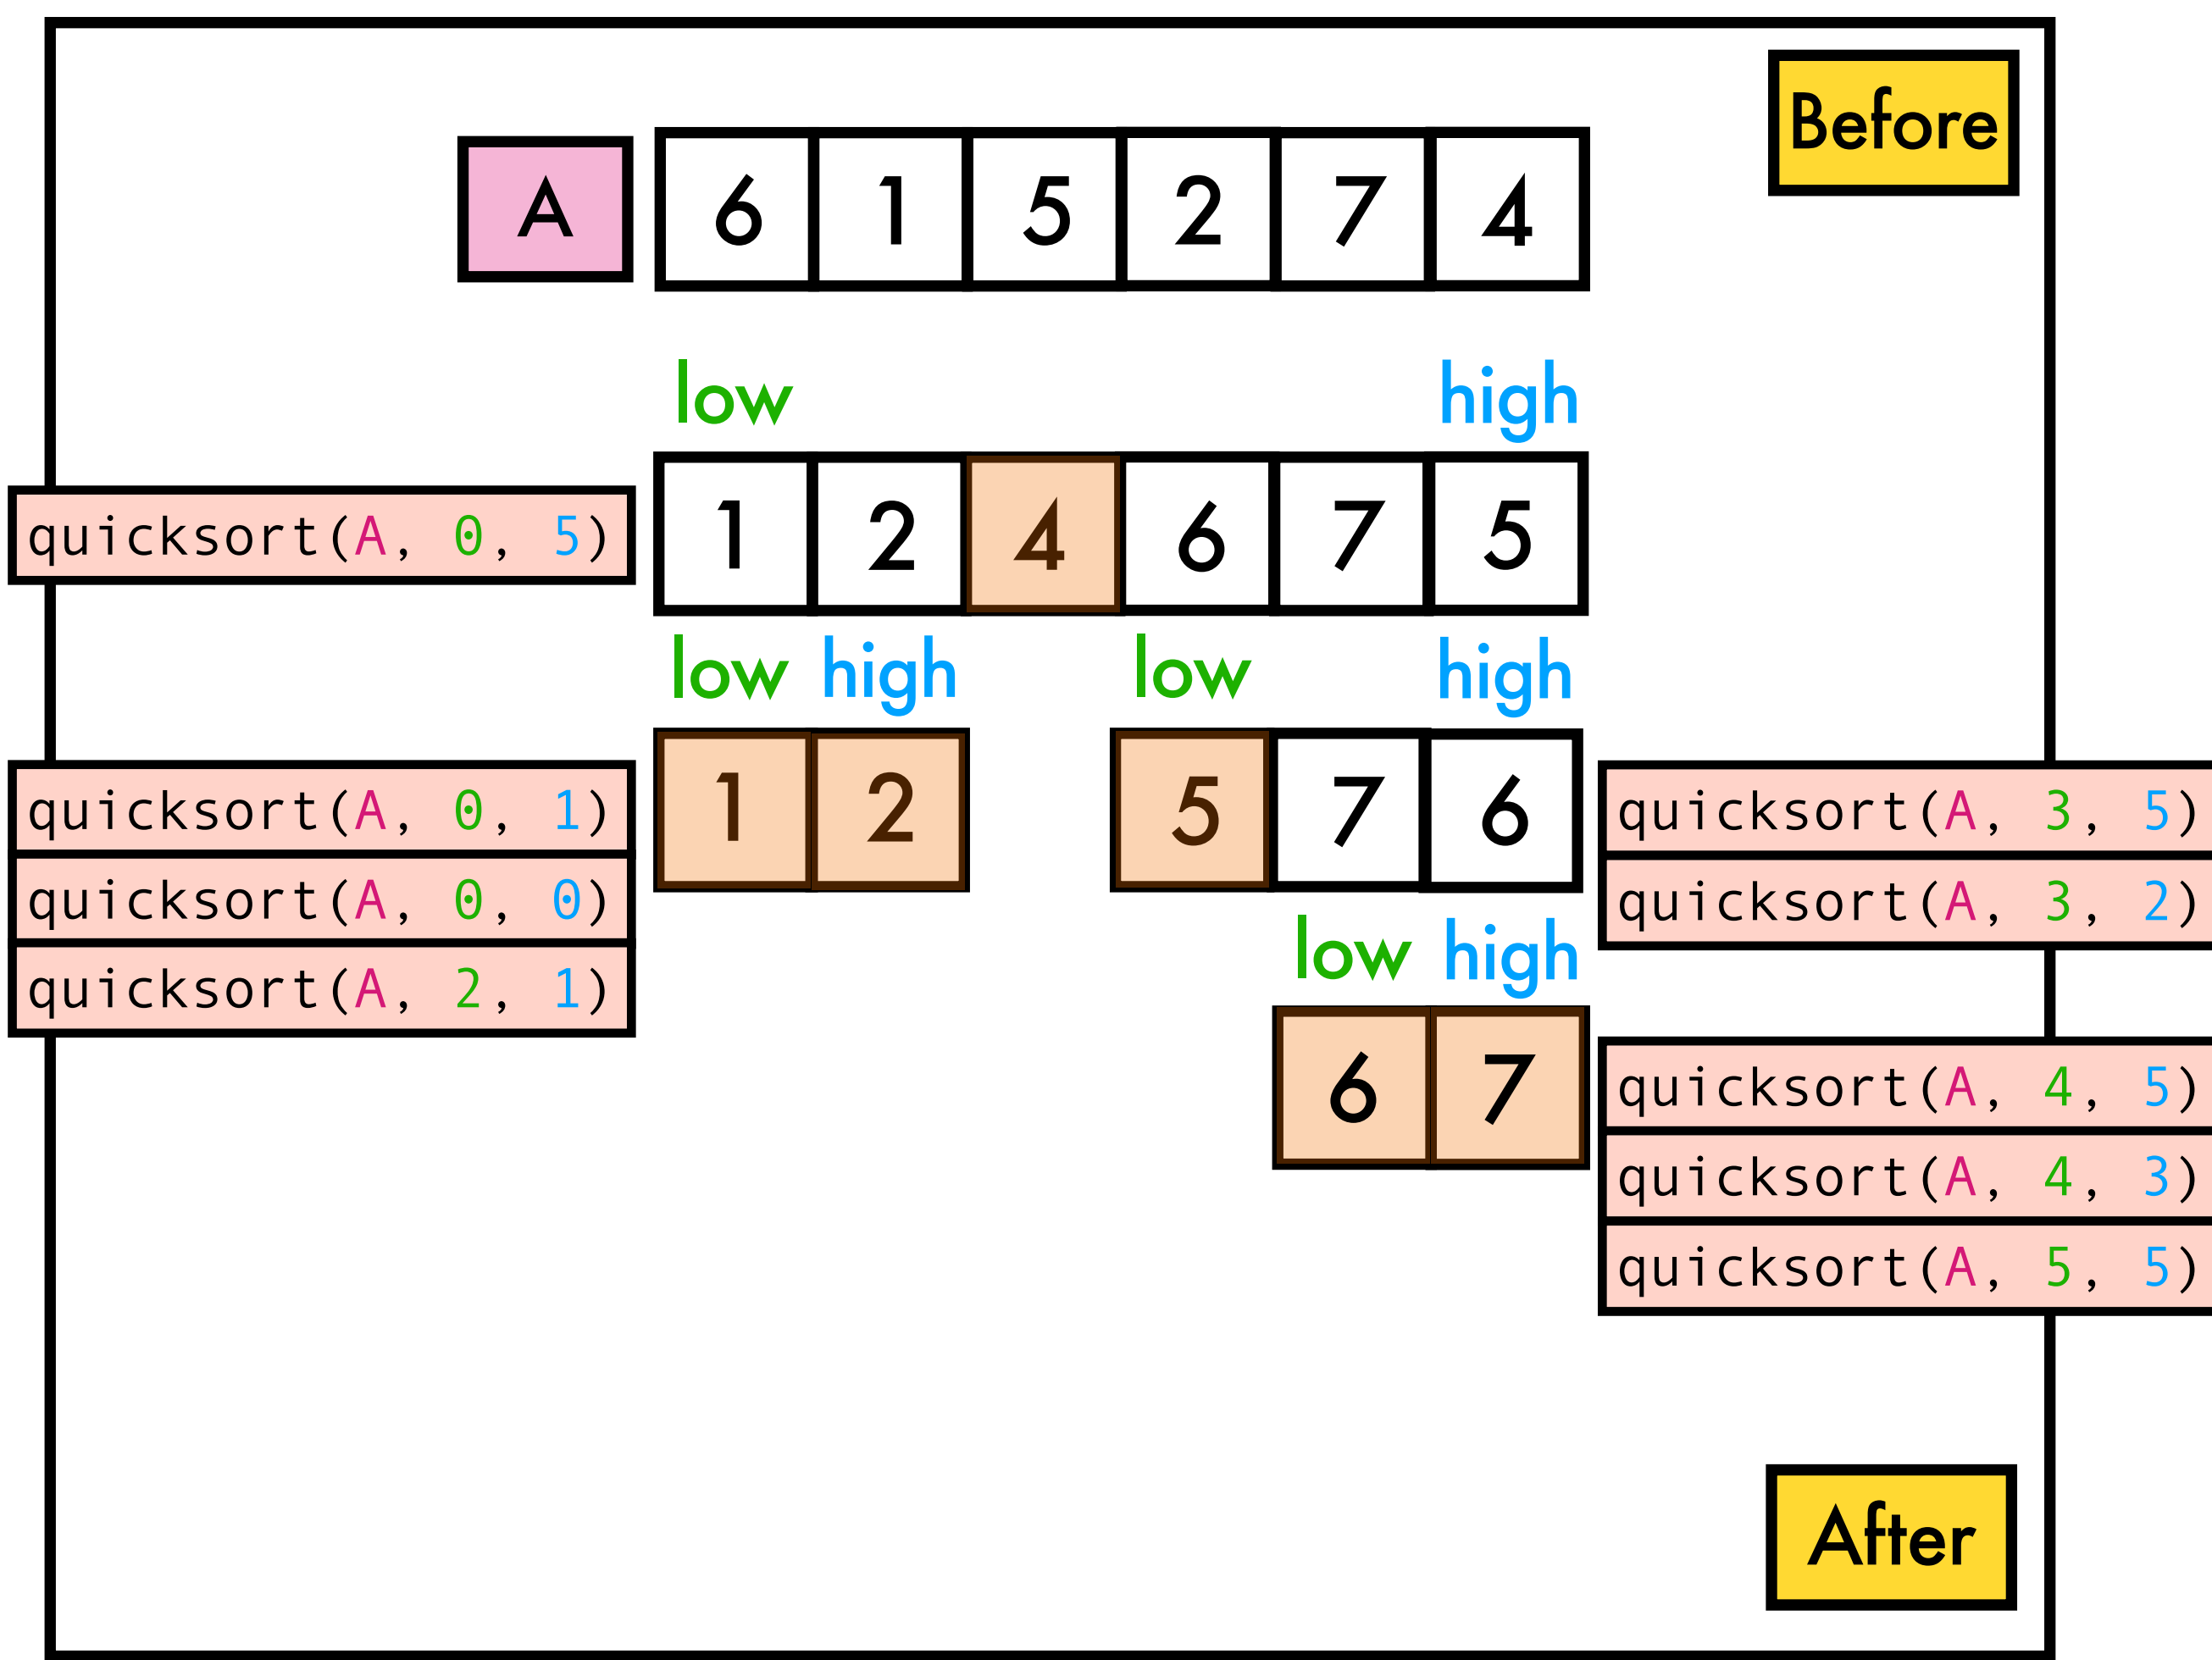
- 1. Select a **pivot** from the array
- 2. **Partition** array into 3 subarrays arranged as:



- 3. Recursively **quicksort** the **first** and **last** subarrays



```
def quicksort(A, low, high):  
    if low < high:  
        pivot = partition(A, low, high)  
        quicksort(A, low, pivot-1) # left subarray  
        quicksort(A, pivot+1, high) # right subarray
```



References:
J. Erickson, "Algorithms", Chap. 2, <http://algorithms.wtf/> (2019)
T. Cormen et al., "Introduction to algorithms", Chap 7.1, MIT press (2022)
L. Xinyu, "Elementary Algorithms", Chap. 13 (2022)


Lomuto Partition

Our first partition function (Lomuto)

partition() does most of the work in **quicksort**

```
def partition(A, low, high): # lomuto partition
    pivot = high
    pivot_val = A[pivot]
    i = low
    for j in range(low, high):
        if A[j] <= pivot_val:
            A[i], A[j] = A[j], A[i]
            i = i + 1
    A[i], A[pivot] = A[pivot], A[i]
    return i # new pivot
```

Complexity: $\Theta(n)$ where $n = \text{high} - \text{low} + 1$



References:

- J. Erickson, "Algorithms", Chap. 2, <http://algorithms.wtf/> (2019)
- T. Cormen et al., "Introduction to algorithms", Chap 7.1, MIT press (2022)
- L. Xinyu, "Elementary Algorithms", Chap. 13 (2022)

Worst-Case Analysis

Unbalanced partitions and quadratic complexity

How can things go badly for **quicksort**?

We have seen that `partition()` is $\Theta(n)$

Since quicksort is **recursive**, can write complexity:

$$T(n) = \underbrace{T(r)}_{\text{left subarray}} + \underbrace{T(n-r-1)}_{\text{right subarray}} + \underbrace{\Theta(n)}_{\text{partition}}$$

where r is the **rank** of the **pivot**

If we choose the **pivot** as smallest value ($r = 0$)

$$T(n) = T(0) + T(n-1) + \Theta(n)$$

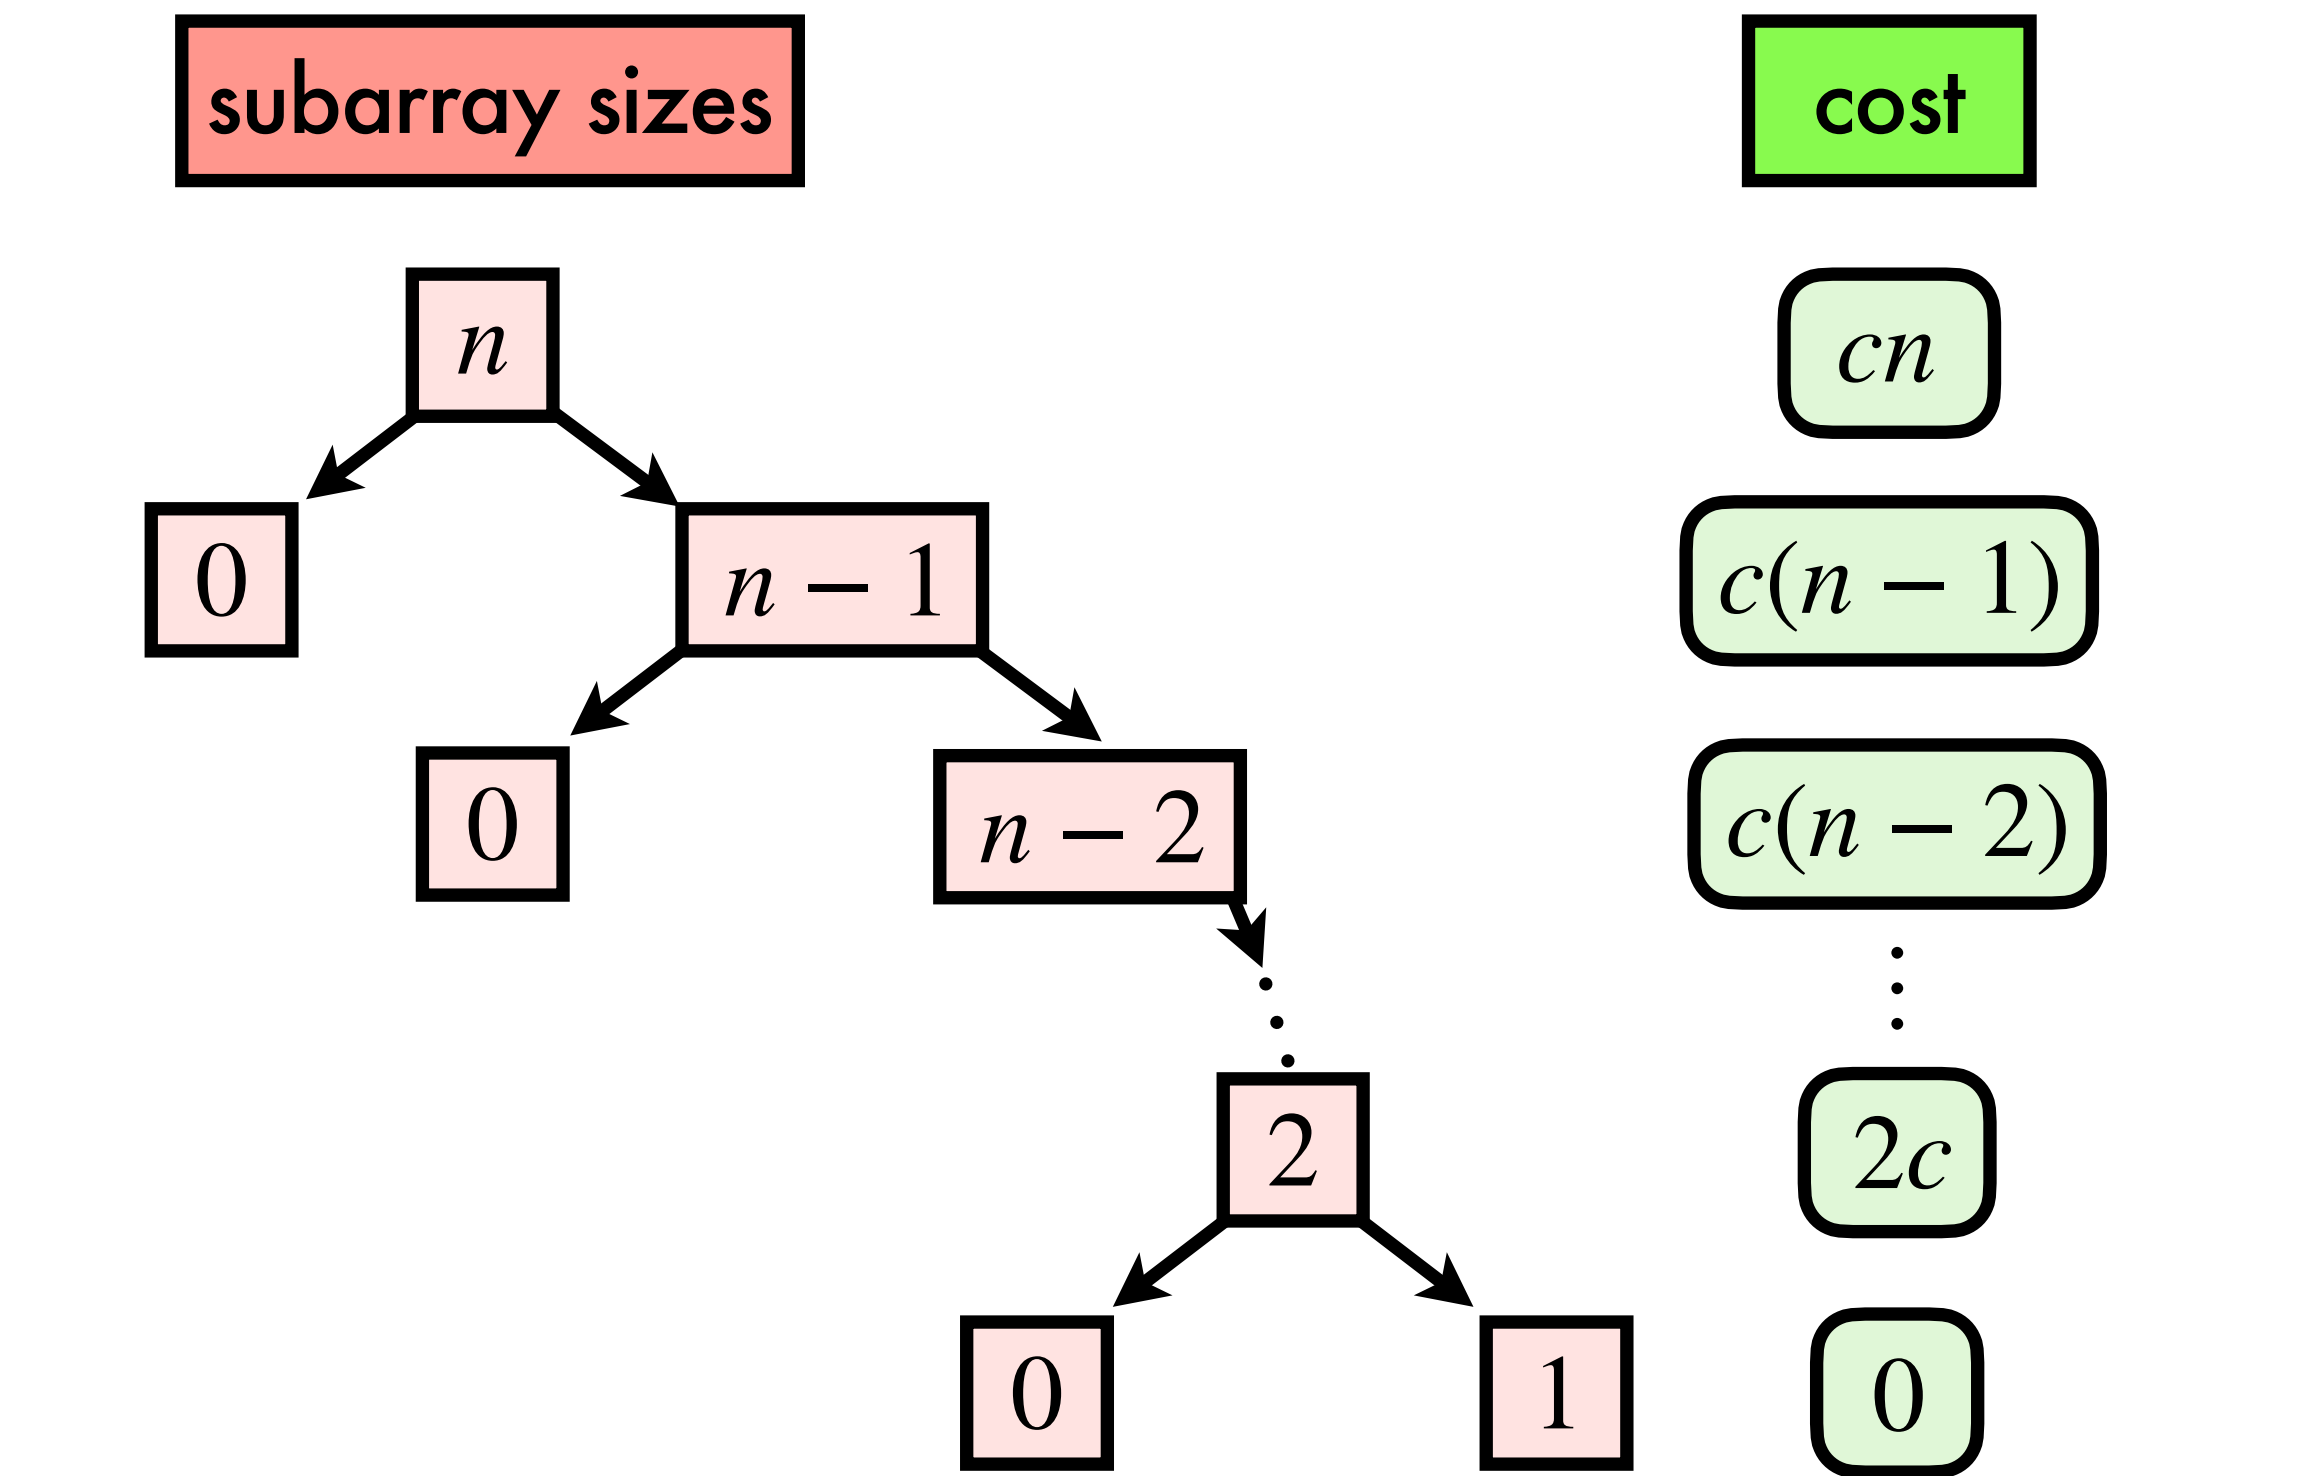
$T(0) = O(1)$ (**empty subarray** - no comparisons)

$$T(n) = T(n-1) + \Theta(n) \implies T(n) = \Theta(n^2)$$

$T(n) = \Theta(n^2)$ for fixed r not dependant on n

Illustrate runtime complexity with $r = 0$

Partition cost $\Theta(n)$ as cn for some **constant** c



$$\text{Total cost} = c(n + (n-1) + (n-2) + \dots + 2 + 0)$$

$$= c \sum_{i=2}^n i = c(n(n+1)/2 - 1) \implies \Theta(n^2)$$

References:

J. Erickson, "Algorithms", Chap. 2, <http://algorithms.wtf/>

T. Cormen et al., "Introduction to algorithms", Chap 7, MIT press (2022)

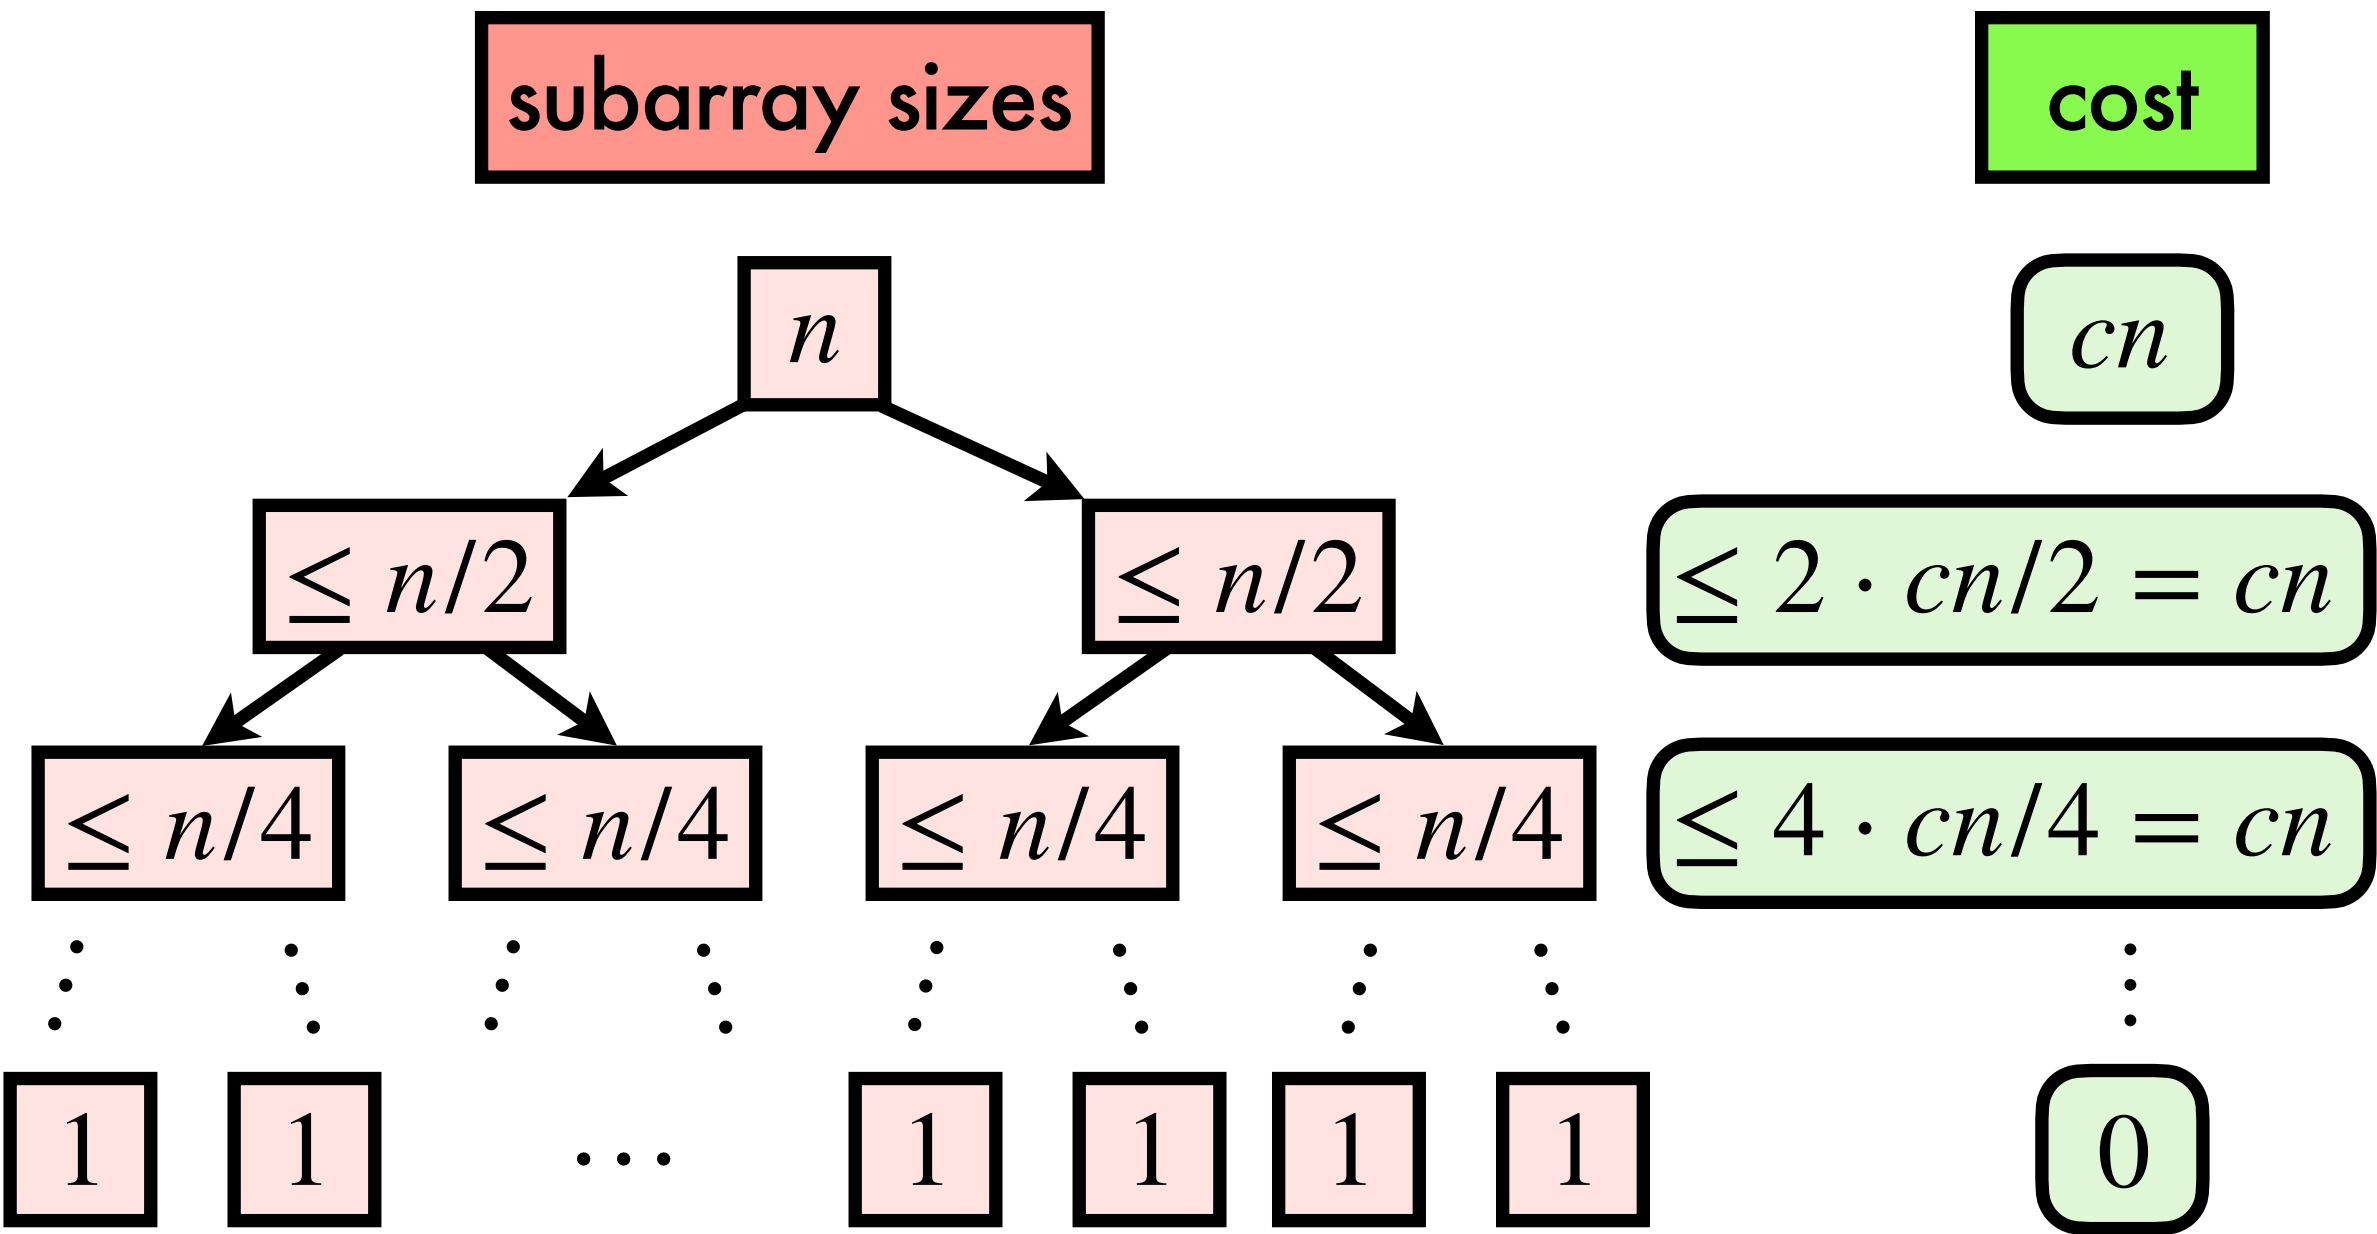
<https://www.khanacademy.org/computing/computer-science/algorithms/quick-sort/a/analysis-of-quicksort>

Balanced Partition Analysis

Quicksort is happy with $r = K \cdot n$ for $0 < K < 1$

By happy, we mean $O(n \log n)$

Consider the simple case when $K = 0.5$:




Tree height is $\log n$ and cost at each level $\leq cn$

Total cost $\leq cn \cdot \log n \implies O(n \log n)$

For complexity, what matters is tree depth growth

Suppose $K = 0.995$ (seems unbalanced!)



Tree depth is $\log_{\frac{1000}{995}} n$

So the complexity is still $O(n \log n)$

Some intuition for why "most" partitions will produce $O(n \log n)$ behaviour

References:
Fig. based on <https://www.khanacademy.org/computing/computer-science/algorithms/quick-sort/a/analysis-of-quicksort>

Engineering Improvements

Behaviour on concerningly common inputs

If array **already sorted**, Lomuto partitioning has **worst case** behaviour $O(n^2)$

If all elements the same, Lomuto is also $O(n^2)$

"Median-of-three"

A **simple heuristic** to get a better pivot:
pick **median** of first, middle and last elements
Useful on **sorted input arrays** (and other inputs)

Does not rule out worst case behaviour

References:

(Median-of-3) R. Sedgewick, "Algorithms in C: Fundamentals, Data Structures, Sorting, Searching" (1998)
(Proof for randomised quicksort) <https://www.cs.cmu.edu/~avrim/451f11/lectures/lect0906.pdf>
T. Cormen et al., "Introduction to algorithms", Chap 7, MIT press (2022)
D. M. McIlroy, "A killer adversary for quicksort", Software: Practice and Experience (1999)
(quickselect) C. A. R. Hoare, "Algorithm 65: Find", Communications of the ACM (1961)

Randomised quicksort

Pick pivot **uniformly at random** and swap it with:
element at index $n - 1$ (if using **Lomuto** partition)

Expected runtime $O(n \log n)$

The full **proof** is quite involved (see refs)

Worst case becomes **very unlikely** (widely used)

Provoking $O(n^2)$ behaviour

"Killer adversary for quicksort" (McIlroy, 1999)

Decide ordering of elements **lazily** during sorting

Achieving $O(n \log n)$ worst case

Can ensure $O(n \log n)$ **worst case** quicksort via
 $O(n)$ **median-of-medians** algorithm

*In practice, this is **too slow**, so not widely used*