

# Comparison Sorting Lower Bounds

What is the **theoretically fastest possible** sort speed?

**Comparison sorting algorithms** sort arrays by  
**comparing elements** (with no extra information)

Examples: **Insertion sort** **Heapsort** **Quicksort**

H. Steinhaus considered a sorting puzzle - **ranking all players** in tennis tournaments (1939) 

L. Ford and S. Johnson introduced **merge-insertion sort** with decision tree analysis (1959) 

Runtime complexity results for comparison sorts

**Lower bound for worst input:**  $\rightarrow \Omega(n \log n)$

**Lower bound for average input:**  $\rightarrow \Omega(n \log n)$

$\Rightarrow O(n \log n)$  sorts such as **heapsort** are  
**asymptotically optimal**

Note: **non-comparison** sorts (e.g. Radix sort) can do better

References/Notes/Image credits:

H. Steinhaus, "Mathematical Snapshots" (1939)  
(Image of Steinhaus) [https://en.wikipedia.org/wiki/Hugo\\_Steinhaus#/media/File:Hugo\\_Steinhaus.jpg](https://en.wikipedia.org/wiki/Hugo_Steinhaus#/media/File:Hugo_Steinhaus.jpg)  
(Bumps photo) <https://www.flickr.com/photos/stanbury/42283400832/in/photostream/>

L. Ford and S. Johnson, "A tournament problem", The American Mathematical Monthly (1959)  
(Image of Ford) [https://en.wikipedia.org/wiki/Lester\\_R.\\_Ford#/media/File:Lester\\_R.\\_Ford.gif](https://en.wikipedia.org/wiki/Lester_R._Ford#/media/File:Lester_R._Ford.gif)  
(Image of Johnson) <https://alchetron.com/Selmer-M-Johnson>

# Comparison Sorting Algorithms and Decision Trees

A **comparison sort** ingests an array  $[a_0, \dots, a_{n-1}]$

It only gains information by **comparing pairs**:

"is  $a_i < a_j$ "?

It outputs a **permutation** that orders the items

For our analysis, assume all elements are **distinct**

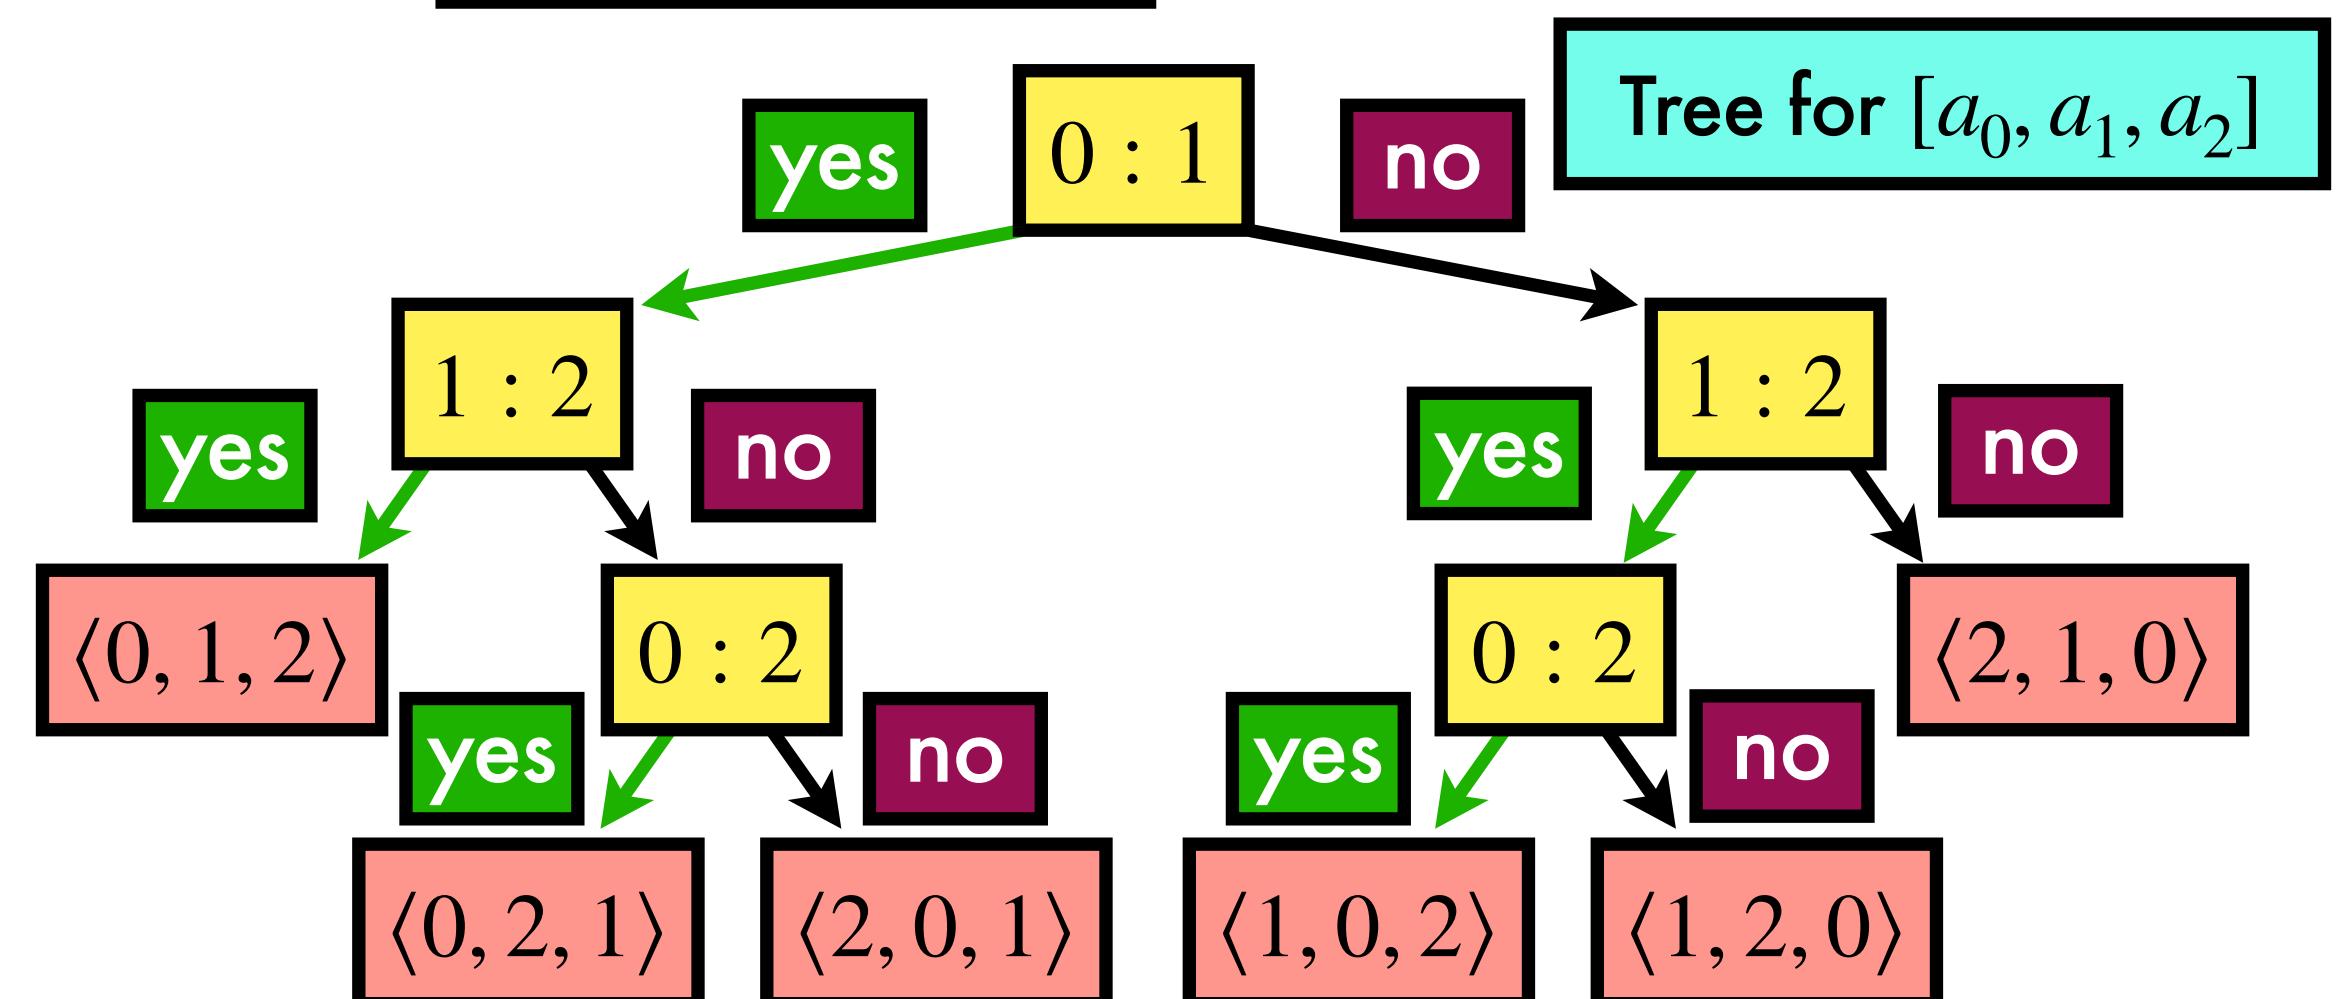
(repeated elements won't affect **lower bounds**)

Decision trees (or "Comparison trees")

Full binary trees where each internal node

$i : j$  represents a **comparison**  $a_i < a_j$

Each leaf  $\langle \pi(0), \dots, \pi(n-1) \rangle$  is an **array permutation**



References:

A. Blum and M. Blum, "Comparison-based Lower Bounds for Sorting", <https://www.cs.cmu.edu/~avrim/451f11/lectures/lect0913.pdf>

D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5.3 (1998)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

J. Erickson, "Lower Bounds", <https://jeffe.cs.illinois.edu/teaching/algorithms/notes/12-lowerbounds.pdf> (2018)

Note: in a "full" binary tree, every node has either 0 or 2 children.

A **comparison sort** follows a path from root to leaf

# Lower Bound On The Worst Case

There are  $n!$  possible permutations (without repetitions) i.e. ways to sort the input

A correct comparison sorting algorithm must be able to produce all of these

→ a comparison sort corresponds to a full binary tree with  $\geq n!$  leaves

Worst case number of comparisons - length of longest simple path from root to leaf

What is the maximum height of a decision tree?

A binary tree with height  $h$  has at most  $2^h$  leaves, so  $n! \leq 2^h$

$$\text{Stirling's approximation: } n! = \left(\frac{n}{e}\right)^n \sqrt{2\pi n} \left(1 + \Theta\left(\frac{1}{n}\right)\right) \geq \left(\frac{n}{e}\right)^n$$

$$h \geq \log(n!) \geq \log\left(\frac{n}{e}\right)^n = n \log n - n \log e = \Omega(n \log n)$$

No comparison sort can be faster than this on its worst case input

References:

- A. Blum and M. Blum, "Comparison-based Lower Bounds for Sorting", <https://www.cs.cmu.edu/~avrim/451f11/lectures/lect0913.pdf>
- D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5.3 (1998)
- T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)
- J. Erickson, "Lower Bounds", <https://jeffe.cs.illinois.edu/teaching/algorithms/notes/12-lowerbounds.pdf> (2018)

# Lower Bound On The Average Case

What is the **smallest** possible **average height** of a decision tree?

The **average height** is smallest when the tree is **completely balanced**

**Completely balanced** means no leaf heights differ by more than 1

If the tree is **completely balanced**, each leaf has depth  $\lceil \log(n!) \rceil$  or  $\lfloor \log(n!) \rfloor$

⇒ smallest possible average height is  $\Omega(n \log n)$

No comparison sort can be faster than this on **average**

## References:

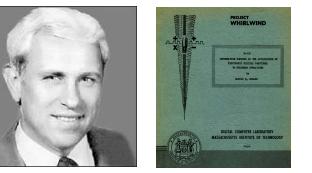
- A. Blum and M. Blum, "Comparison-based Lower Bounds for Sorting", <https://www.cs.cmu.edu/~avrim/451f11/lectures/lect0913.pdf>
- D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5.3 (1998)

# Counting Sort

## Counting Sort

Fast **sorting algorithm** that assumes keys are **integers** in a **fixed range**

Introduced by **H. Seward** (1954)



Not a **comparison sort** (uses key values)

**Stable** - preserves order of repeated elements

## Counting sort complexity

For  $n$  items with **integer keys** (in 0 to  $k - 1$ ):

Running time **for any input**:  $\rightarrow \Theta(n + k)$

Note: counting sort avoids the  **$\Omega(n \log n)$  lower bound** because it is not a comparison sort

**Storage:**  $\Theta(n + k)$  in addition to  $\Theta(n)$  for input

Uses of **counting sort**:

As a building block for **radix sort**

Very fast when  $k = O(n) \rightarrow$  run time is  **$\Theta(n)$**

### References/Notes/Image credits:

H. Seward, "Information sorting in the application of electronic digital computers to business operations", Dissertation MIT (1954)

(Image of H. Seward) <https://www.legacy.com/us/obituaries/bostonglobe/name/harold-seward-obituary?id=20744826>

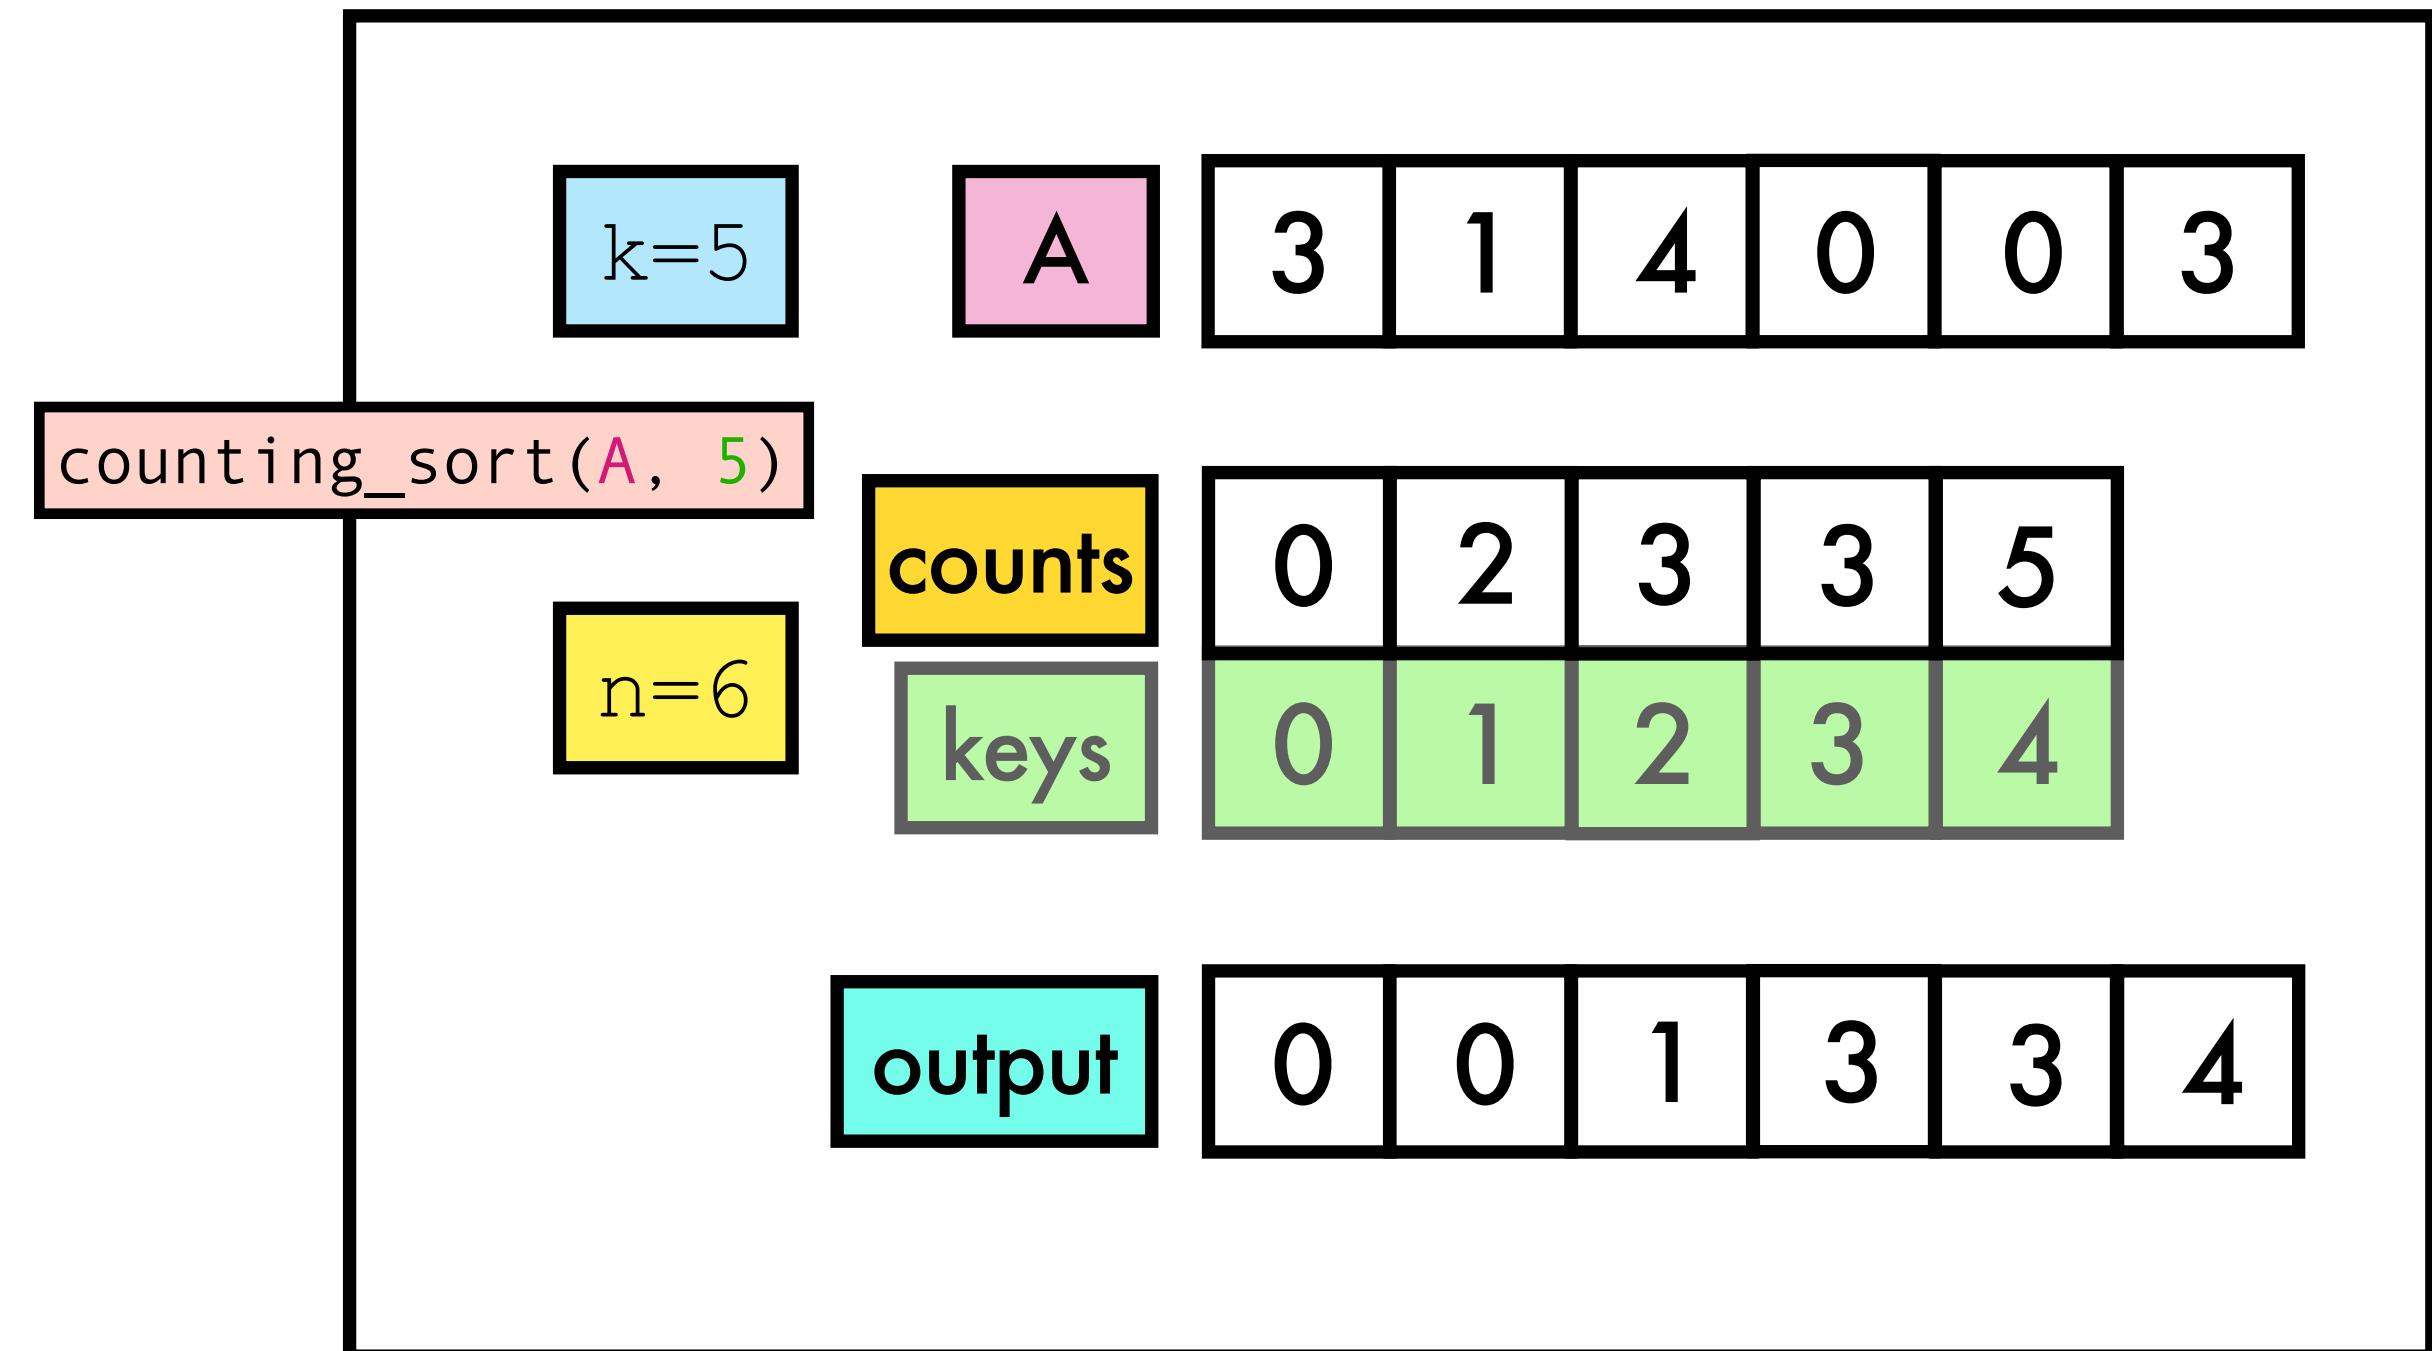
(Image of thesis) [http://bitsavers.trailing-edge.com/pdf/mit/whirlwind/R-series/R-232\\_Information\\_Sorting\\_in\\_the\\_Application\\_of\\_Electronic\\_Digital\\_Computers\\_to\\_Business\\_Operations\\_May54.pdf](http://bitsavers.trailing-edge.com/pdf/mit/whirlwind/R-series/R-232_Information_Sorting_in_the_Application_of_Electronic_Digital_Computers_to_Business_Operations_May54.pdf)

D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5.2 (1998)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

# Counting Sort Implementation

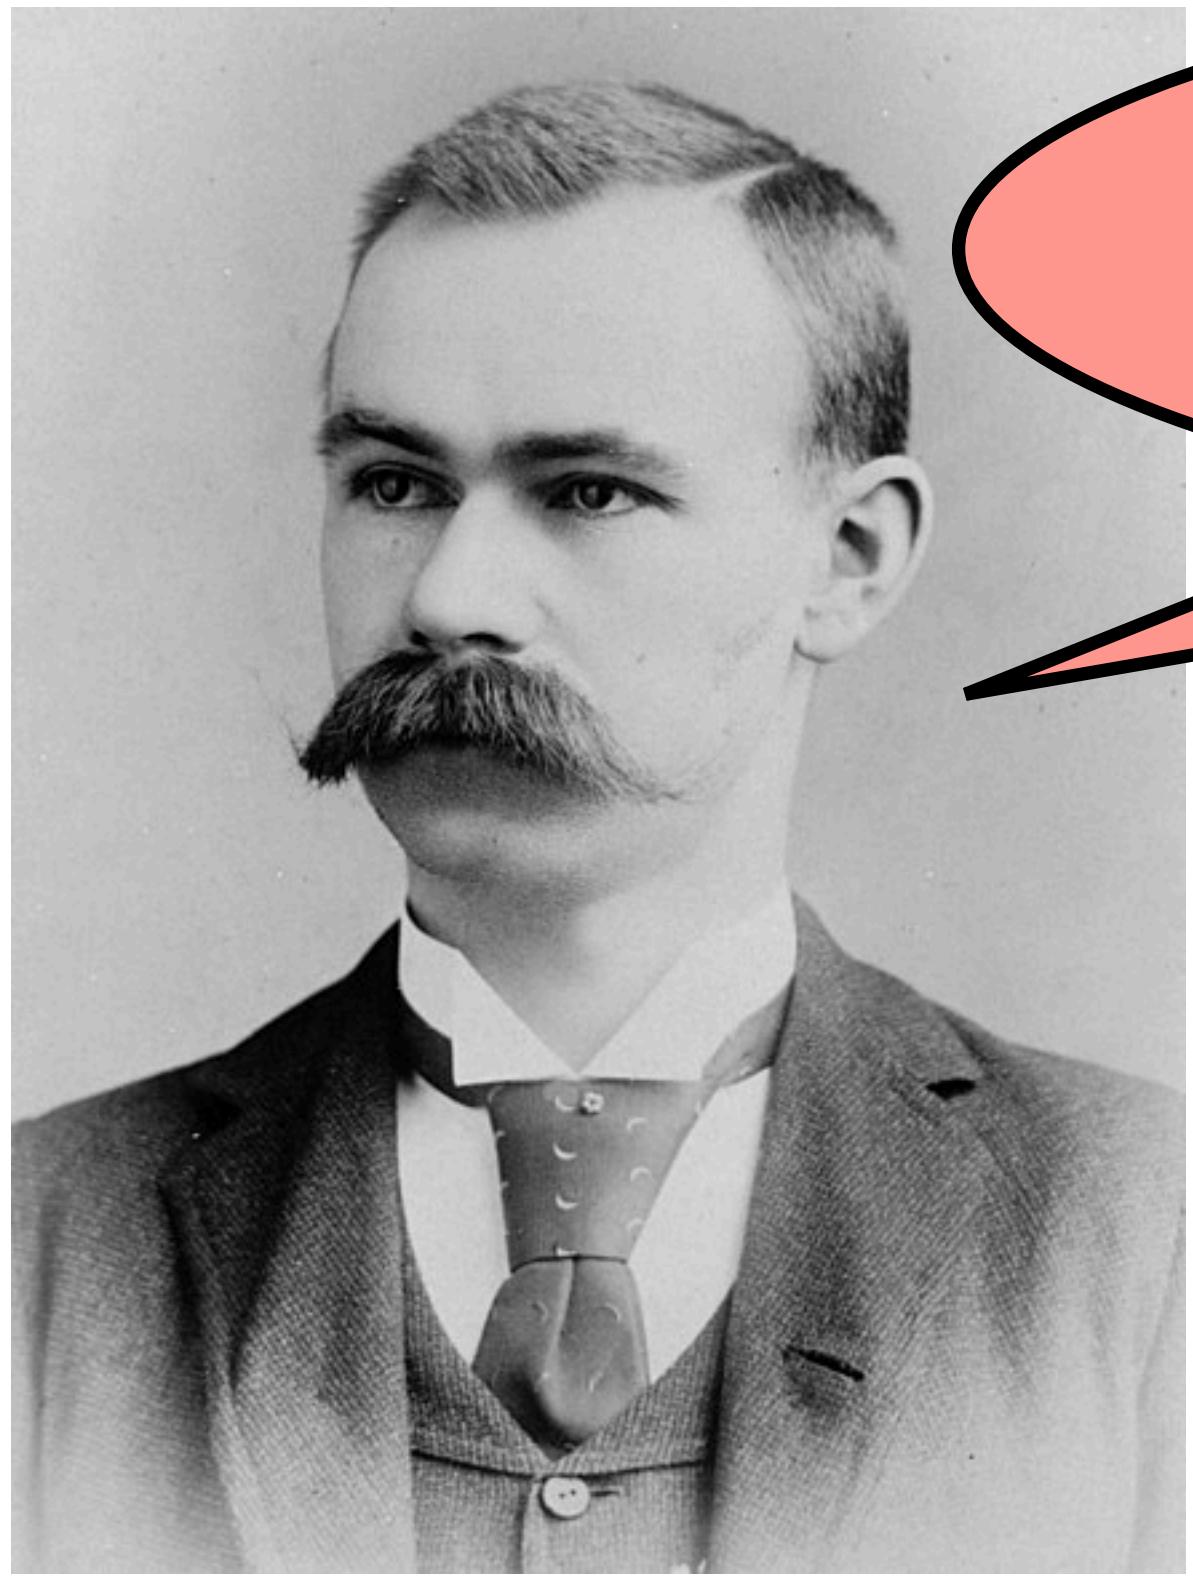
```
def counting_sort(A, k): # keys from 0 to k-1
    n = len(A)
    counts = [0 for _ in range(k)]
    output = [None for _ in range(n)]
    for key in A:
        counts[key] += 1
    for i in range(1, k):
        counts[i] += counts[i-1]
    for key in reversed(A):
        output[counts[key]-1] = key
        counts[key] = counts[key] - 1
    return output
```



References:

[https://en.wikipedia.org/wiki/Counting\\_sort](https://en.wikipedia.org/wiki/Counting_sort)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)



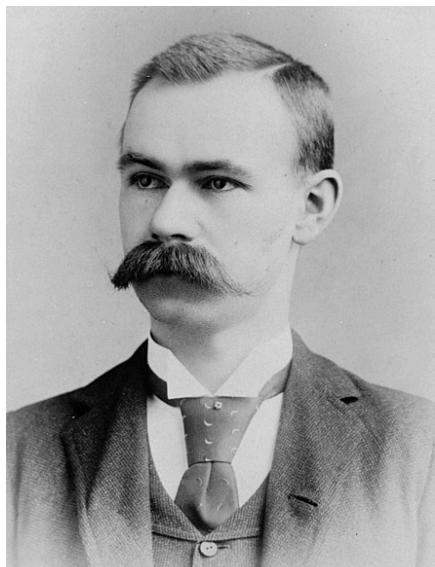
What is the best way to sort data *with machines*?

There is no "best" way to sort.

It depends..

# Radix Sort

valuable for **certain kinds** of keys



History

Runtime complexity

$$\Theta(d(n + k))$$

Properties

```
def radix_sort(A, d, k):  
    ...
```

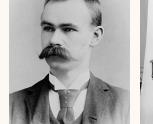
Python Implementation

# Radix Sort

## Radix Sort

Non-comparative sorting algorithm

It orders keys by values, not by comparisons

H. Hollerith card sorting (1901)   IBM

L. J. Comrie published LSD variant (1929) 

H. Seward improved storage (1954) 

Radix sort applies to data that can be sorted

lexicographically (integers, words, cards etc.)

## Radix sort runtime complexity

Two radix sort variants:

Least Significant Digit (LSD)

Most Significant Digit (MSD)

For  $n$  items with  $d$ -digit integer keys (in 0 to  $k - 1$ ):

### LSD Radix Sort

Running time for any input:  $\rightarrow \Theta(d(n + k))$  when using counting sort for inner loop

Storage:  $\Theta(n + k)$  in addition to  $\Theta(dn)$  for input

### MSD Radix Sort (worst case runtime matches LSD)

Average runtime:  $\Theta(n \log_k n)$  on random inputs

Storage:  $\Theta(n + pk)$  in addition to  $\Theta(dn)$  for input,

where  $p$  is length of longest key prefix match

References/Notes/Image credits:

(Early history) <https://www.ibm.com/ibm/history/ibm100/us/en/icons/tabulator/>

D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5 (1998)

(Hollerith image) [https://en.wikipedia.org/wiki/Herman\\_Hollerith#/media/File:Hollerith.jpg](https://en.wikipedia.org/wiki/Herman_Hollerith#/media/File:Hollerith.jpg)

(Hollerith machine) [https://americanhistory.si.edu/collections/search/object/nmah\\_694412](https://americanhistory.si.edu/collections/search/object/nmah_694412)

(CTR logo) [https://en.wikipedia.org/wiki/Computing-Tabulating-Recording\\_Company#/media/File:CTR\\_Company\\_Logo.png](https://en.wikipedia.org/wiki/Computing-Tabulating-Recording_Company#/media/File:CTR_Company_Logo.png)

(IBM logo) [https://en.wikipedia.org/wiki/File:IBM\\_logo.svg](https://en.wikipedia.org/wiki/File:IBM_logo.svg)

L. J. Comrie, "The Hollerith and Powers Tabulating Machines", Transactions of the Office Machinery Users Assoc. (1929)

H. Seward, "Information sorting in the application of electronic digital computers to business operations", Dissertation MIT (1954)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

R. Sedgewick, "MSD Radix Sort", <https://es.coursera.org/lecture/algorithms-part2/msd-radix-sort-gFwG> (2007)

# Least And Most Significant Digit Radix Sorts

Should we **radix sort** digits from **right** to **left** or from **left** to **right**?

Least Significant Digit - LSD (**right** to **left**)

Most Significant Digit - MSD (**left** to **right**)



## LSD radix sort

Assumes all keys have **same length**

Approach used in **card-sorting machines**

Requires a **stable sort** in inner loop

**Iteratively** sorts  $d$ -digit data in  $d$  passes

## MSD radix sort

Supports keys with **different lengths**

Sorts by **recursion** (exits early if  $< 2$  keys share digit)

Runtime depends heavily on the **input data distribution**

Can create lots of (small and inefficient) **subarray sorts**

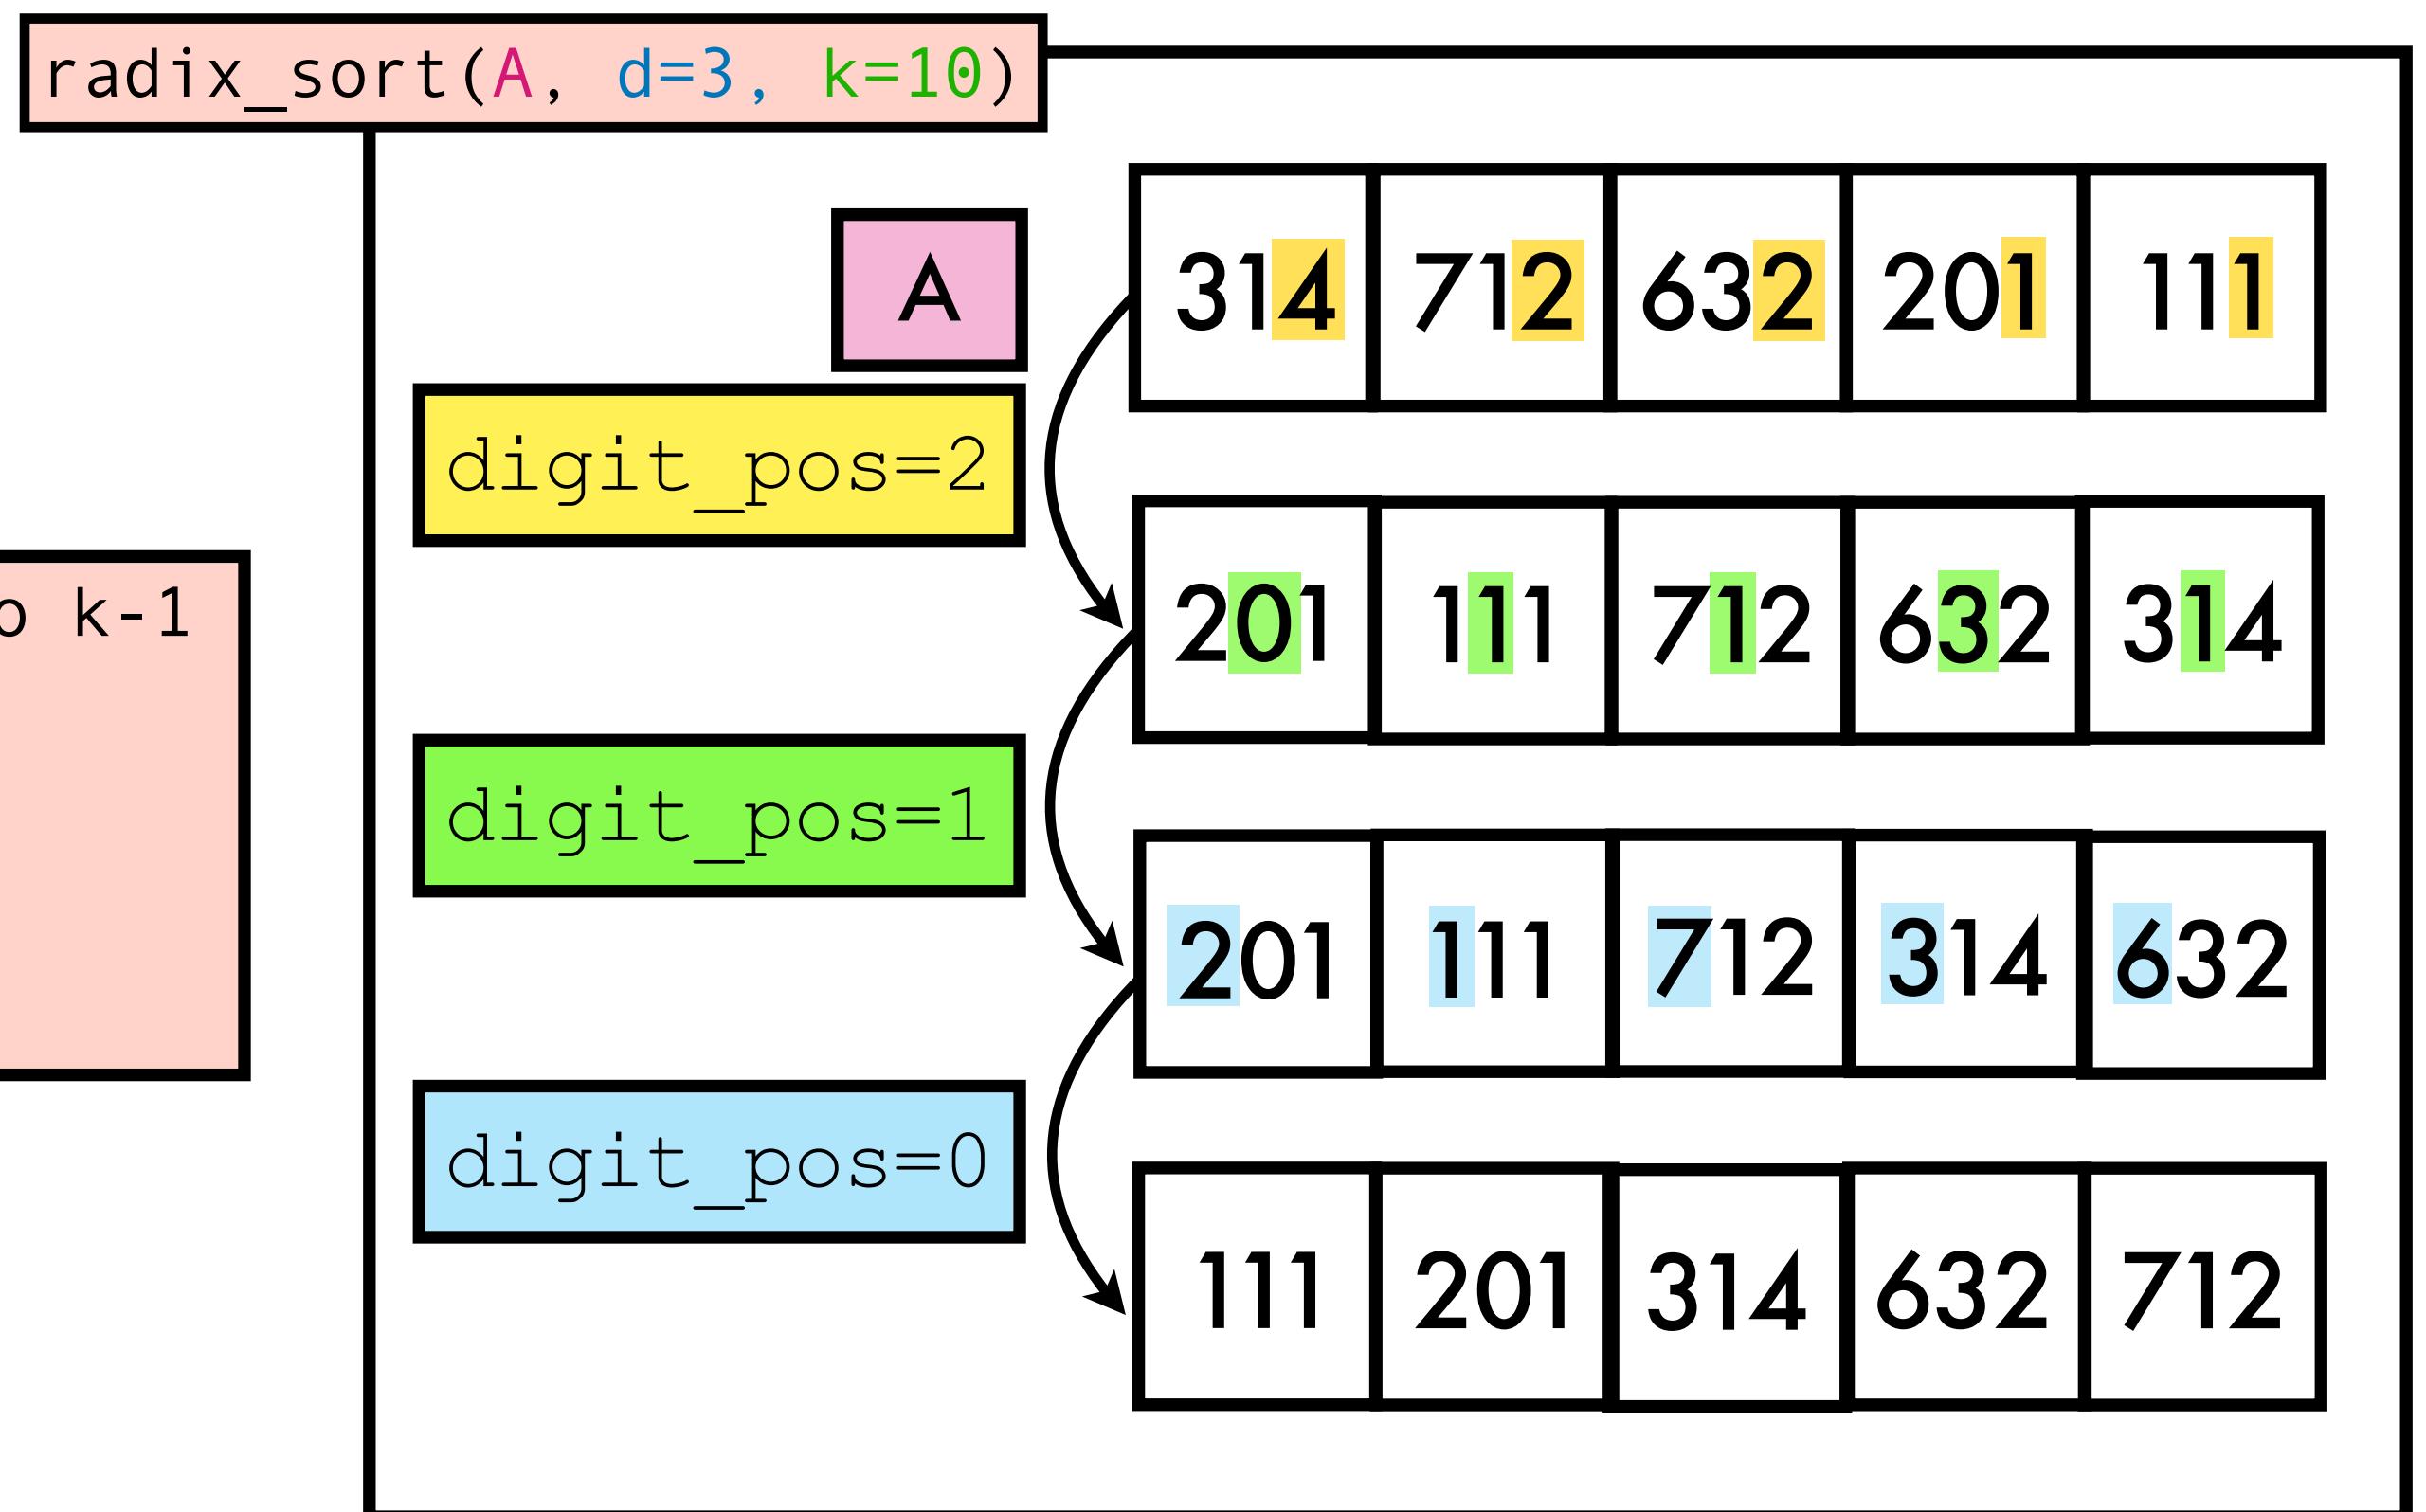
This can be addressed with **cutoff** to **insertion sort**

### Reference:

R. Sedgewick et al., "Radix Sorts", <https://www.cs.princeton.edu/courses/archive/spr07/cos226/lectures/11RadixSort.pdf> (2007)

# LSD Radix Sort Implementation

```
def radix_sort(A, d, k): # keys w. d digits from 0 to k-1
    # loop from least to most significant digit
    for digit_pos in range(d-1, -1, -1):
        # typically use counting sort as stable sort
        stable_sort(A, k, digit_pos)
```



Reference:

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

# Digit Size Selection

The influence of digit size

For  $n$  numbers with  $b$ -bits and  $r \leq b$  (digit size)

LSD radix sort runtime is  $\Theta\left(\frac{b(n + 2^r)}{r}\right)$

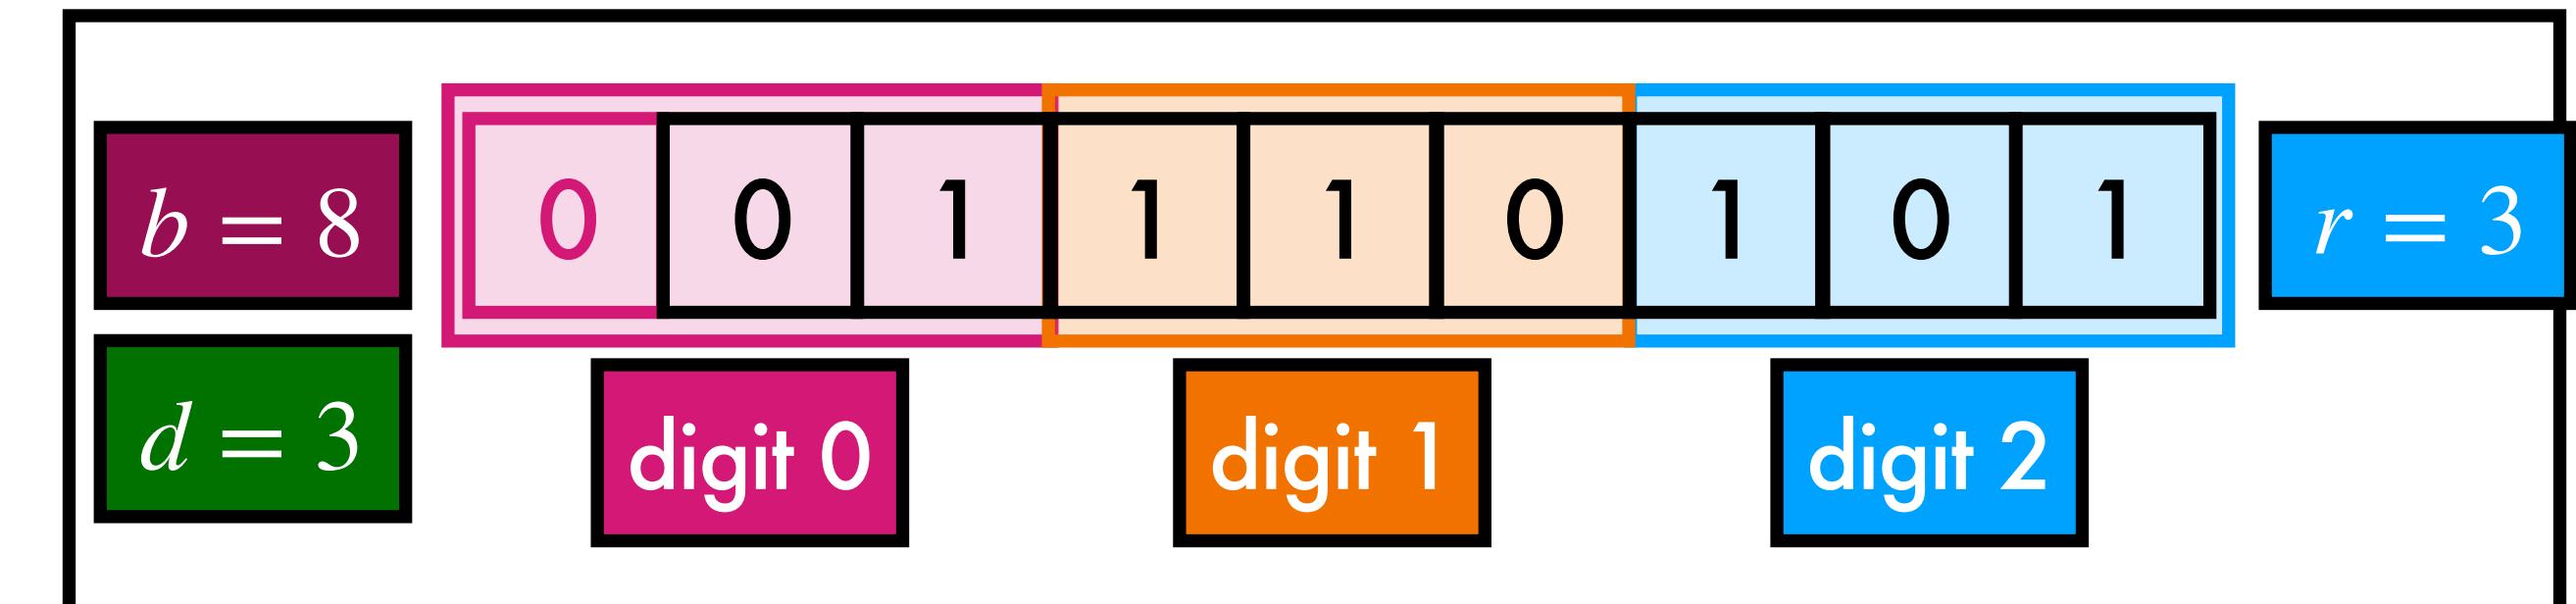
(assumes counting sort in inner loop)

**Proof:** interpret keys as  $d = \lceil b/r \rceil$  digits

where each digit has  $r$  bits

Use counting sort with  $k = 2^r$

Runtime is  $\Theta(d(n + k)) = \Theta\left(\frac{b(n + 2^r)}{r}\right)$



How should we choose digit size ( $r$  bits)?

Goal: minimise  $\frac{b(n + 2^r)}{r}$

If  $b < \lfloor \log n \rfloor$ : pick  $r = b \Rightarrow \frac{b(n + 2^b)}{b} = \Theta(n)$

If  $b \geq \lfloor \log n \rfloor$ : pick  $r = \lfloor \log n \rfloor$  (optimal)

$$\Rightarrow \frac{b(n + 2^{\lfloor \log n \rfloor})}{\lfloor \log n \rfloor} = \Theta\left(\frac{n}{\log n}\right) \text{ (best we can do)}$$

References:

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

# Comparison To Quicksort

## LSD radix sort vs quicksort

Quicksort runtime is typically  $\Theta(n \log n)$

LSD radix sort  $\Theta(dn)$  when  $k = O(n)$

Quicksort - typically more cache-friendly & in-place

LSD radix sort - less cache-friendly & not in-place

Best choice of sorting algorithm will depend on

data distribution

storage requirements

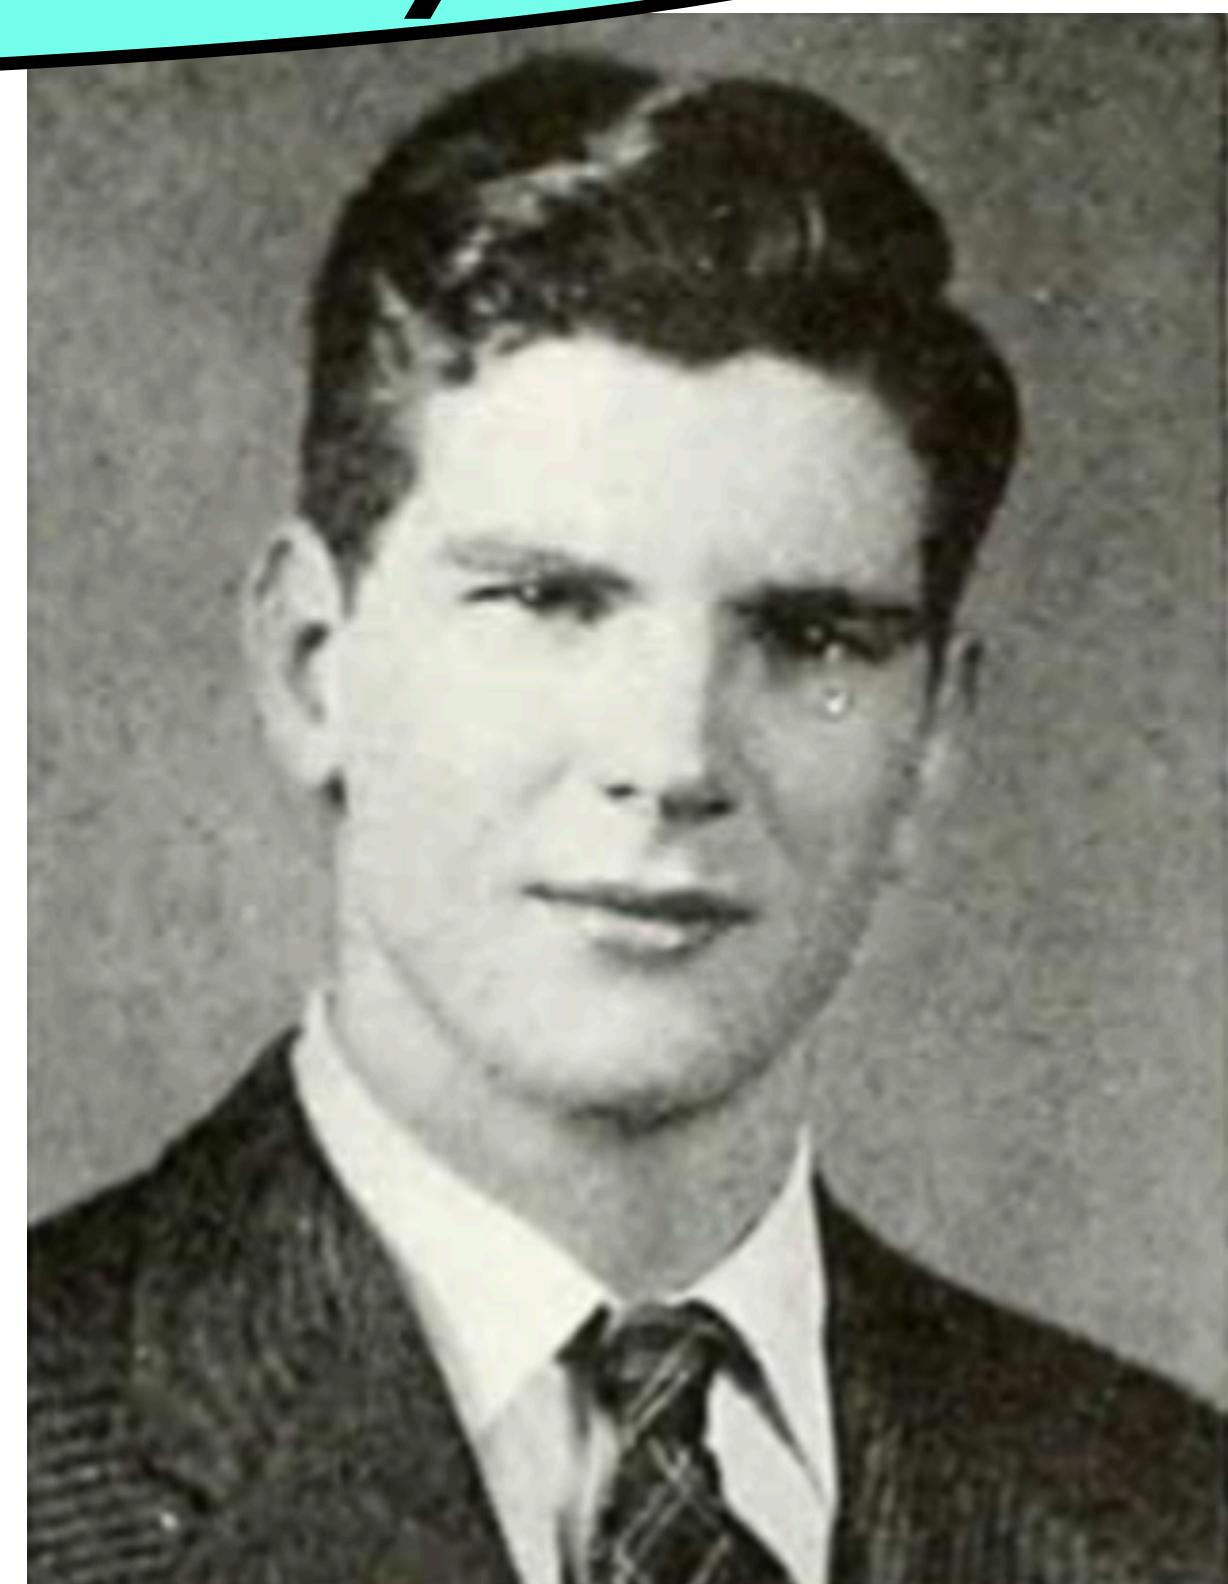
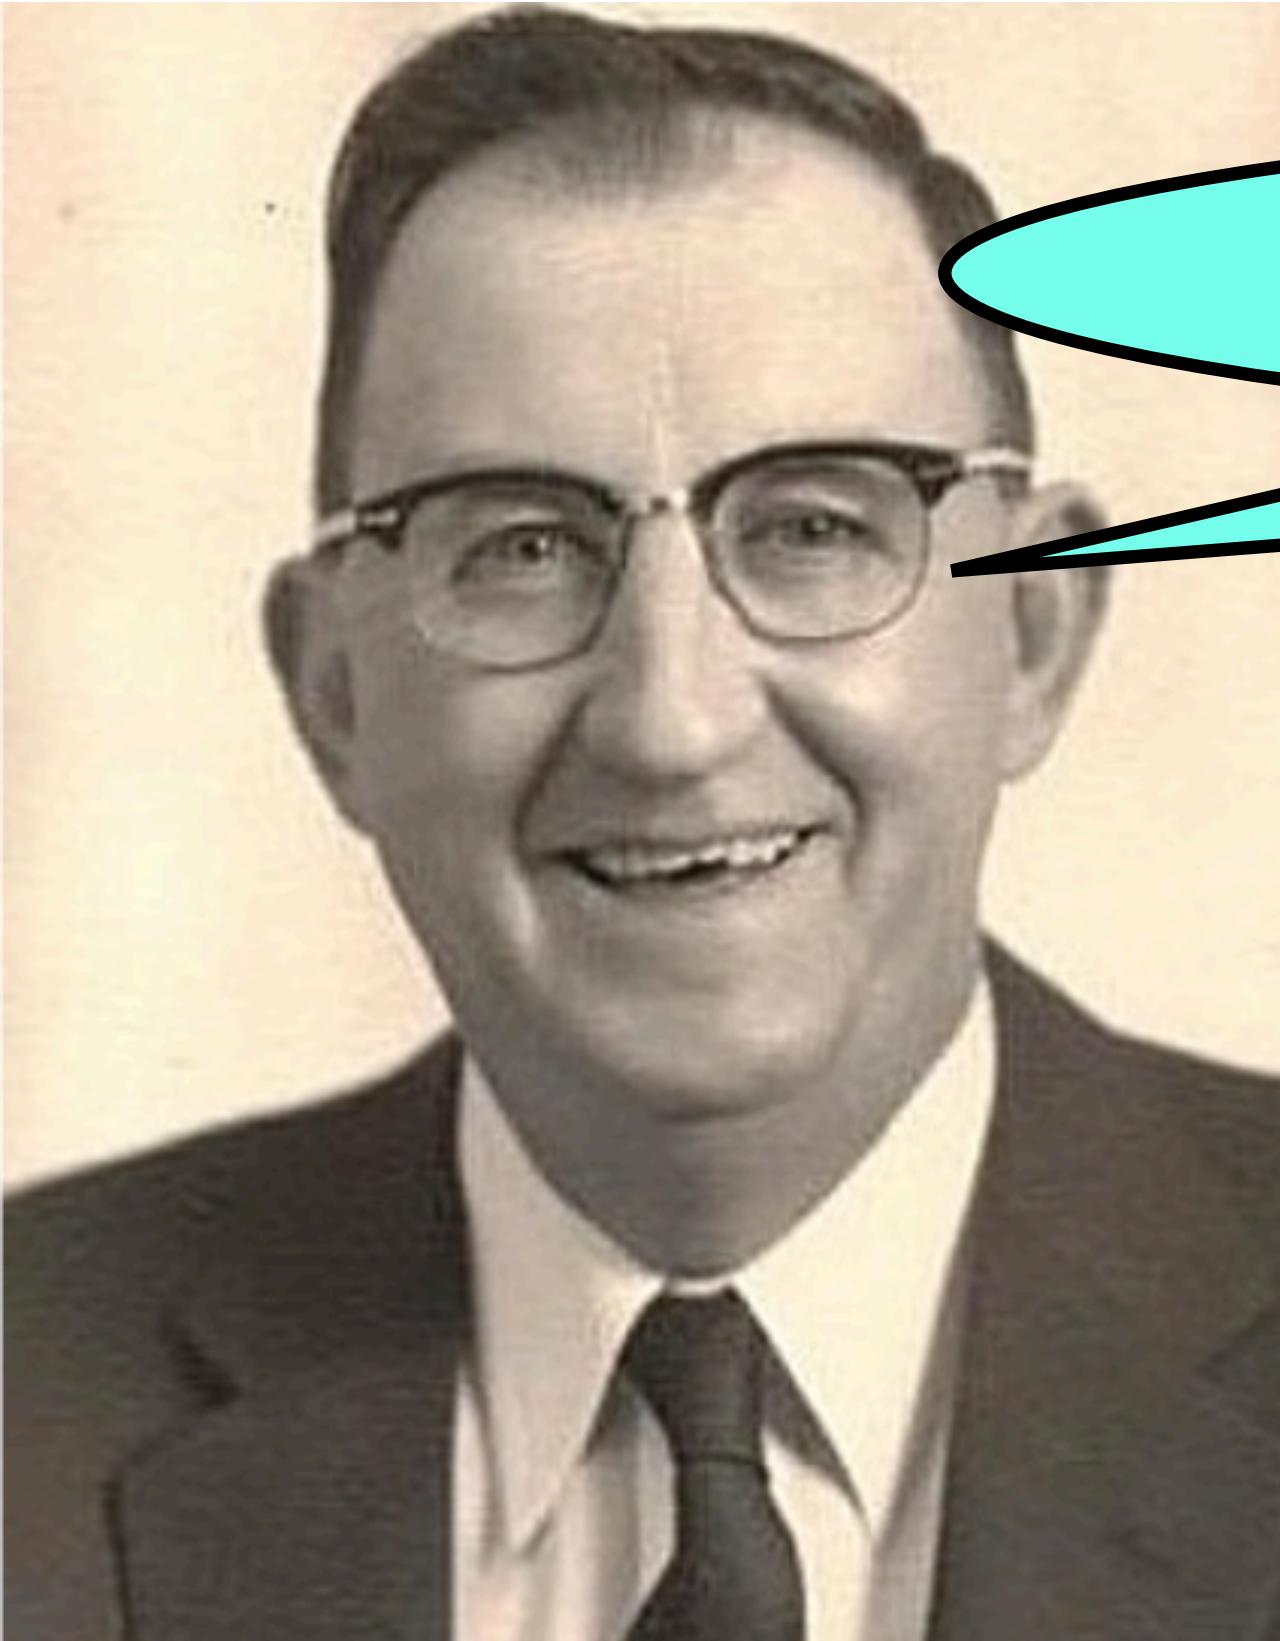
machine details

### References:

R. Sedgewick et al., "Radix Sorts", <https://www.cs.princeton.edu/courses/archive/spr07/cos226/lectures/11RadixSort.pdf> (2007)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

D. E. Knuth, "The art of computer programming, vol. 3: sorting and searching", Chap 5 (1998)



How can we *sort data* efficiently?

(when we know the  
*data distribution*)

# Bucket Sort

useful when *data distribution* is known

Runtime complexity

Average case  $\Theta(n)$

Properties

```
def bucket_sort(A):  
    ...
```

Python Implementation

# Bucket Sort

## Bucket Sort

Distribution sort algorithm with three stages:

1. Scatter distribute keys to buckets
2. Sort sort keys within buckets
3. Gather gather sorted keys (in order)

Effective with knowledge of key distribution

Key idea - E. Isaac & R. Singleton (1955)



### Bucket sort complexity

Assume uniformly distributed keys

Runtime for  $n$  keys and  $b$  buckets:

Average-case:  $\rightarrow \Theta\left(n + \frac{n^2}{b} + b\right)$

Worst-case:  $\rightarrow \Theta(n^2 + b)$  (all keys in same bucket)

Storage:  $\Theta(n)$  for items, sort requires:  $\Theta(n + b)$

Typically, we choose  $b \approx n$ :

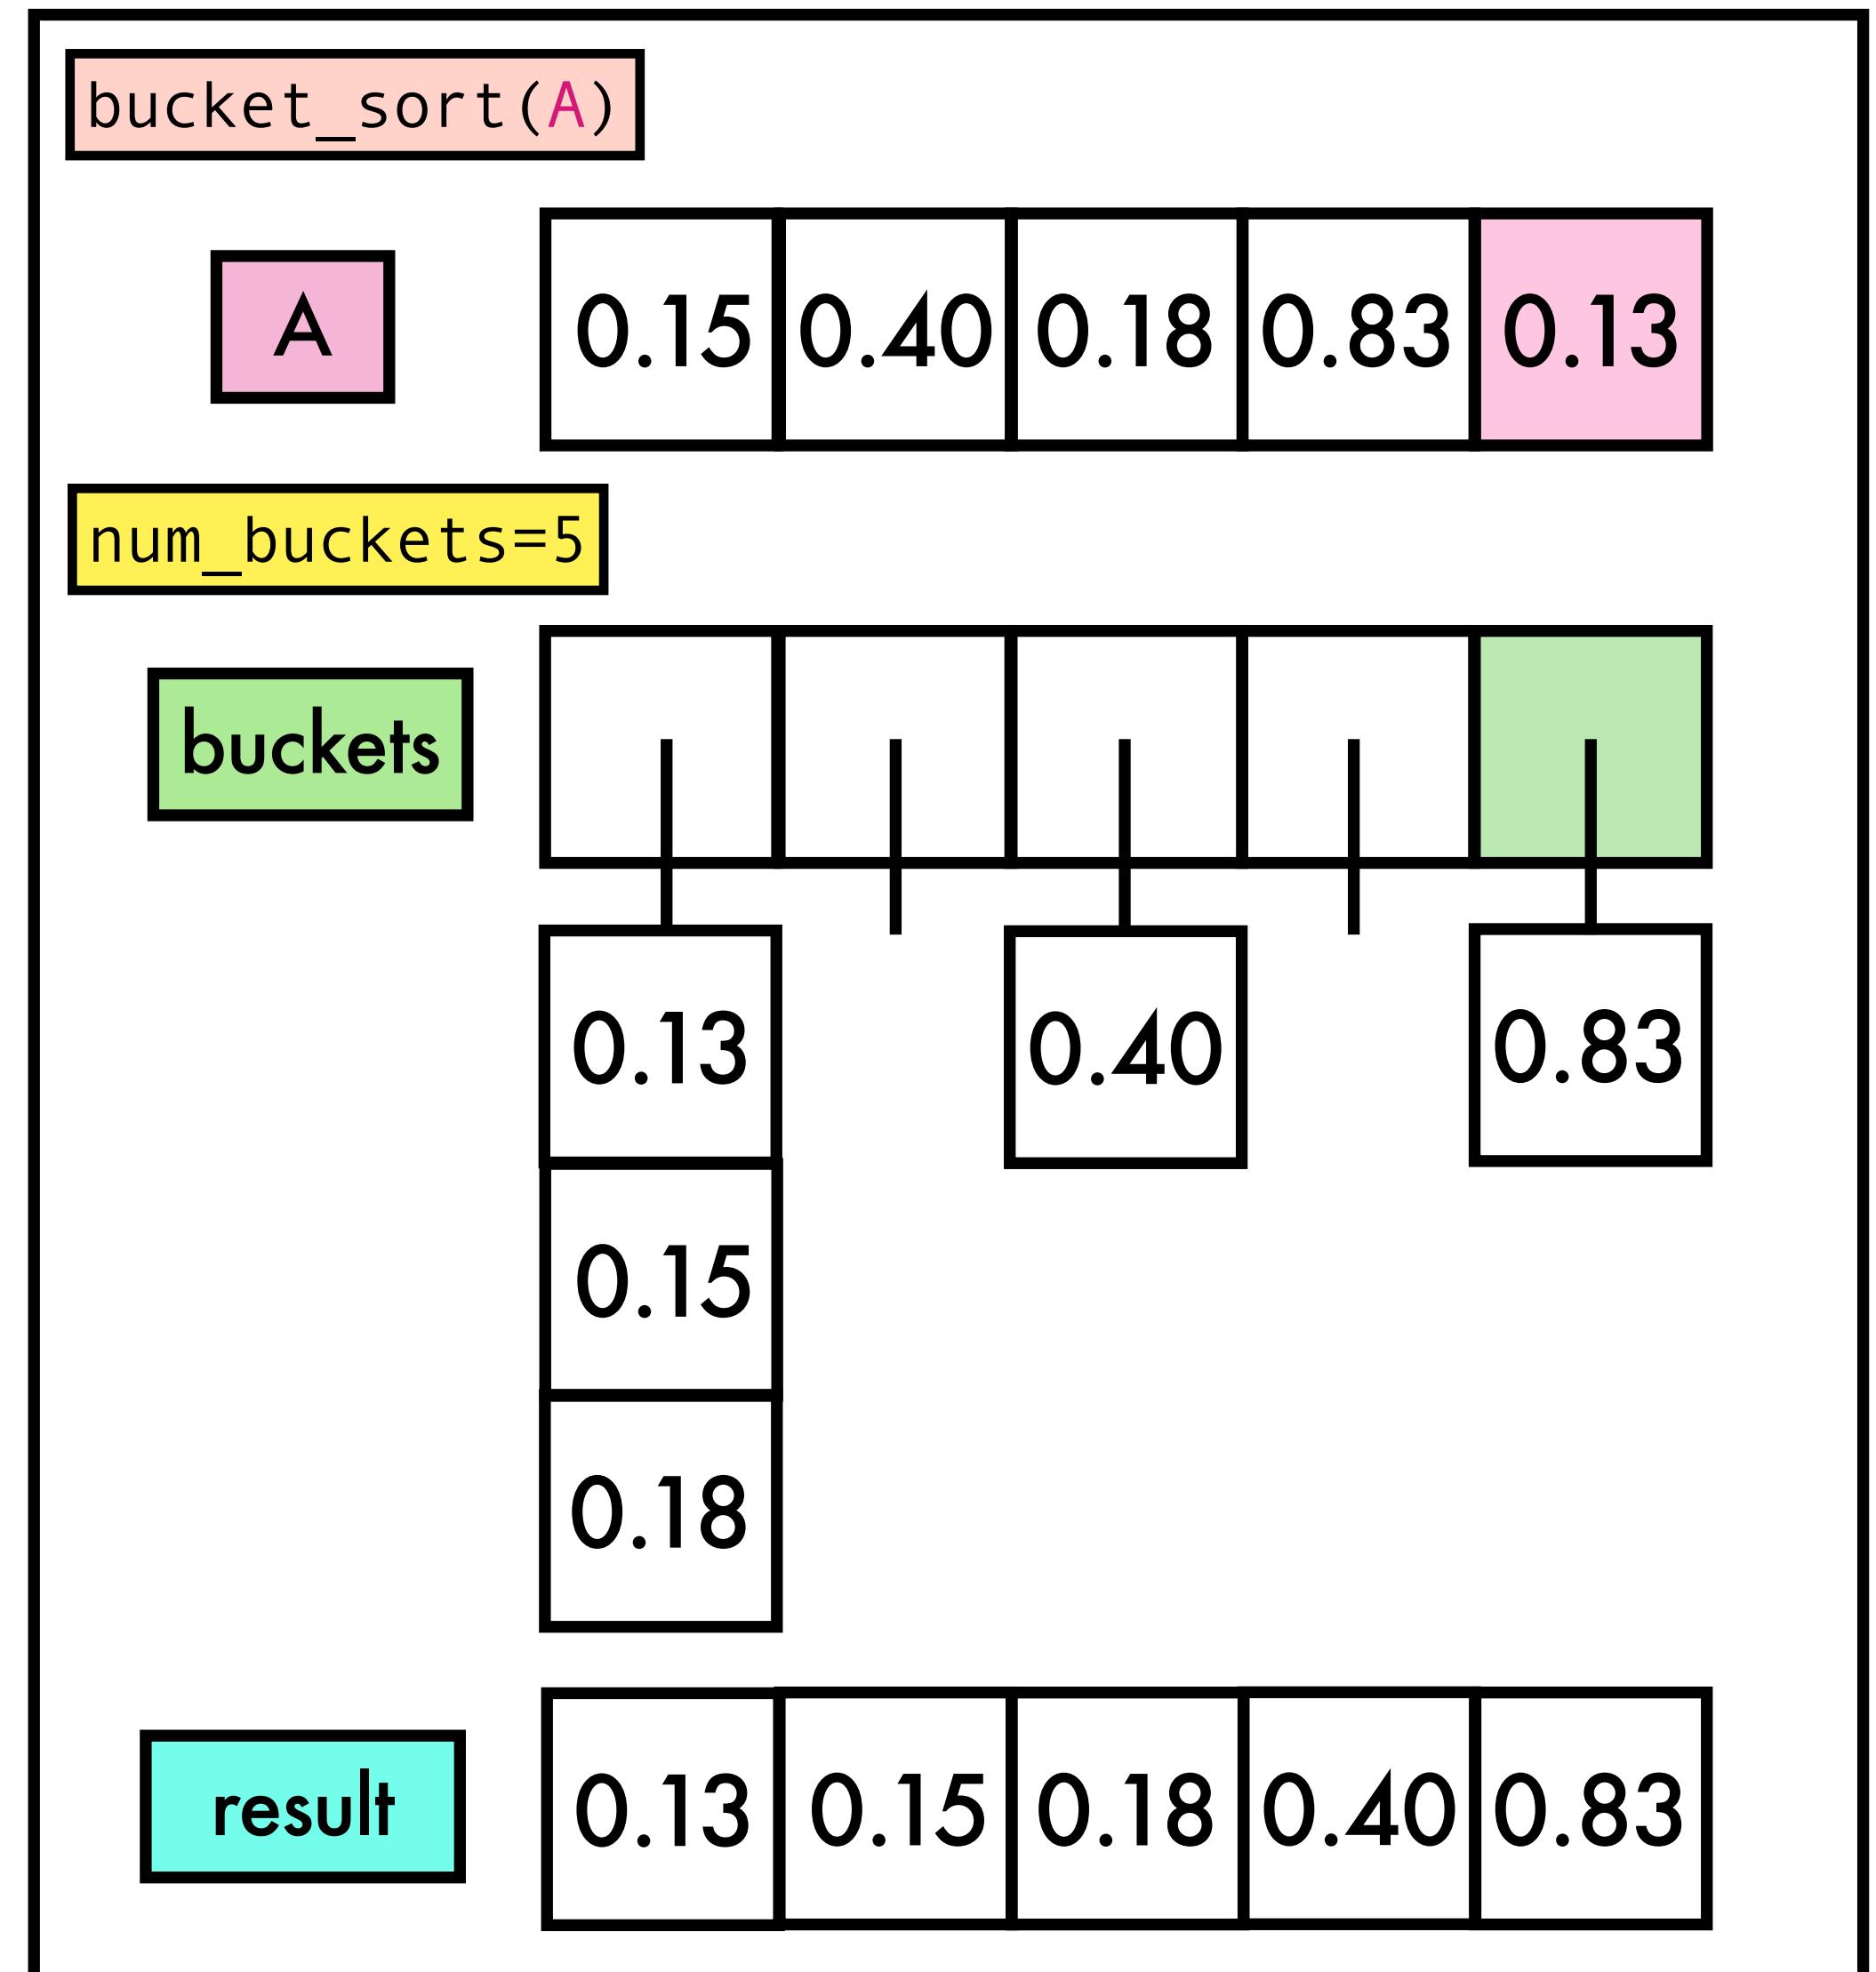
Average-case:  $\rightarrow \Theta(n)$

References/Notes/Image credits:

E. J. Isaac and R. C. Singleton, "Sorting by address calculation", *Journal of the ACM* (1956)  
(Earl Isaac image source) <https://www.finnotes.org/people/earl-isaac>  
(R. Singleton image source) Massachusetts Institute of Technology Yearbook, Vol. 65 (1949)  
[https://en.wikipedia.org/wiki/Bucket\\_sort](https://en.wikipedia.org/wiki/Bucket_sort)

# Bucket Sort Python Implementation

```
def bucket_sort(A: list): # all keys fall in [0, 1)
    num_buckets = len(A)
    buckets = [[] for _ in range(num_buckets)]
    for key in A: # scatter
        buckets[int(num_buckets * key)].append(key)
    for bucket in buckets:
        insertion_sort(bucket)
    return [x for bucket in buckets for x in bucket] # gather
```



Reference:

[https://en.wikipedia.org/wiki/Bucket\\_sort](https://en.wikipedia.org/wiki/Bucket_sort)

T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

# Bucket Sort Runtime Complexity

Runtime for **average case** when num. keys,  $n$  equals num. buckets,  $b$

The **scatter** and **gather** operations involve simple for loops - each are  $\Theta(n)$

Rest of cost is from calling **insertion\_sort** on  $n$  buckets

Let  $m_i$  denote the **number of keys** in **bucket  $i$**

Cost of  $n$  **insertion\_sort** calls:  $\sum_{i=0}^{n-1} O(m_i^2)$

$$\sum_{i=0}^{n-1} m_i = n$$

Take **expectation** w.r.t key distribution  $\mathbb{E}\left[\sum_{i=0}^{n-1} O(m_i^2)\right] = \sum_{i=0}^{n-1} O(\mathbb{E}[m_i^2])$

$$\text{Var}[m_i] = \mathbb{E}[m_i^2] - \mathbb{E}^2[m_i]$$

$$\mathbb{E}[m_i^2] = \text{Var}[m_i] + \mathbb{E}^2[m_i] = (1 - 1/n) + 1 = 2 - 1/n$$

Expected cost of  $n$  **insertion\_sort** calls:  $\sum_{i=0}^{n-1} O(2 - 1/n) = O(n)$

$m_i$  has **binomial distribution**  $\text{Bin}(n, p)$

Probability of falling in **bucket  $i$** :  $p = 1/n$

$$\mathbb{E}[m_i] = n \cdot 1/n = 1$$

$$\text{Var}[m_i] = 1 - 1/n$$

References:

(CLRS) T. Cormen et al., "Introduction to algorithms", Chap 8, MIT press (2022)

For the general case when  $n \neq b$ , see [https://en.wikipedia.org/wiki/Bucket\\_sort](https://en.wikipedia.org/wiki/Bucket_sort)

Worst case (all in 1 bucket) - degrades to **insertion\_sort**:  $\Theta(n^2)$

# A Variant Of Bucket Sort

a "surprisingly easy improvement"

Switching the execution order

Conventional bucket sort:      Scatter      Sort      Gather

Variation of bucket sort:      Scatter      Gather      Sort

This can be faster (due to greater cache-friendliness)

Note: whether this helps depends on implementation/hardware details

References:

E. Corwin et al., "Sorting in linear time - variations on the bucket sort", Journal of Computing Sciences in Colleges (2004)